Detecting Performance Anti-patterns for Applications Developed Using Object-Relational Mapping

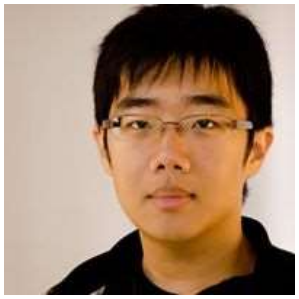

Tse-Hsun(Peter) Chen

Weiyi Shang

Zhen Ming Jiang

Ahmed E. Hassan

Queen's University

Queen's University

YORK UNIVERSITÉ UNIVERSITY

Queen's University

Mohamed Nasser, Parminder Flora

BlackBerry

Databases are essential in large-scale software systems

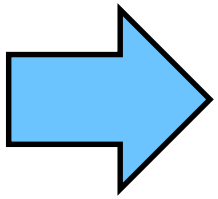

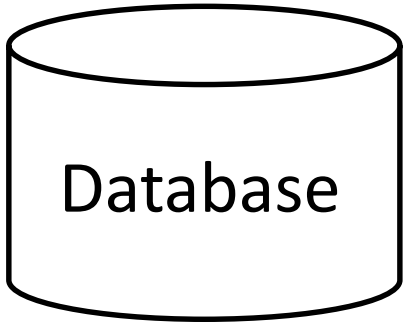

Application developers work with objects

More intuitive if we can map objects directly to DB

Object-Relational Mapping eliminates the gap between objects and SQL

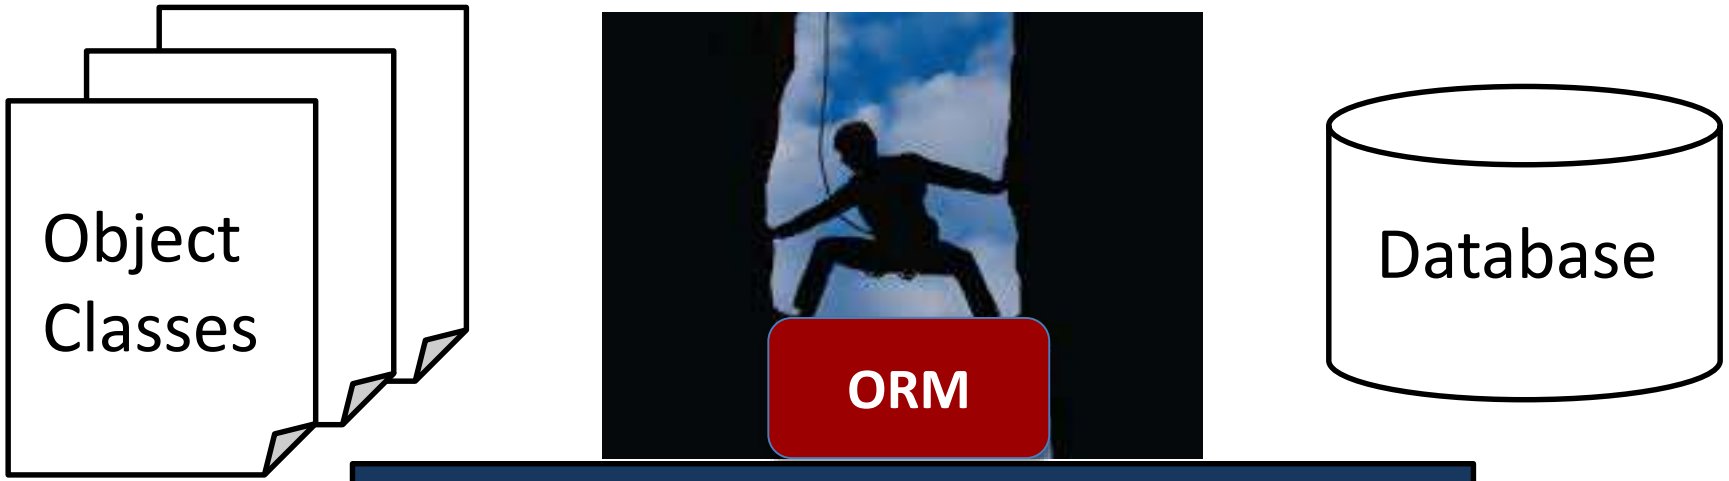

Much less code and shorter development time

Problem of using raw SQLs

- Lots of **boilerplate code**
- Need to **manage object-DB translations** manually

ORM is widely used in practice

- Java Hibernate has ***more than 8 million*** downloads
- In 2013, ***15% of the 17,000*** Java developer jobs require ORM experience (dice.com)

Different ORM technologies

An example class with ORM code

User class is mapped to "*user*" table in DB

id is mapped to the column *"id"* in the user table

A user can belong to multiple teams

Eagerly retrieve associated teams when retrieving a user object

User.java

```
@Entity
@Table(name = "user")
public class User{
    @Column(name="id")
    private int id;

    @Column(name="name")
    String userName;

@OneToMany(fetch=FetchType.EAGER)
    List<Team> teams;

    public void setName(String n){
        userName = n
    }
... other getter and setter methods
```

Accessing the database using ORM

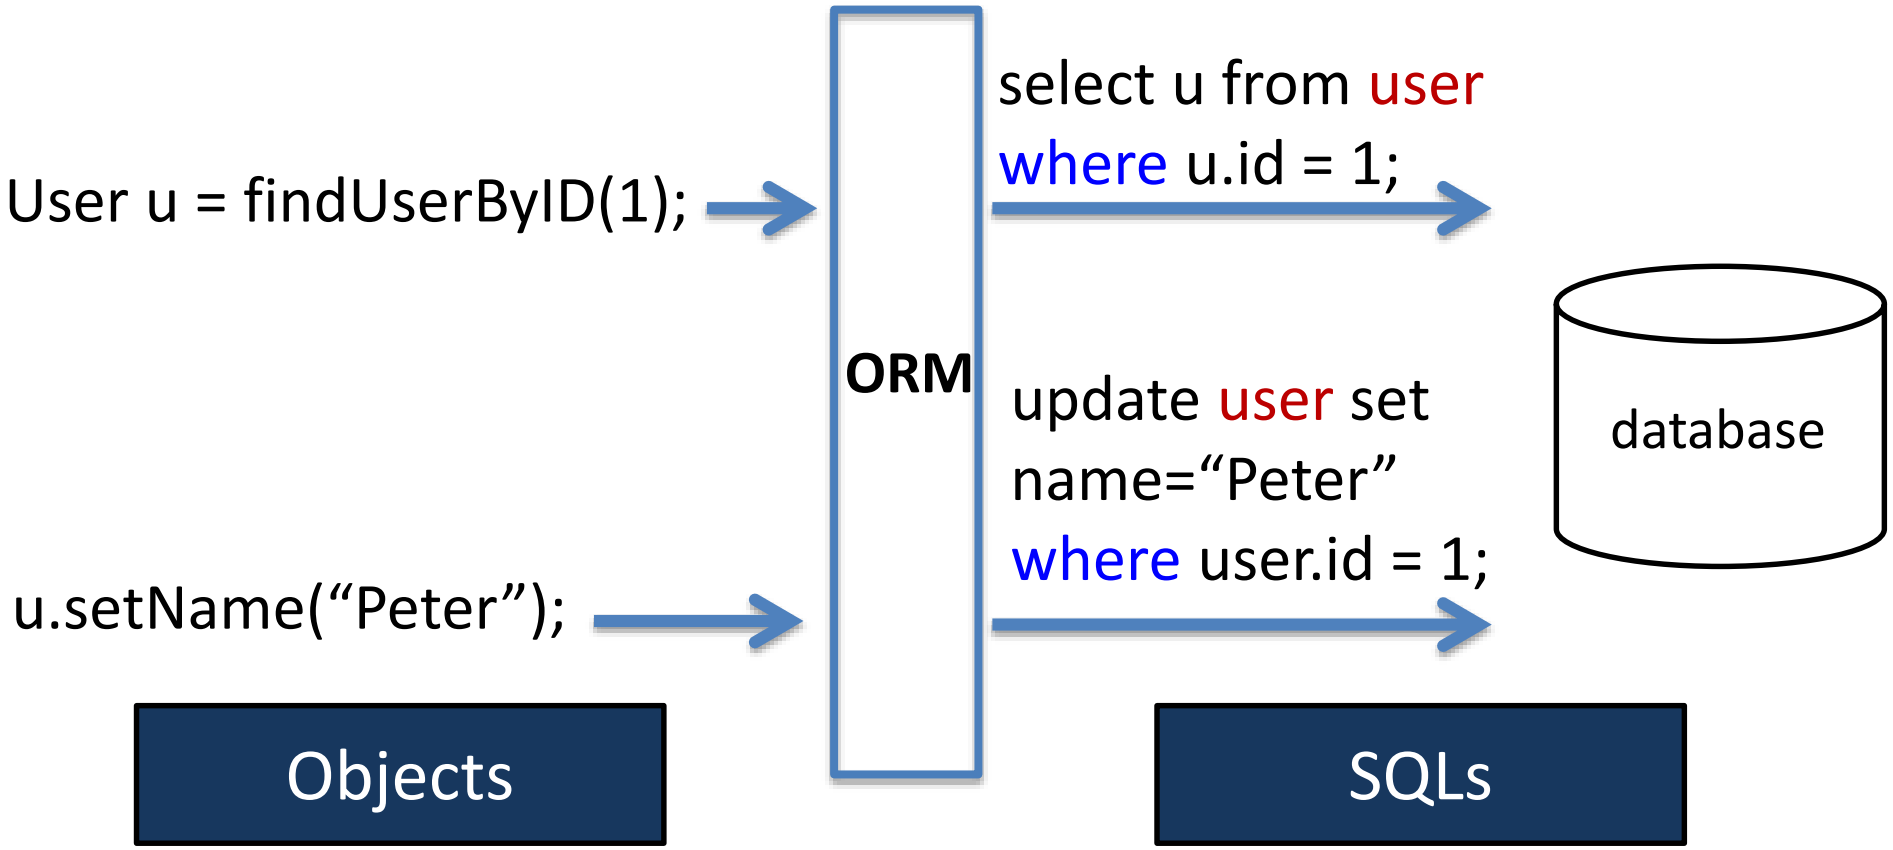

Developers may not be aware of database access

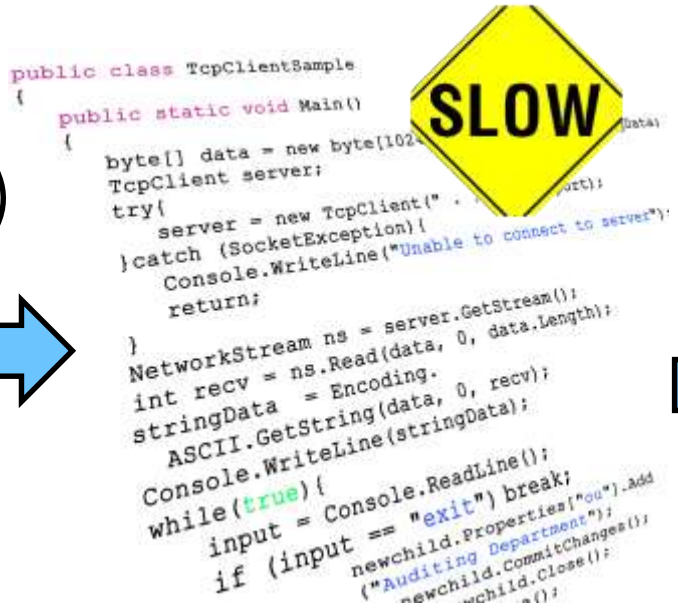

ORM code with performance anti-patterns

Bad system performance

The performance difference can be LARGE!

Performance anti-pattern detection framework

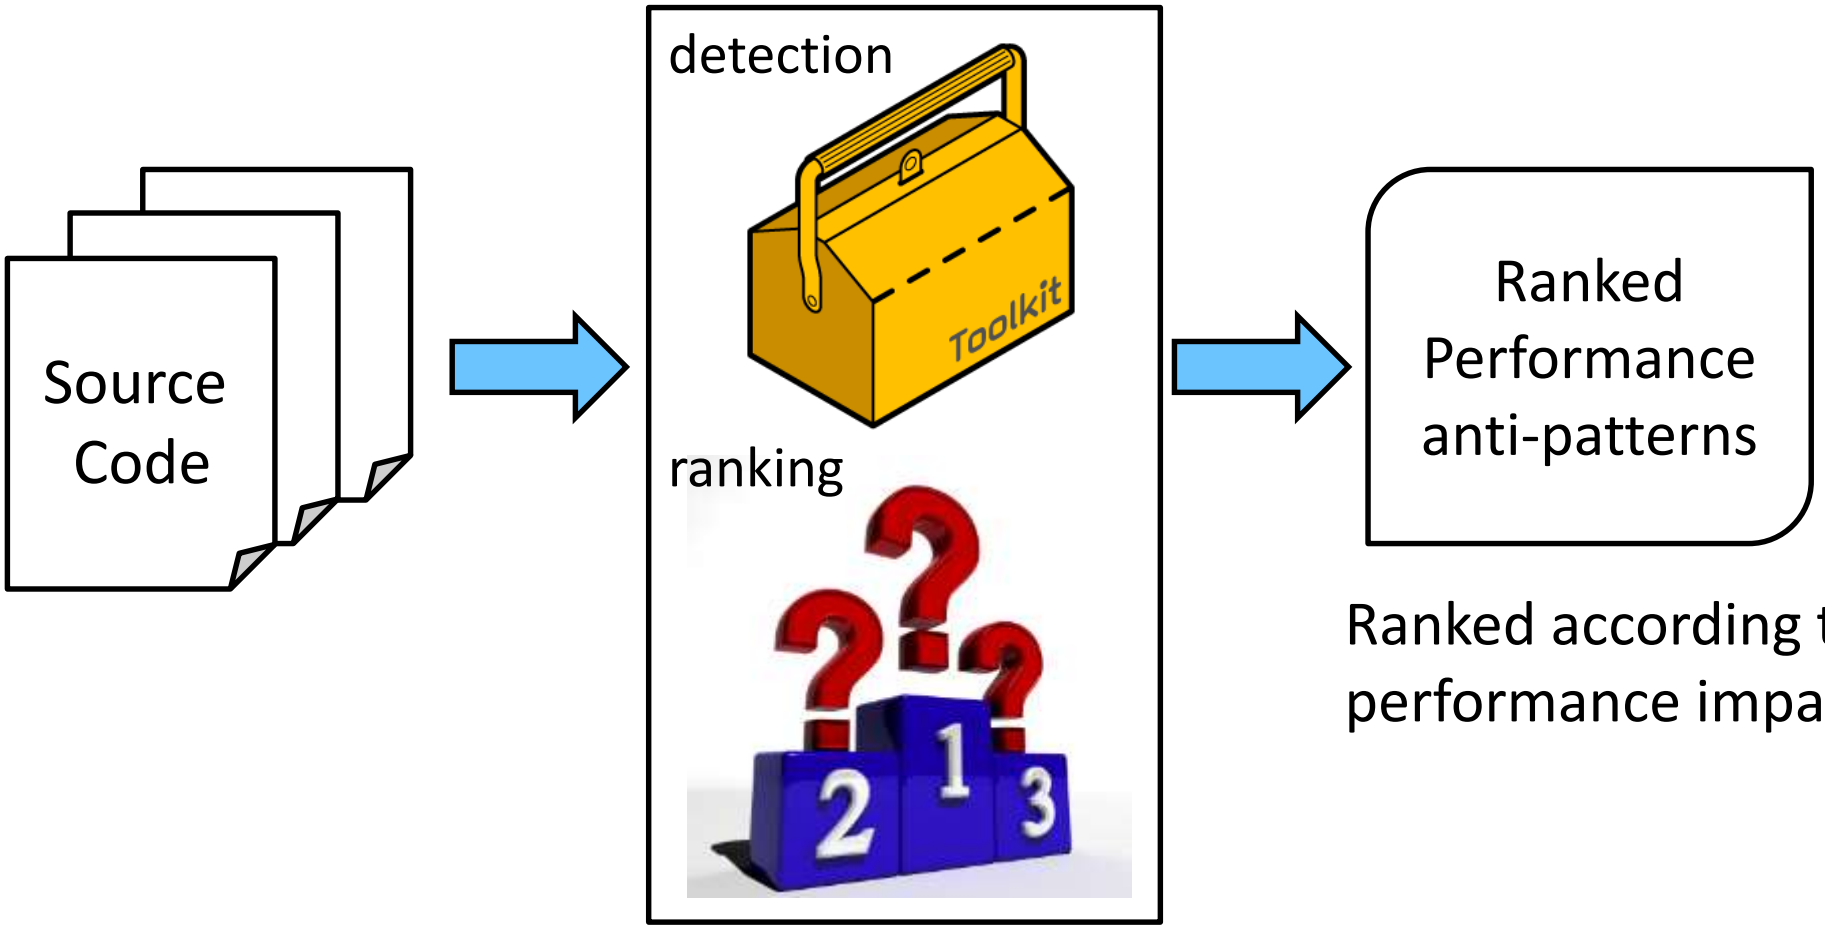

Performance anti-pattern detection and ranking framework

Performance anti-pattern detection framework

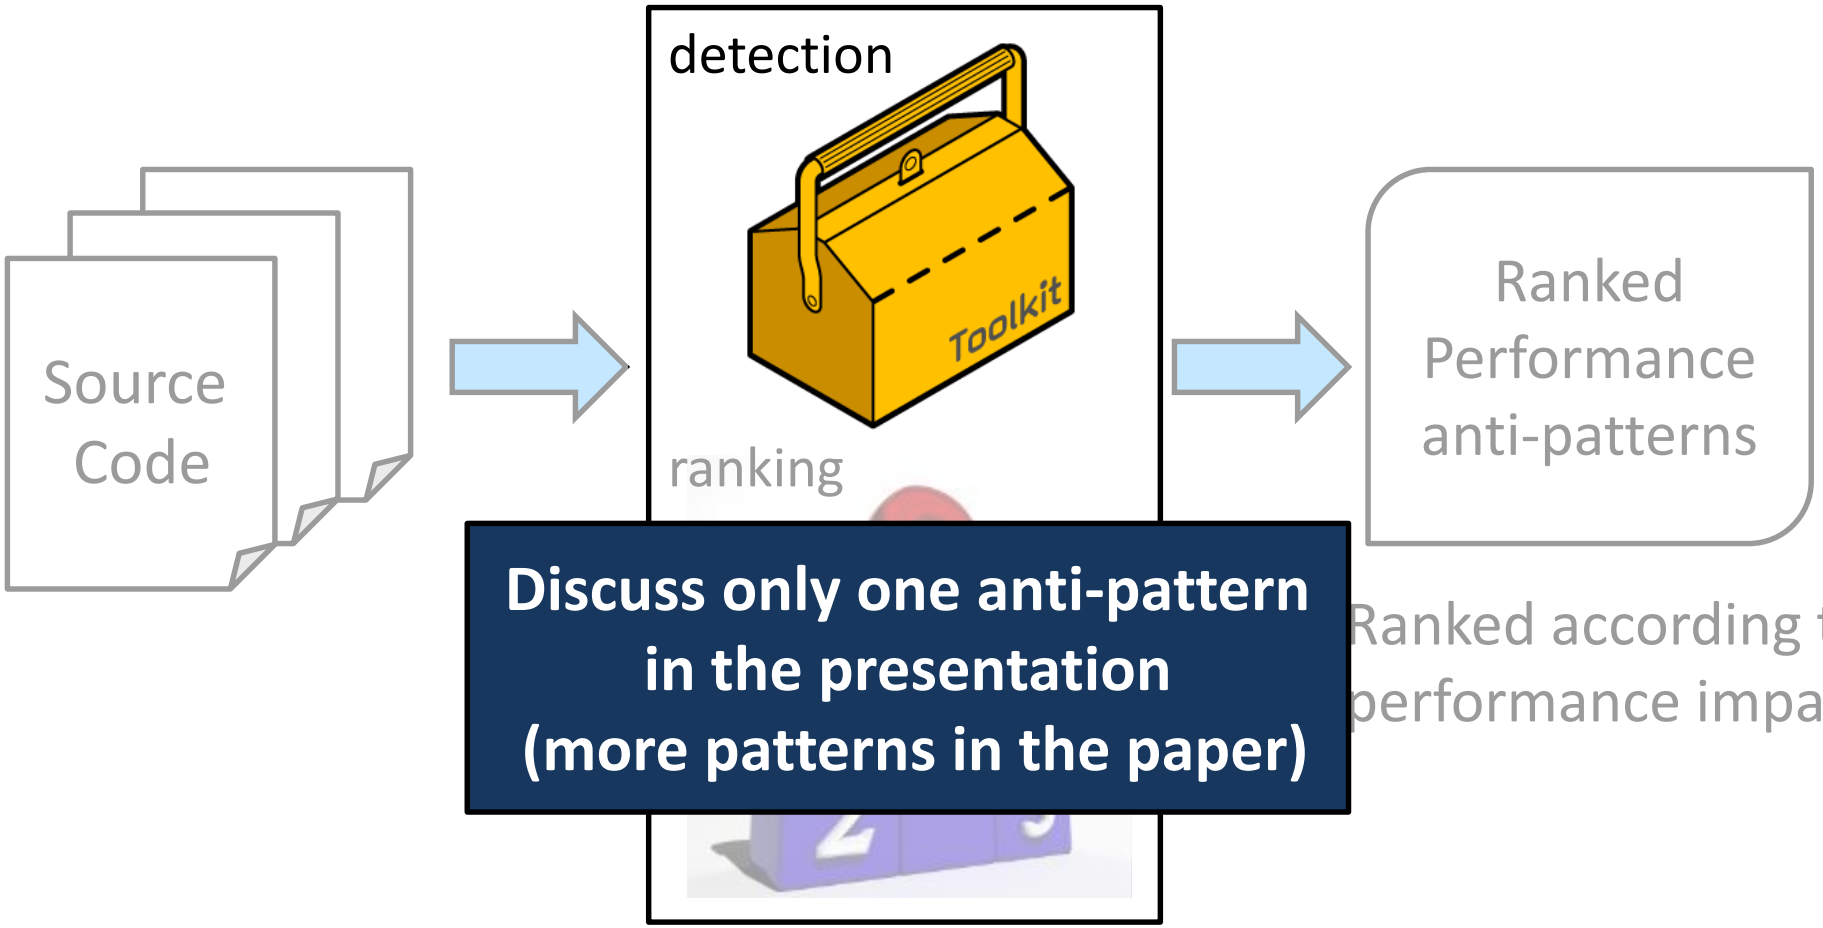

Performance anti-pattern detection and ranking framework

ORM excessive data anti-pattern

```
Class User{
        @EAGER
        List<Team> teams;
}
```

Eagerly retrieve teams from DB

Objects

```
User u = findUserById(1);
u.getName();
EOF
```

SQL

User Table *join* **Team Table**

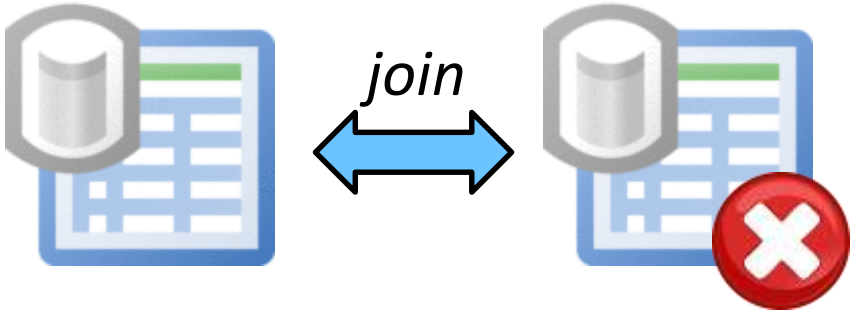

Team data is never used!

Detecting excessive data using static analysis


```
Class User{
        @EAGER
        List<Team> teams;
}
```

First find all the objects that eagerly retrieve data from DB

```
User user = findUserByID(1);
```

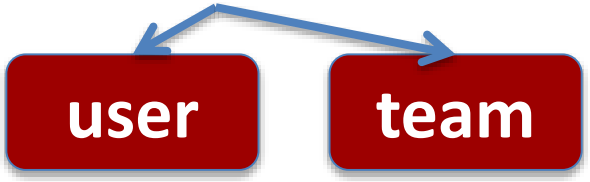

Identify all the data usages of objects

```
user.getName();
```

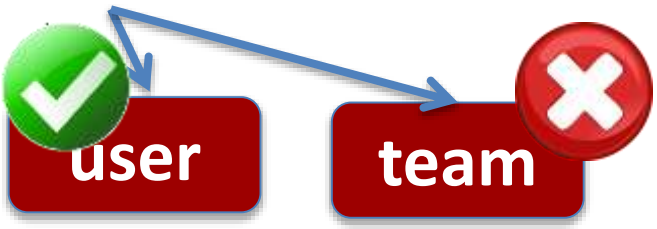

Check if the retrieved data is ever used

Performance anti-pattern detection framework

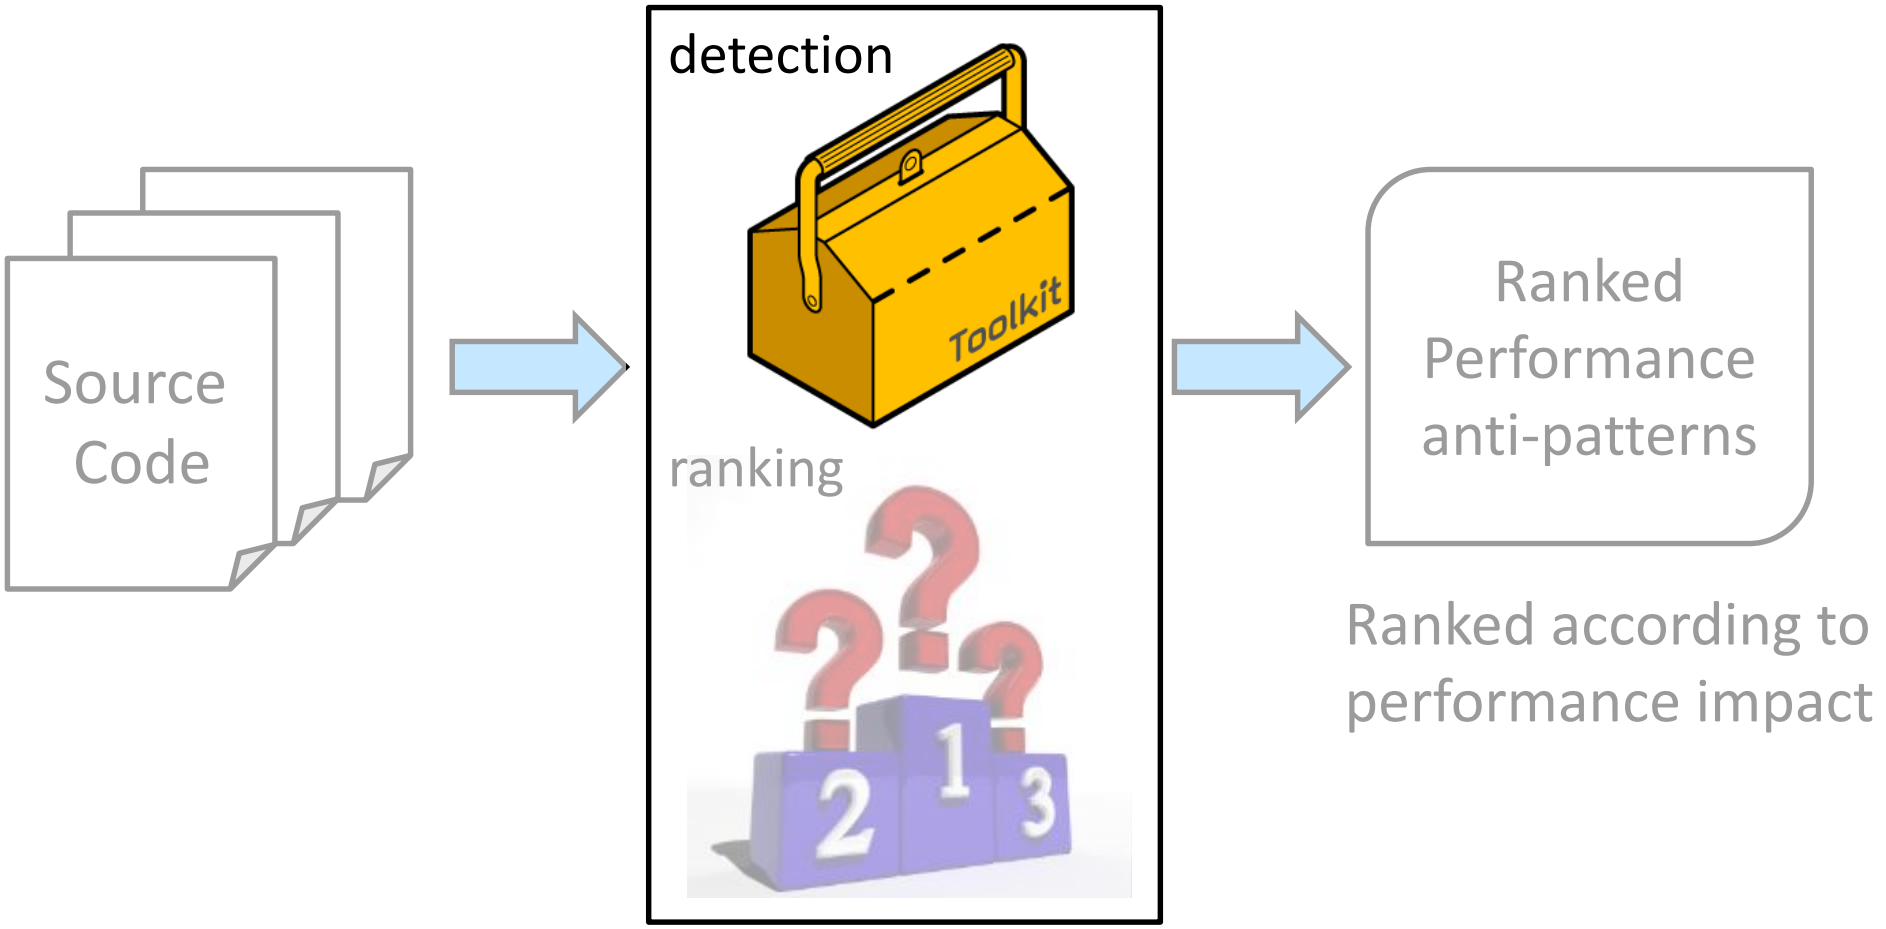

Performance anti-pattern detection and ranking framework

Performance anti-pattern detection framework

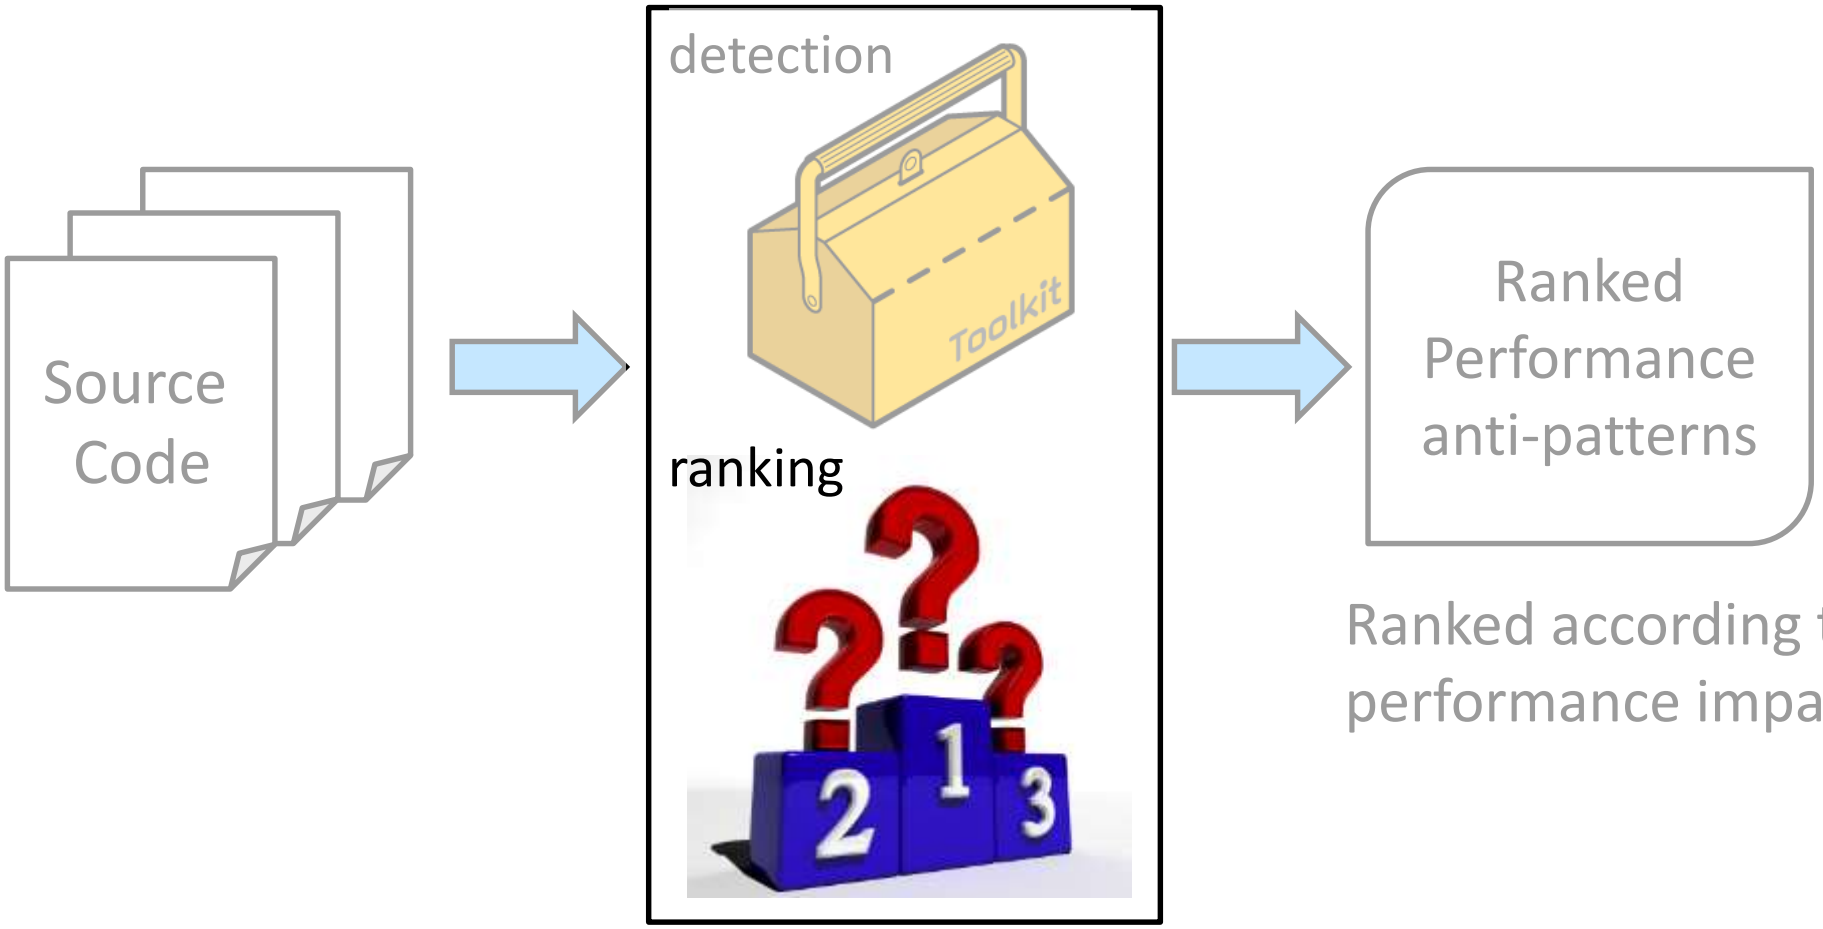

Performance anti-pattern detection and ranking framework

Performance anti-patterns have different impacts

User user_in_1_team = findUserByID(1);

Retrieving ***1 user*** and ***1 team***

User user_in_100_teams = findUserByID(100);

Retrieving ***1 user*** and ***100 teams!***

One can only reveal performance impact by execution

Measuring the impact using repeated measurements and effect sizes

Performance measurements are unstable:

We repeat each test 30 times to obtain stable measurement results

Size of performance impact is not defined:

We use ***effect sizes (Cohen's D)*** to measure the performance impact

Effect sizes =

$$\begin{cases} \text{trivial} & \text{if } Cohen's\ d \leq 0.2 \\ \text{small} & \text{if } 0.2 < Cohen's\ d \leq 0.5 \\ \text{medium} & \text{if } 0.5 < Cohen's\ d \leq 0.8 \\ \text{large} & \text{if } 0.8 < Cohen's\ d \end{cases}$$

Studied systems and detection results

Large open-source e-commence system
> 1,700 files
> 206K LOC

482 excessive data

Enterprise system
> 3,000 files
> 300K LOC

> 10 excessive data

Spring open-source system
Online system for a pet clinic
51 files
3.3K LOC

10 excessive data

Research questions

Performance impact

Ranks of the anti-patterns at different scales

Research questions

Performance impact

Ranks of the anti-patterns at different scales

Assessing anti-pattern impact by fixing the anti-patterns

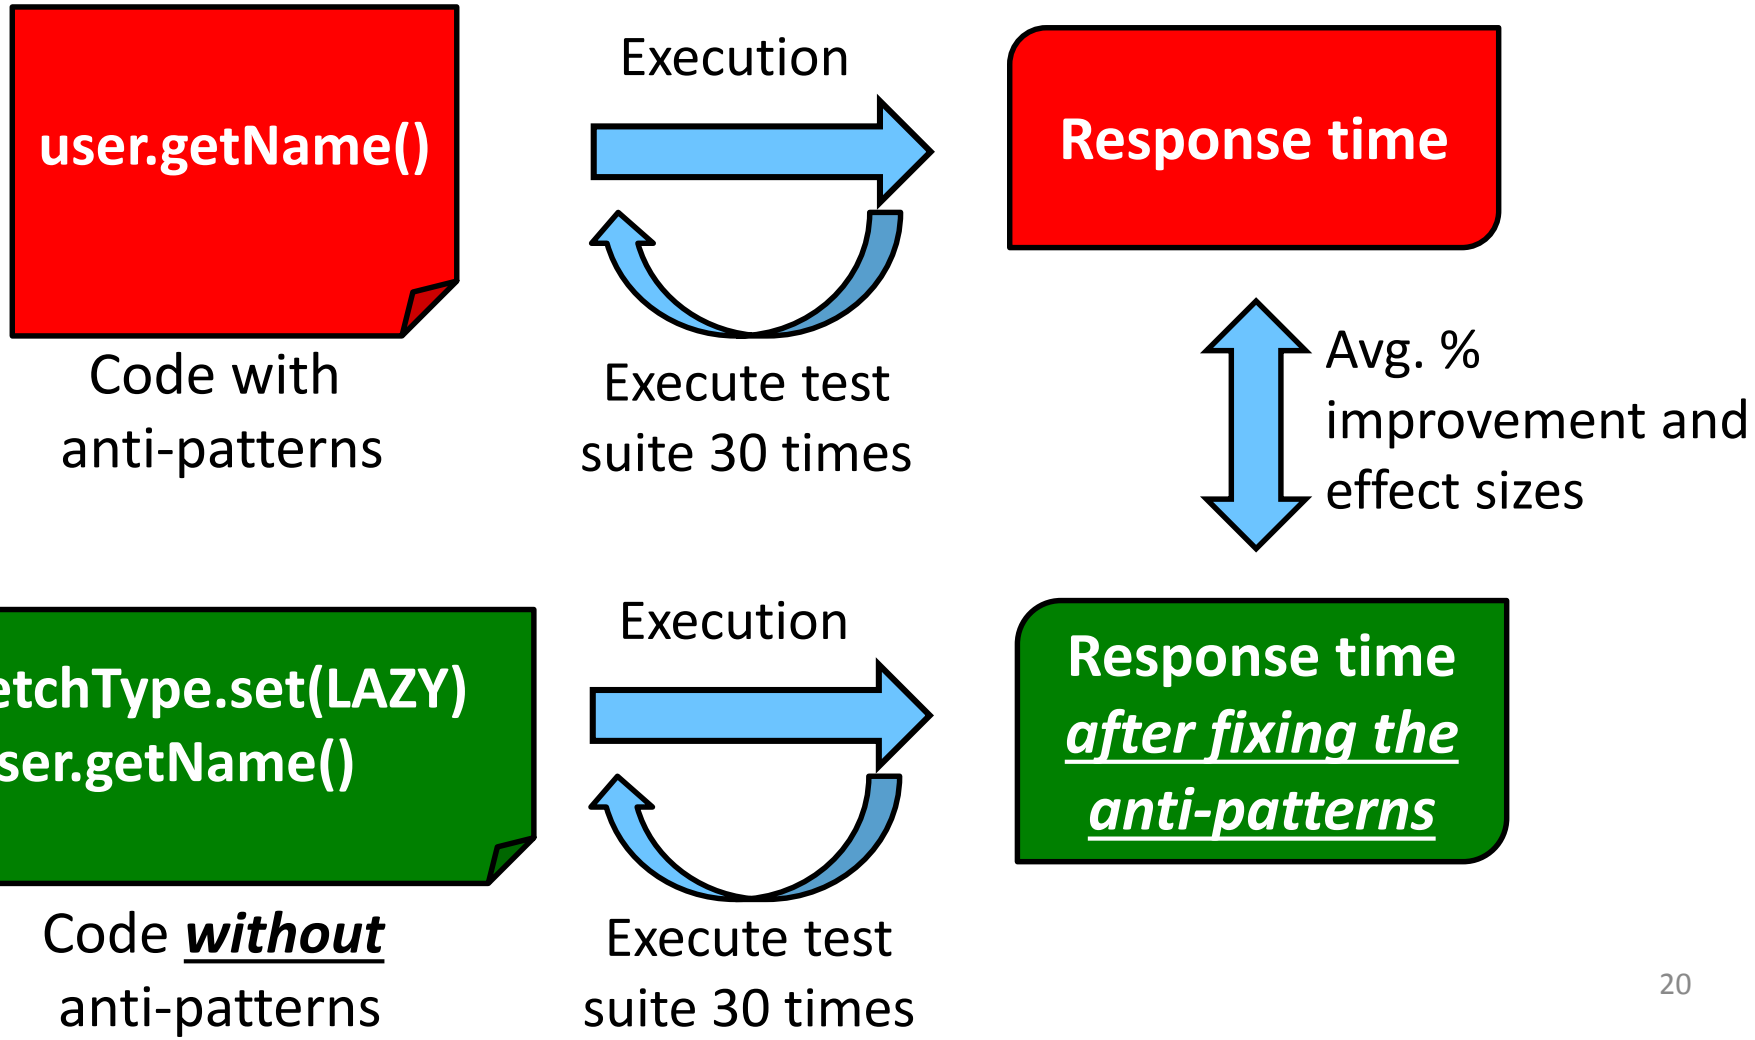

Performance anti-patterns have medium to large effect sizes

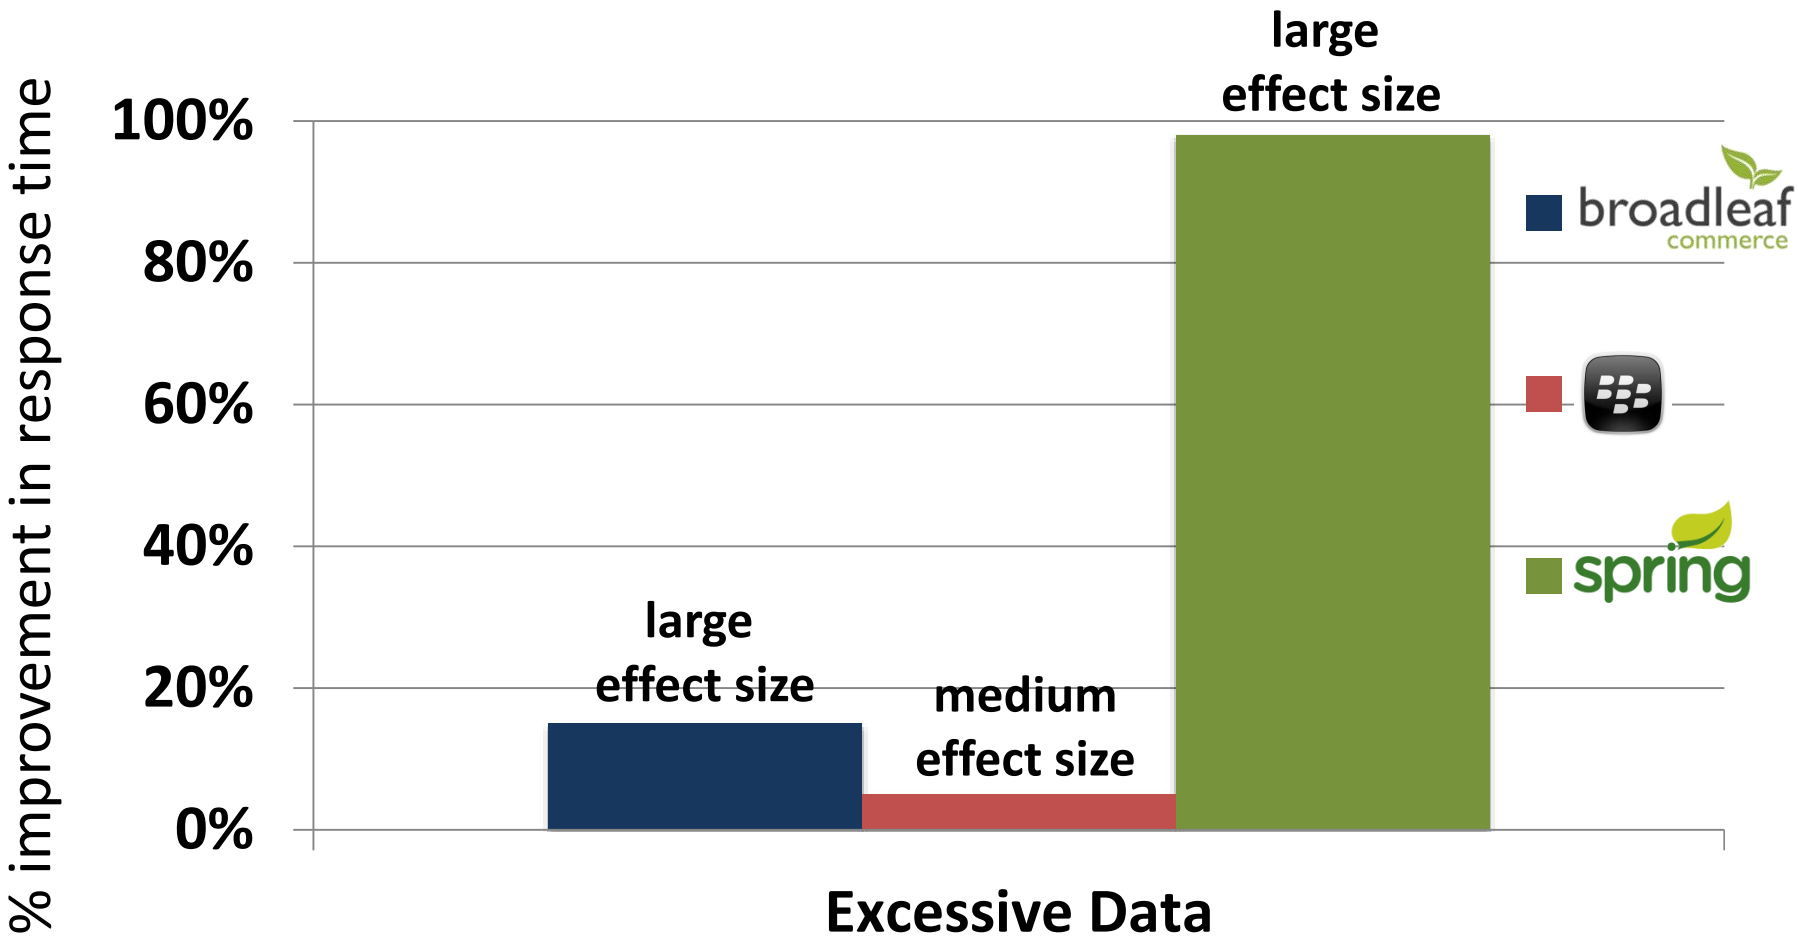

Research questions

Performance impact

Ranks of the anti-patterns at different scales

Removing anti-pattern improves response by ~35%

Research questions

Performance impact

Ranks of the anti-patterns at different scales

Removing anti-pattern improves response by ~35%

Performance problems usually arise under large load

Performance problems revealed at small scales may be more serious

Different input scales

Input scales may have ***exponential*** effects on performance

Performance at different input scales

We should first fix the anti-patterns that have larger effects at smaller scales

Comparing ranked anti-patterns at different data scales

Small size input

Large size input

detection

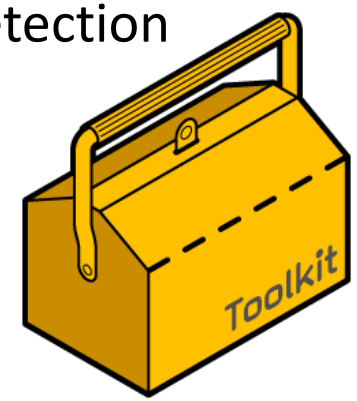

ranking

Ranked Performance anti-patterns from small data

?

Ranked Performance anti-patterns from large data

Anti-patterns have large effects on performance even at smaller data scales

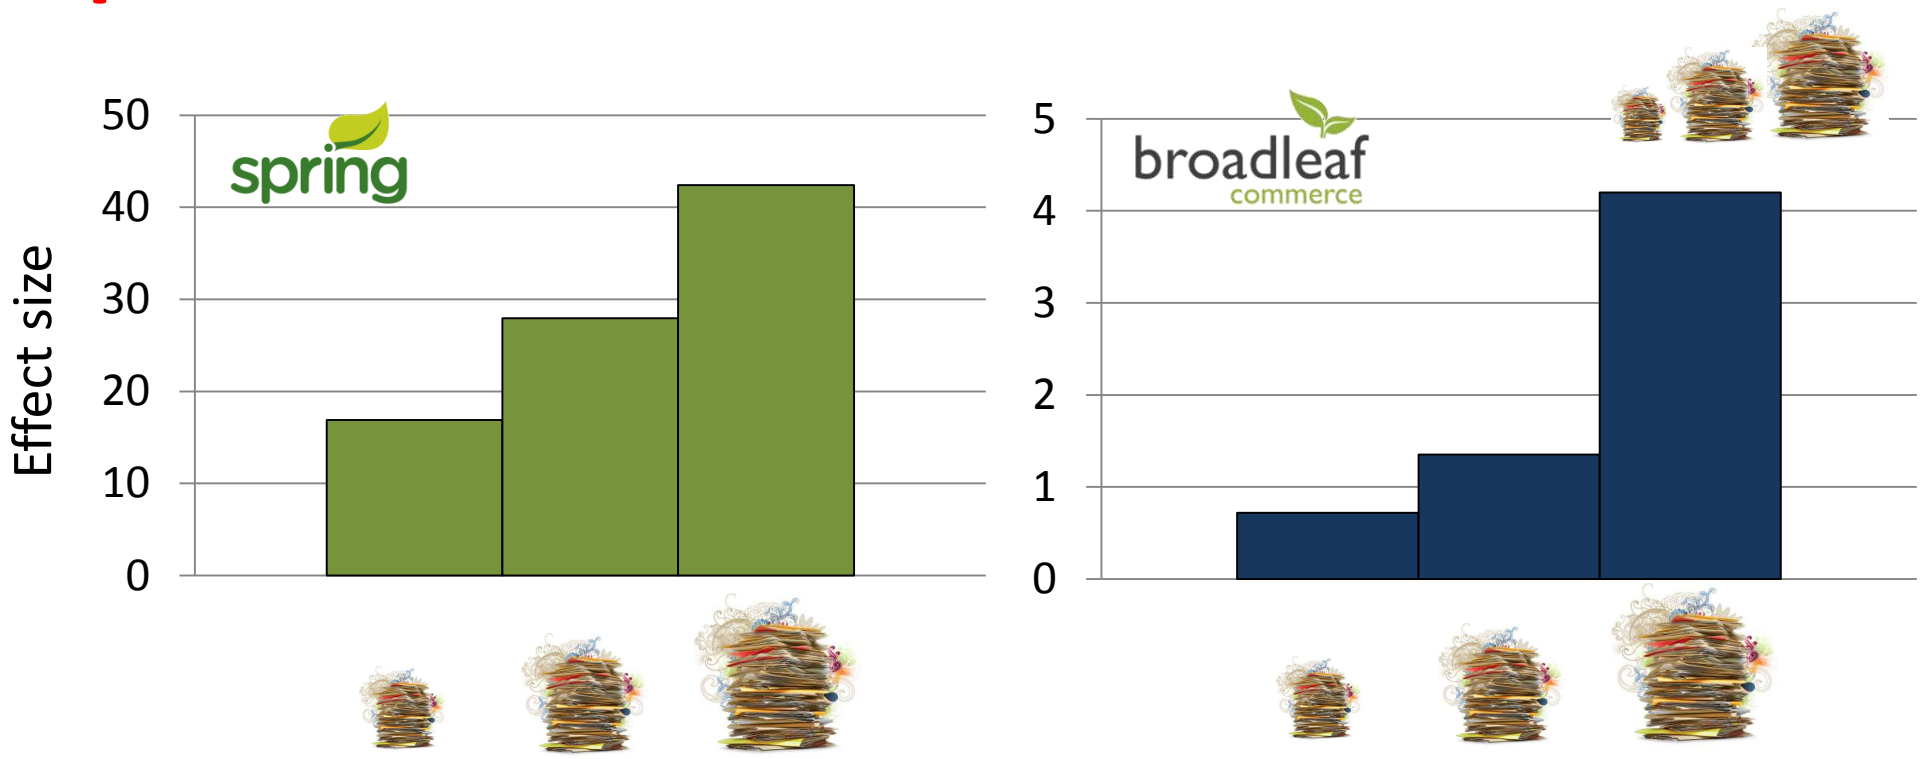

Effect sizes and the ranks of the anti-patterns are consistent in different data scales

Research questions

Performance impact

Ranks of the anti-patterns at different scales

Removing anti-pattern improves response by ~35%

Ranks of the anti-patterns are consistent in different data scales

Object-Relational Mapping eliminates the gap between objects and SQL

Problem of using raw SQLs

- Lots of **boilerplate code**
- Need to **manage object-DB translations** manually

Accessing the database using ORM

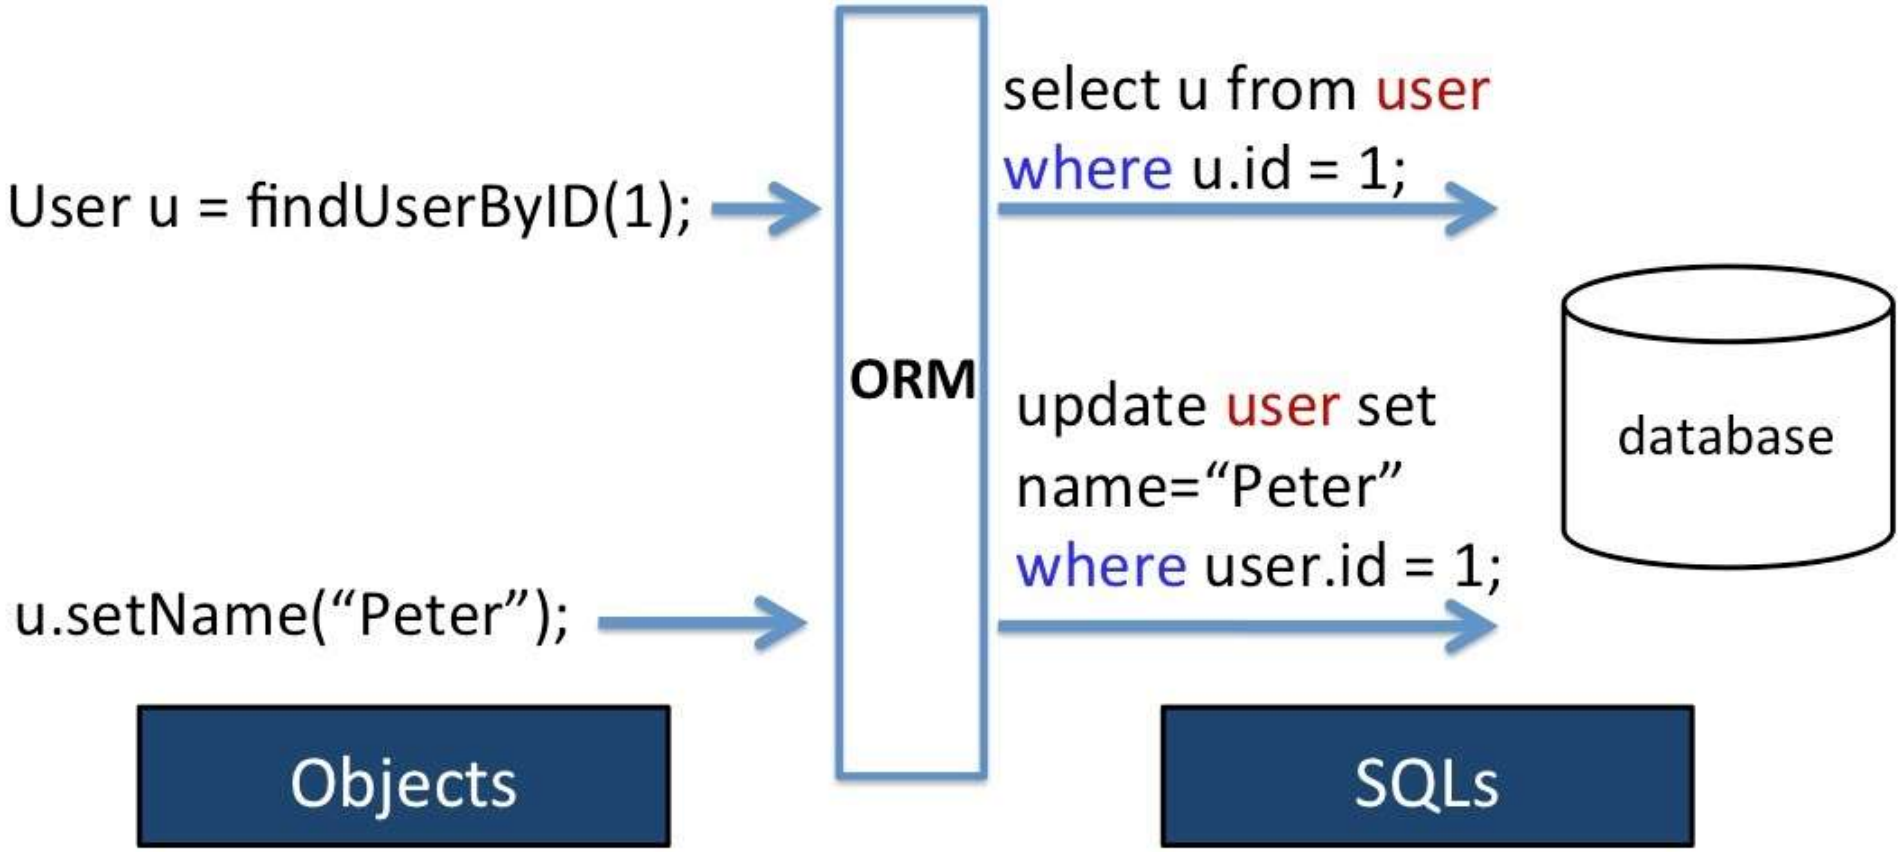

Performance anti-pattern detection framework

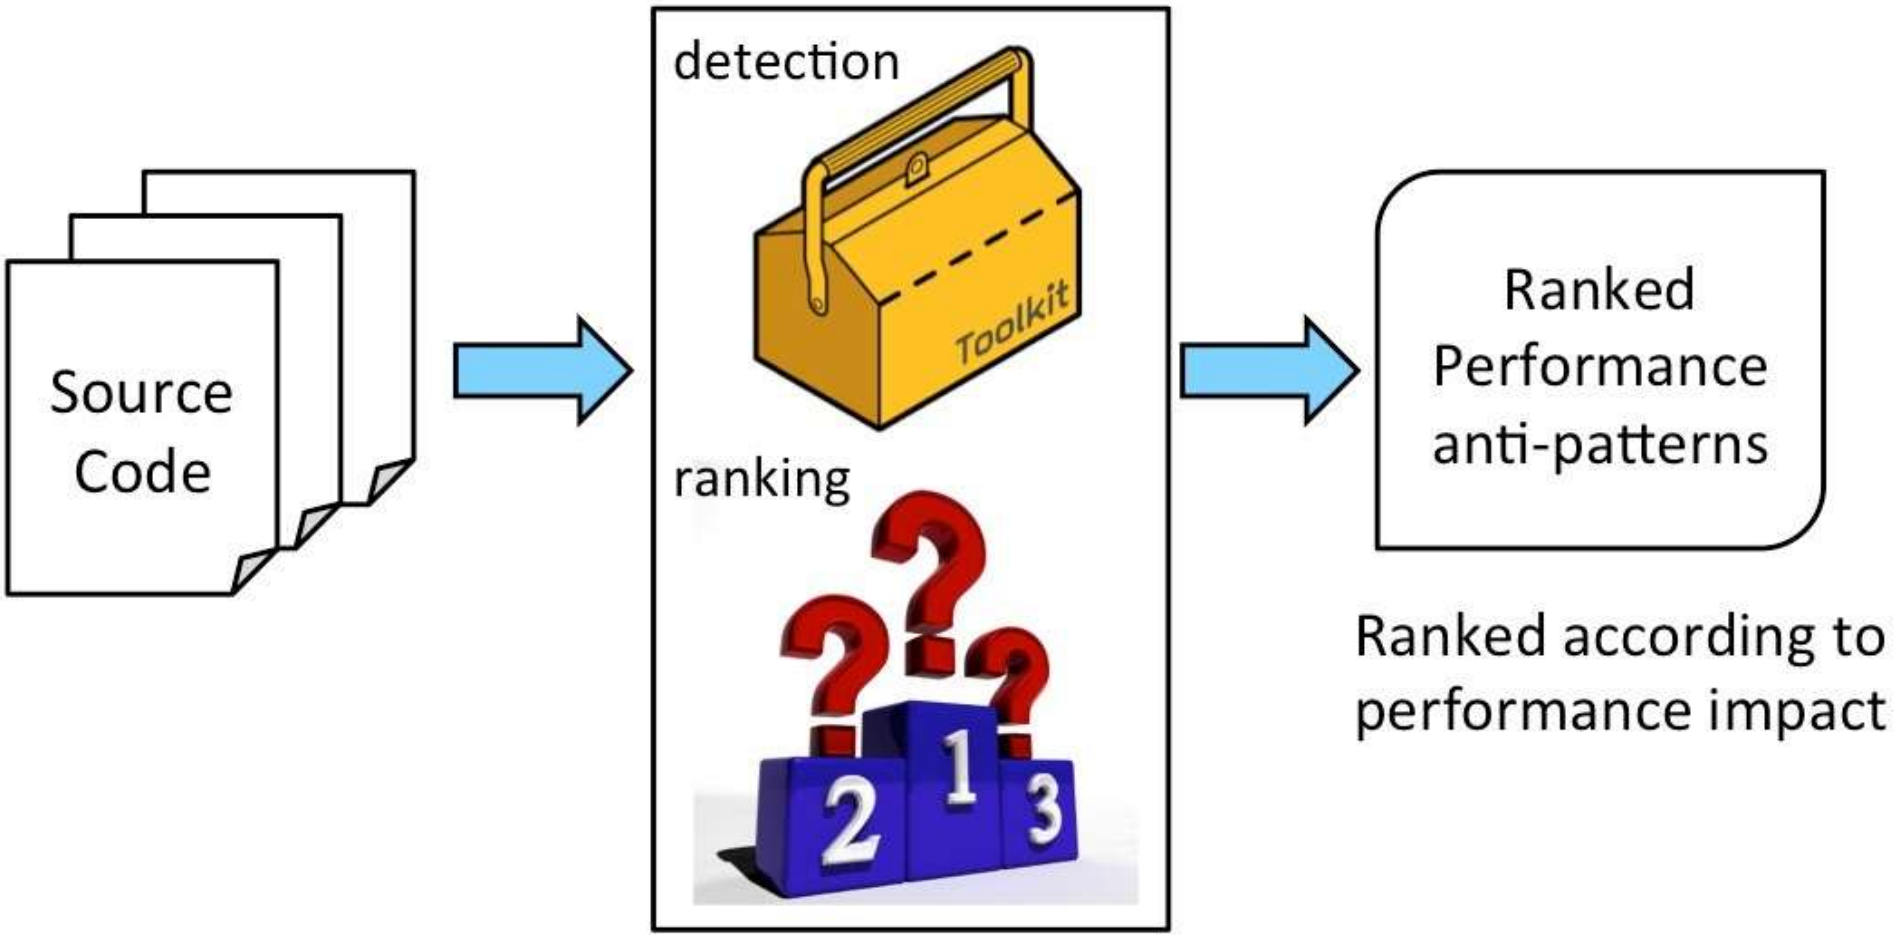

Performance anti-pattern detection and ranking framework

Research questions

Performance impact

Ranks of the anti-patterns at different scales

Removing anti-pattern improves response by ~35%

Ranks of the anti-patterns are consistent in different data scales

Object-Relational Mapping eliminates the gap between objects and SQL

Much less code and shorter development time

Problem of using raw SQLs

- Lots of **boilerplate code**
- N… **manage object-DB tr…**

Accessing the database using ORM

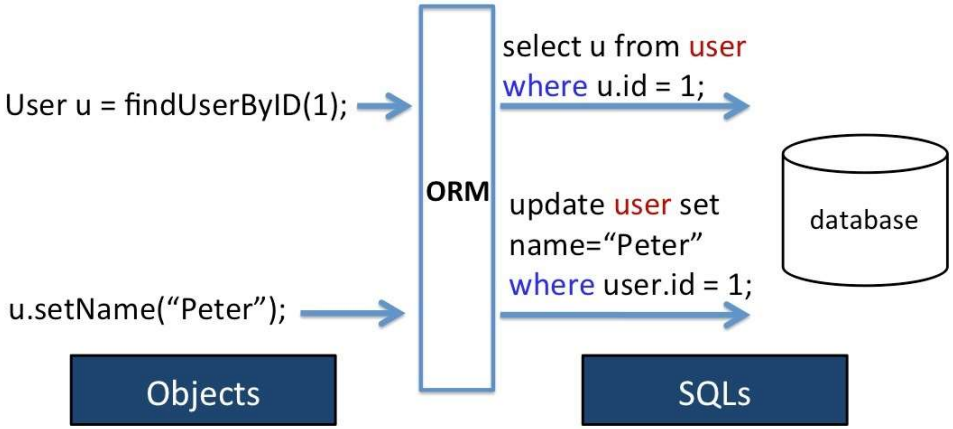

tsehsun@cs.queensu.ca

Perform… detection framework …estions

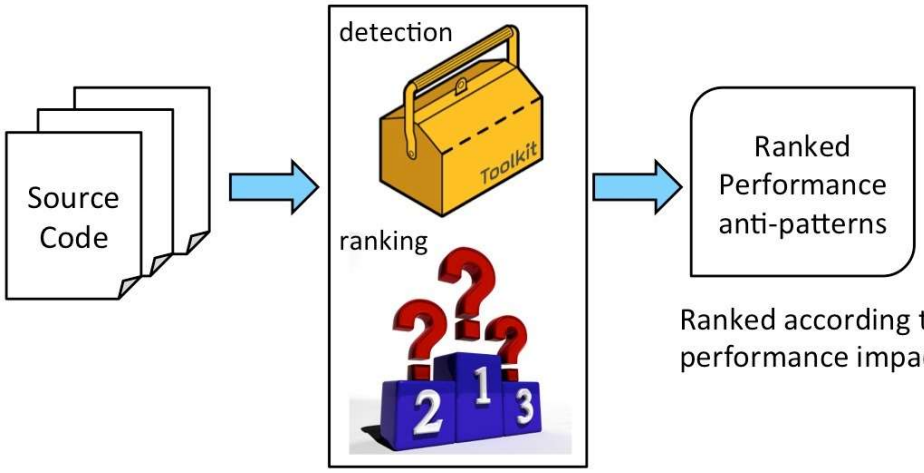

Ranked according to performance impact

Performance anti-pattern detection and ranking framework

Performance impact

Ranks of the anti-patterns at different scales

Removing anti-pattern improves response by ~35%

Ranks of the anti-patterns are consistent in different data scales