

DCatch: Automatically Detecting Distributed Concurrency Bugs in Cloud Systems

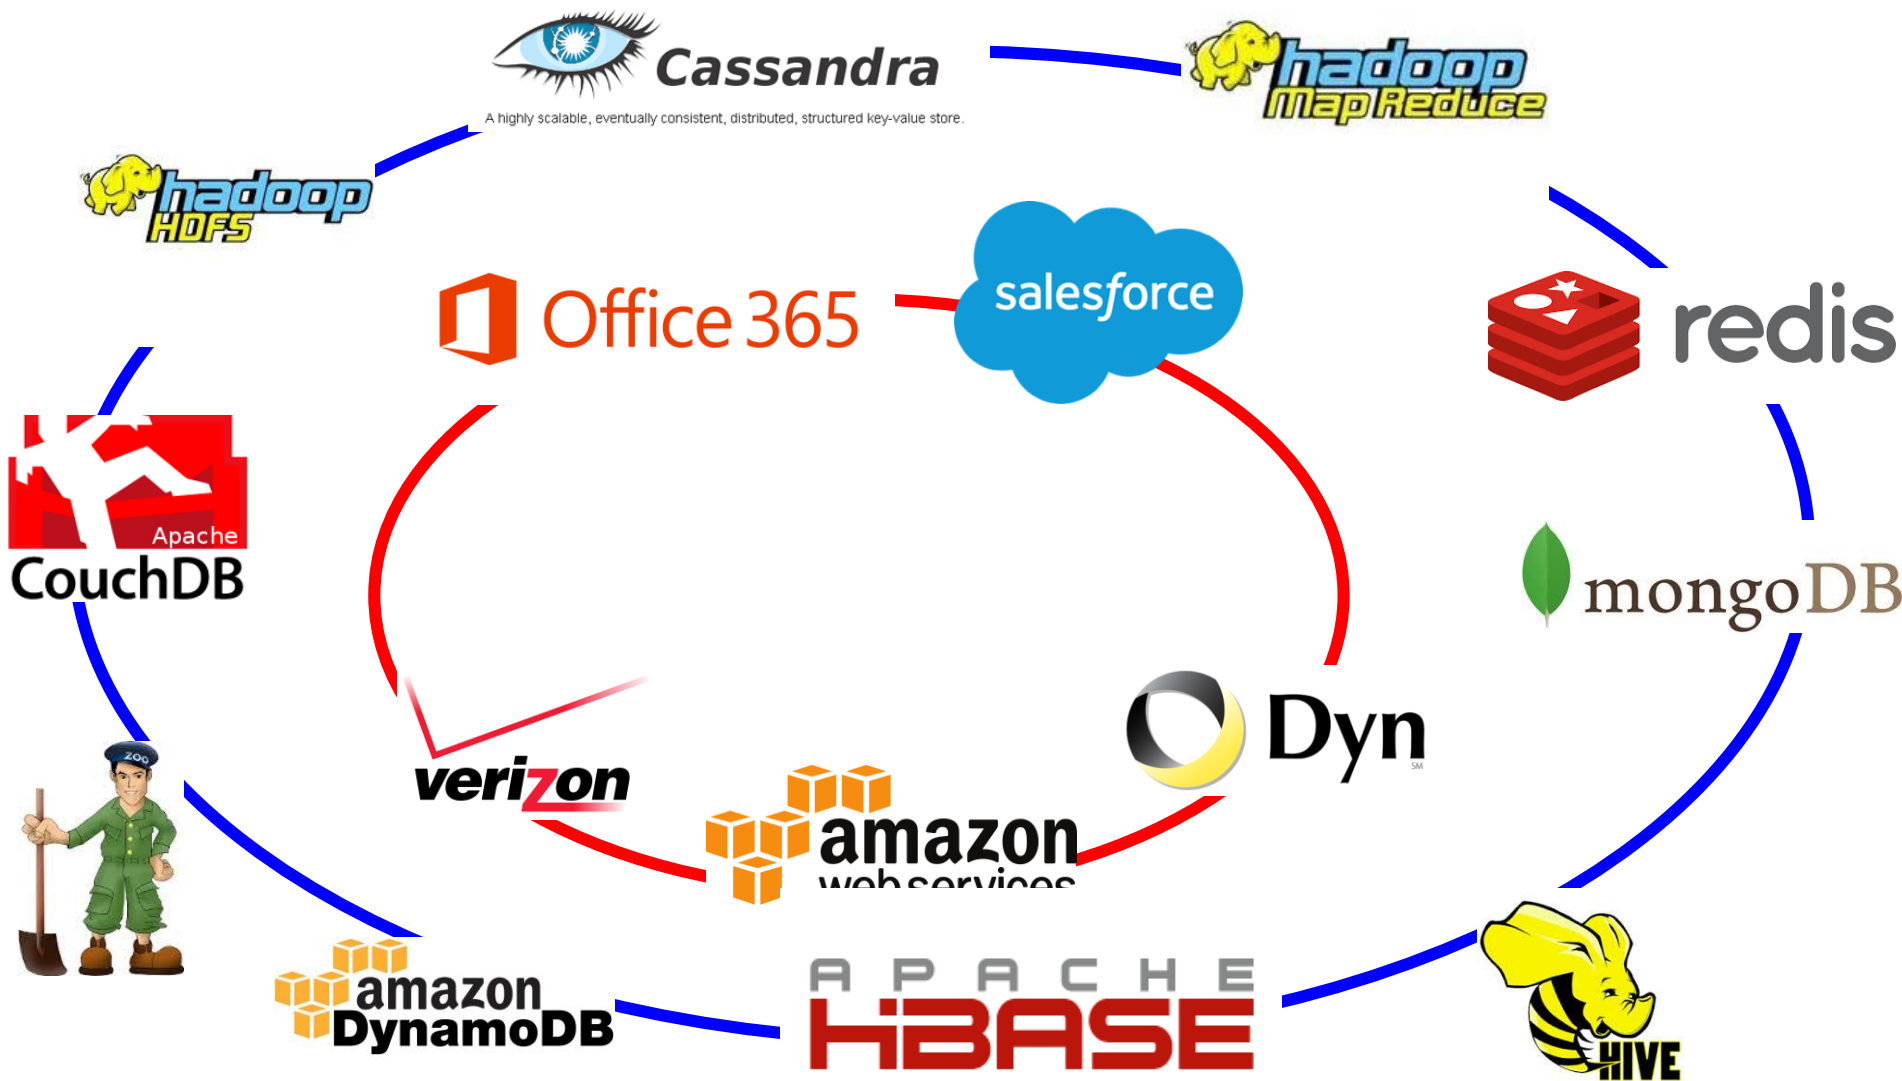
Haopeng Liu, Guangpu Li, Jeffrey Lukman, Jiaxin Li,
Shan Lu, Haryadi Gunawi, and Chen Tian*



THE UNIVERSITY OF
CHICAGO



Cloud systems



Cloud systems



Cassandra

A highly scalable, eventually consistent, distributed, structured key-value store.



The 10 Biggest Cloud Outages of 2016

redis

www.crn.com/slide-shows/cloud/.../the-10-bigges

Dec 29, 2016 - The biggest cloud outages of 2016 incl

large, are increasingly vulnerable from downtime.

mongoDB



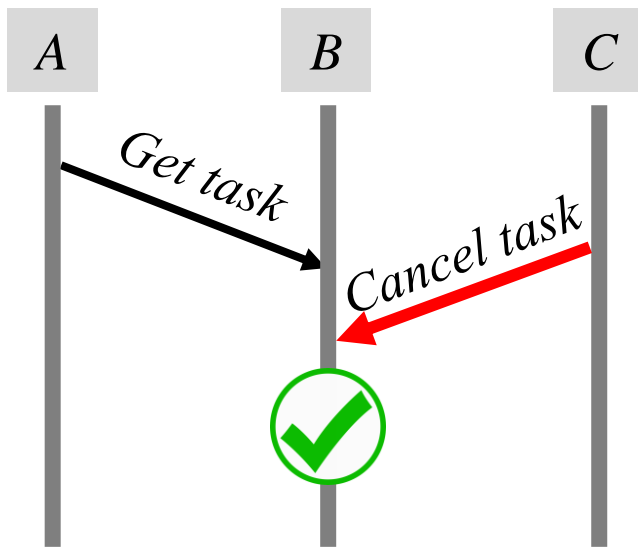
Distributed concurrency bugs (DCbugs)

Distributed concurrency bugs (DCbugs)

- Unexpected timing among distributed operations

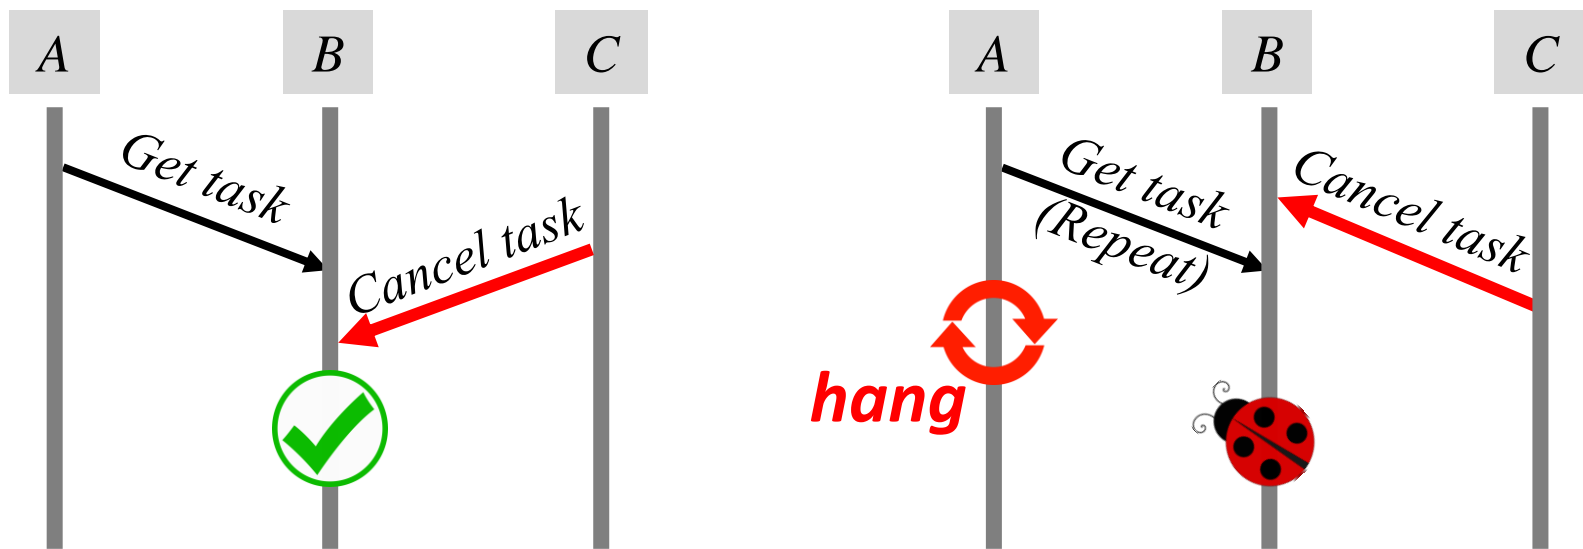
Distributed concurrency bugs (DCbugs)

- Unexpected timing among distributed operations
- Example



Distributed concurrency bugs (DCbugs)

- Unexpected timing among distributed operations
- Example



DCbugs need to be tackled

- Common in distributed systems [1, 2, 3]
 - 26% failures caused by non-deterministic [1]
 - 6% software bugs in clouds system [2]

[1] Yuan. Simple Testing Can Prevent Most Critical Failures. In OSDI'14

[2] Gunawi. What Bugs Live in the Cloud?. In SoCC'14

[3] Leesatapornwongsa. TaxDC. In ASPLOS'16

DCbugs need to be tackled

- Common in distributed systems [1, 2, 3]
- Difficult to avoid, expose and diagnose

[1] Yuan. Simple Testing Can Prevent Most Critical Failures. In OSDI'14

[2] Gunawi. What Bugs Live in the Cloud?. In SoCC'14

[3] Leesatapornwongsa. TaxDC. In ASPLOS'16

DCbugs need to be tackled

- Common in distributed systems [1, 2, 3]
- Difficult to avoid, expose and diagnose



[Hadoop Map/Reduce](#) / [MAPREDUCE-3274](#)

“That is one monster of a race!”

- [1] Yuan. Simple Testing Can Prevent Most Critical Failures. In OSDI'14
- [2] Gunawi. What Bugs Live in the Cloud?. In SoCC'14
- [3] Leesatapornwongsa. TaxDC. In ASPLOS'16

DCbugs need to be tackled

- Common in distributed systems [1, 2, 3]
- Difficult to avoid, expose and diagnose



Hadoop Map/Reduce / [MAPREDUCE-3274](#)



HBase / [HBASE-4397](#)

“There isn’t a week going by without new bugs about races.”

- [1] Yuan. Simple Testing Can Prevent Most Critical Failures. In OSDI’14
[2] Gunawi. What Bugs Live in the Cloud?. In SoCC’14
[3] Leesatapornwongsa. TaxDC. In ASPLOS’16

DCbugs need to be tackled

- Common in distributed systems [1, 2, 3]
- Difficult to avoid, expose and diagnose



Hadoop Map/Reduce / MAPREDUCE-3274



HBase / HBASE-4397



HBase / HBASE-6147

“We have already fix many cases,
however it seems exist many other [racing] cases.”

[1] Yuan. Simple Testing Can Prevent Most Critical Failures. In OSDI'14

[2] Gunawi. What Bugs Live in the Cloud?. In SoCC'14

[3] Leesatapornwongsa. TaxDC. In ASPLOS'16

DCbugs need to be tackled

- Common in distributed systems [1, 2, 3]
- Difficult to avoid, expose and diagnose



Hadoop Map/Reduce / MAPREDUCE-3274



HBase / HBASE-4397



HBase / HBASE-6147



Hadoop Map/Reduce / MAPREDUCE-4819

“This has become quite messy, **sigh.**”

[1] Yuan. Simple Testing Can Prevent Most Critical Failures. In OSDI'14

[2] Gunawi. What Bugs Live in the Cloud?. In SoCC'14

[3] Leesatapornwongsa. TaxDC. In ASPLOS'16

DCbugs need to be tackled

- Common in distributed systems [1, 2, 3]
- Difficult to avoid, expose and diagnose



Hadoop Map/Reduce / MAPREDUCE-3274



HBase / HBASE-4397



HBase / HBASE-6147



Hadoop Map/Reduce / MAPREDUCE-4819



Hadoop Map/Reduce / MAPREDUCE-4099

“Great catch, Sid! Apologies for missing the race condition.”

[1] Yuan. Simple Testing Can Prevent Most Critical Failures. In OSDI'14

[2] Gunawi. What Bugs Live in the Cloud?. In SoCC'14

[3] Leesatapornwongsa. TaxDC. In ASPLOS'16

DCbugs need to be tackled

- Common in distributed systems [1, 2, 3]
- Difficult to avoid, expose and diagnose



Hadoop Map/Reduce / MAPREDUCE-3274



HBase / HBASE-4397



HBase / HBASE-6147



Hadoop Map/Reduce / MAPREDUCE-4819



Hadoop Map/Reduce / MAPREDUCE-4099



Hadoop Map/Reduce / MAPREDUCE-3634

“We [prefer] debug crashes instead of hanging jobs.”

[1] Yuan. Simple Testing

[2] Gunawi. What Bugs Live in the Cloud?. In SoCC'14

[3] Leesatapornwongsa. TaxDC. In ASPLOS'16

DCbugs need to be tackled

- Common in distributed systems [1, 2, 3]
- Difficult to avoid, expose and diagnose



Hadoop Map/Reduce / MAPREDUCE-3274



HBase / HBASE-4397



HBase / HBASE-6147

Can we detect DCbugs before they manifest?



Hadoop Map/Reduce / MAPREDUCE-3634

“We [prefer] debug crashes instead of hanging jobs.”

[1] Yuan. Simple Testing

[2] Gunawi. What Bugs Live in the Cloud?. In SoCC'14

[3] Leesatapornwongsa. TaxDC. In ASPLOS'16

Previous work

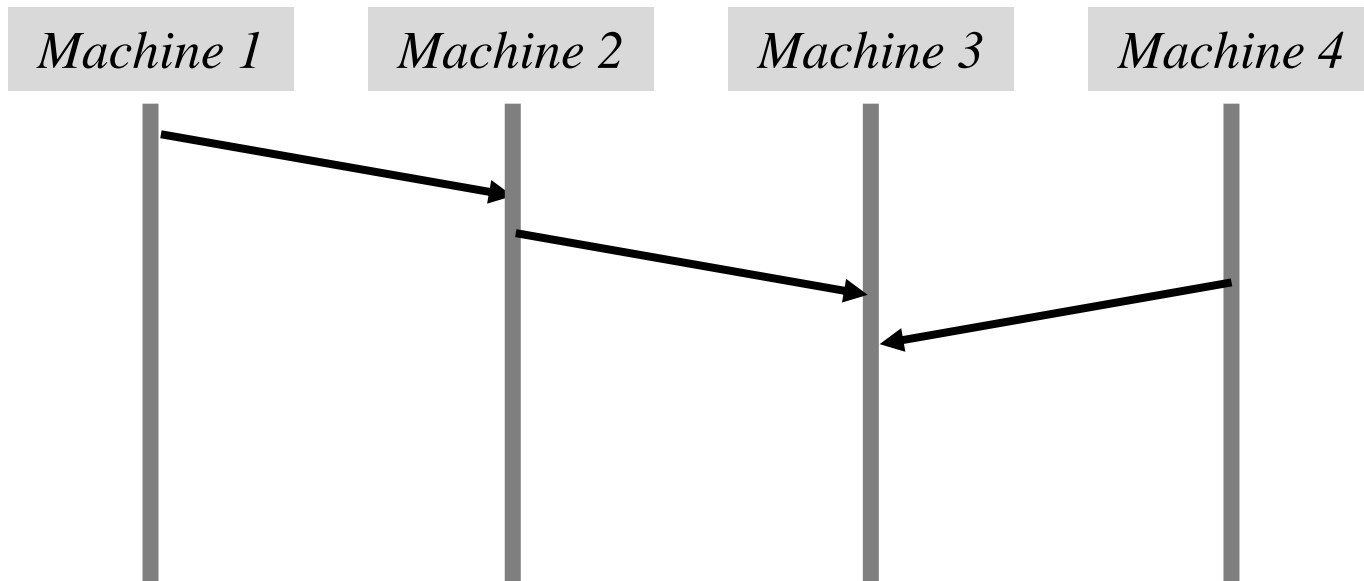
- Model checking
 - Work on abstracted models
 - Face state-space explosion issue

Our idea

- Follow the philosophy of traditional concurrency bug detection

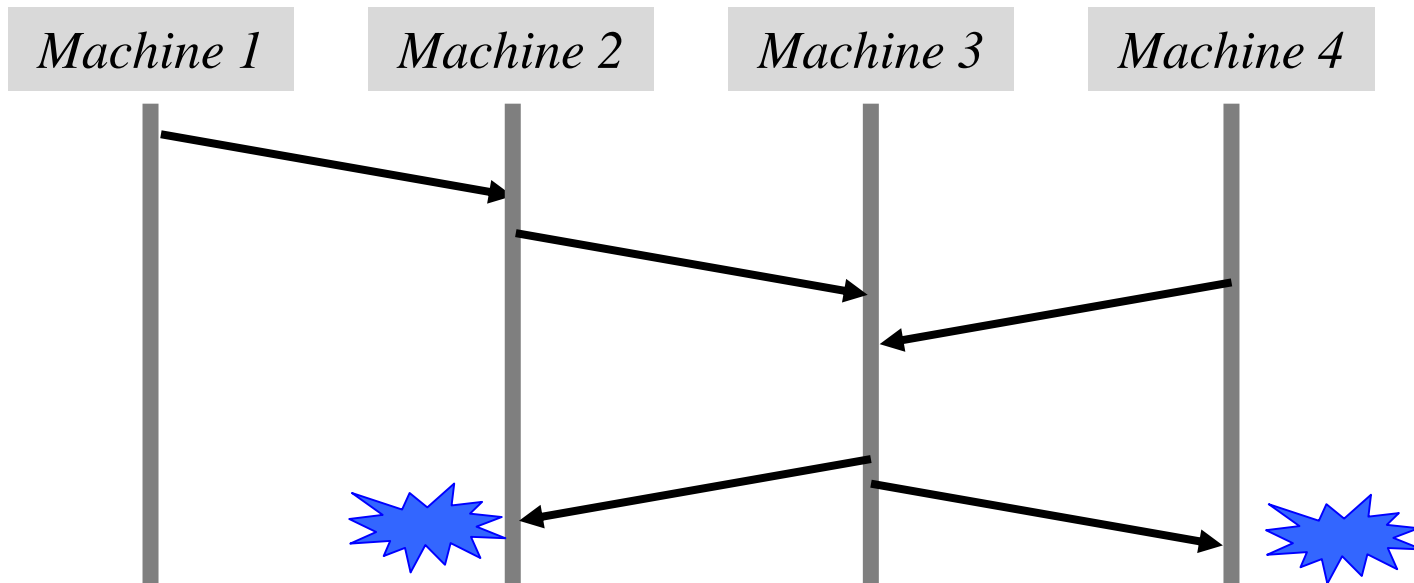
Our idea

- Follow the philosophy of traditional concurrency bug detection



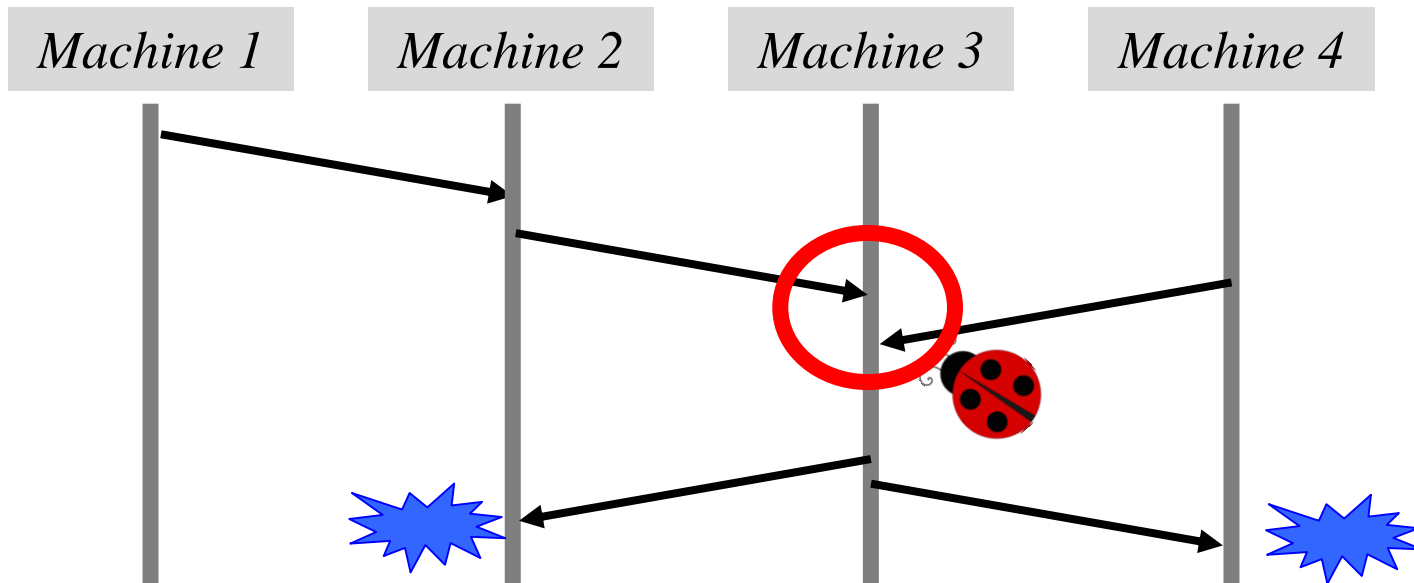
Our idea

- Follow the philosophy of traditional concurrency bug detection

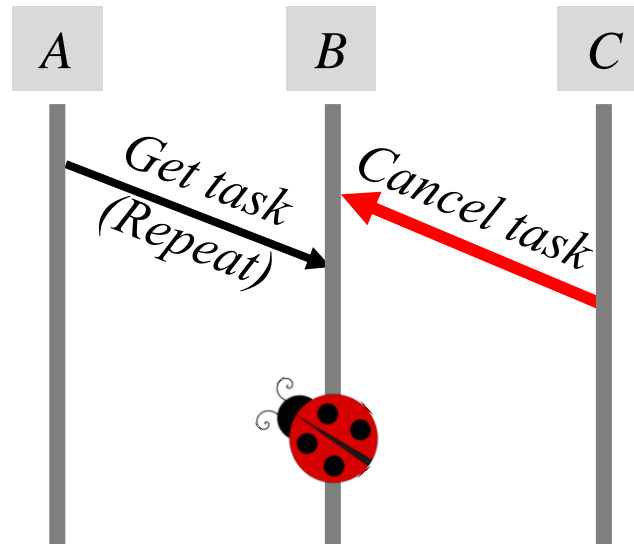


Our idea

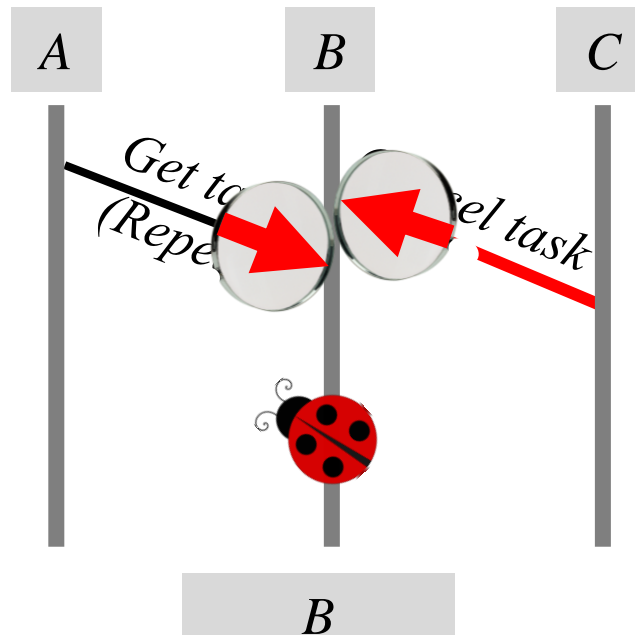
- Follow the philosophy of traditional concurrency bug detection



Example



Example



```
//RPC thread
Task getTask(jID) {
    ...
    return jMap.get(jID);
}
```

```
//UnReg thread
void unReg(jID) {
    jMap.remove(jID);
    ....
}
```

Local concurrency bug detection

ASPLOS

ASPLOS [Architectural Support for Programming](#)

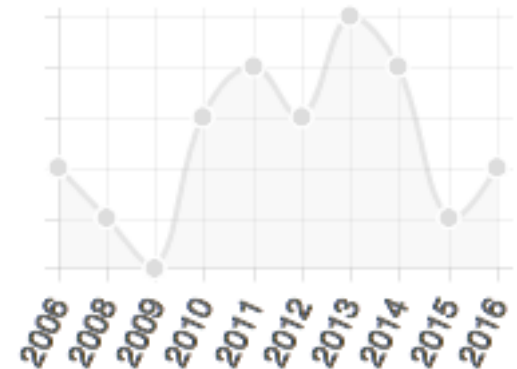
ASPLOS is a multi-disciplinary conference for researchers, compilers, languages, operating systems, network and engineers to present their latest research findings in computer systems innovations of the past two decades. It focuses on multiprocessors, clusters and networks-of-workstations.

This conference occurs at a time when computer processor performance scaling and to new demands are increasingly important as boundaries between hardware capabilities of computing devices become [expand](#)

Search within ASPLOS:

35 results found

Refine by Publication Year



Published Since 2006



Local concurrency bug detection

ASPLOS

ASPLOS [Architectural Support for Programming](#)

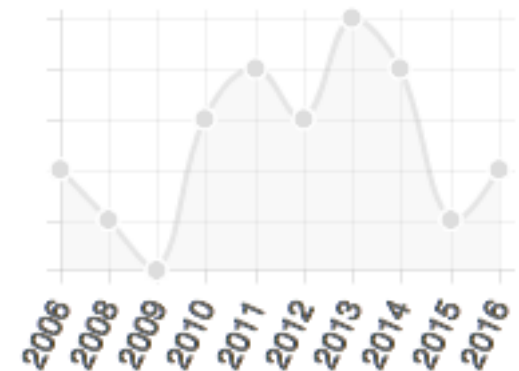
ASPLOS is a multi-disciplinary conference for researchers in compilers, languages, operating systems, networks, and engineers to present their latest research findings on computer systems innovations of the past two decades. It focuses on multiprocessors, clusters and networks-of-workstations.

This conference occurs at a time when computer processor performance scaling and to new demands are increasingly important as boundaries between hardware capabilities of computing devices become expand

Search within ASPLOS: "concurrency bug" "race"

35 results found

Refine by Publication Year



Published Since 2006



Local concurrency bug detection

ASPLOS

ASPLOS Architectural Support for Programming

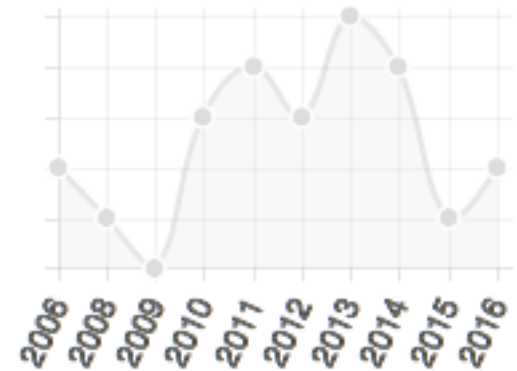
ASPLOS is a multi-disciplinary conference for researchers in compilers, languages, operating systems, networks, and engineers to present their latest research findings and computer systems innovations of the past two decades. The conference focuses on multiprocessors, clusters and networks-of-workstations.

This conference occurs at a time when computer processor performance scaling and to new demands are increasingly important as boundaries between hardware and software capabilities of computing devices become expand

Search within ASPLOS: "concurrency bug" "race"

35 results found

Refine by Publication Year

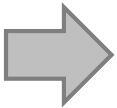
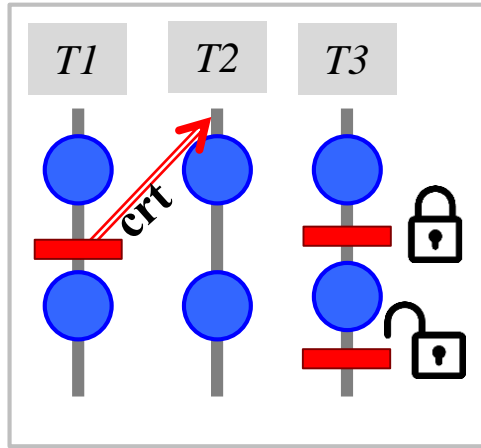


Published Since 2006

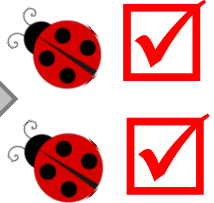
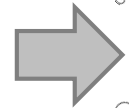


Is the problem solved?

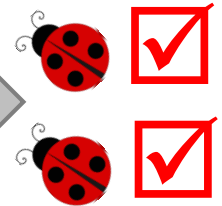
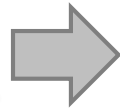
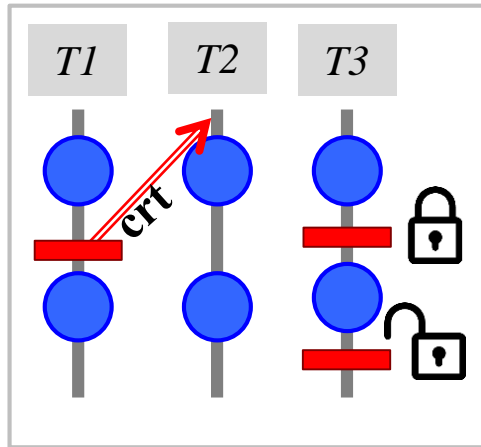
Local concurrency bug detection



Trace



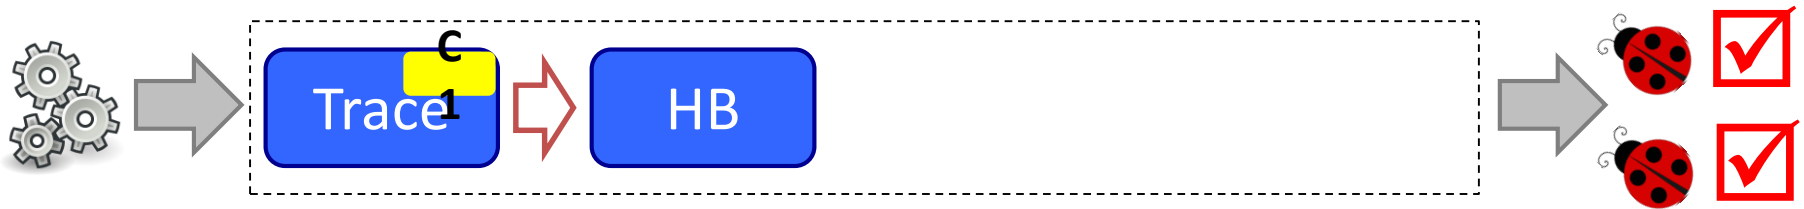
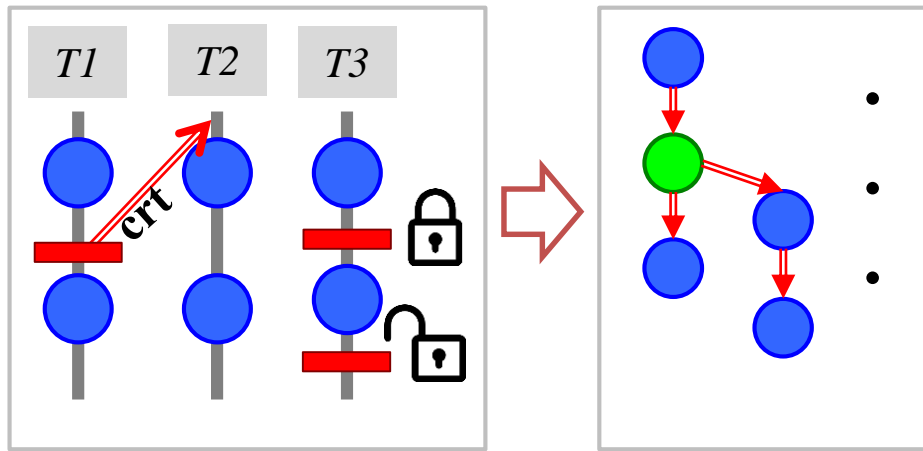
Local concurrency bug detection



C1: How to handle the huge amount of mem accesses?

Challenges

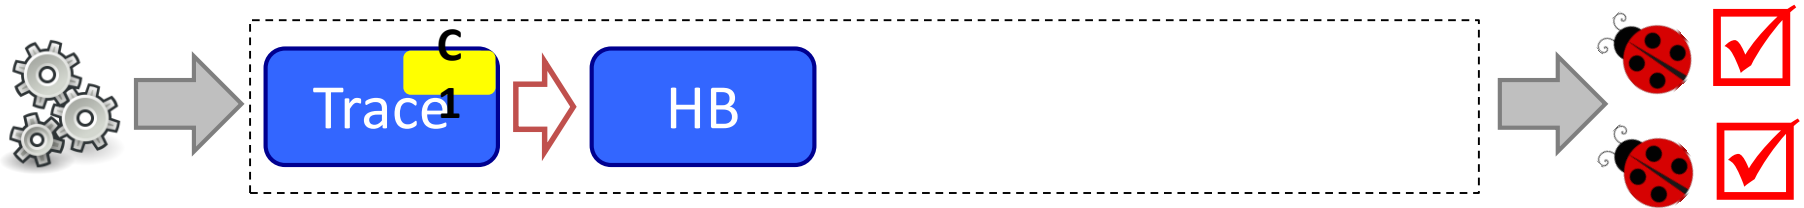
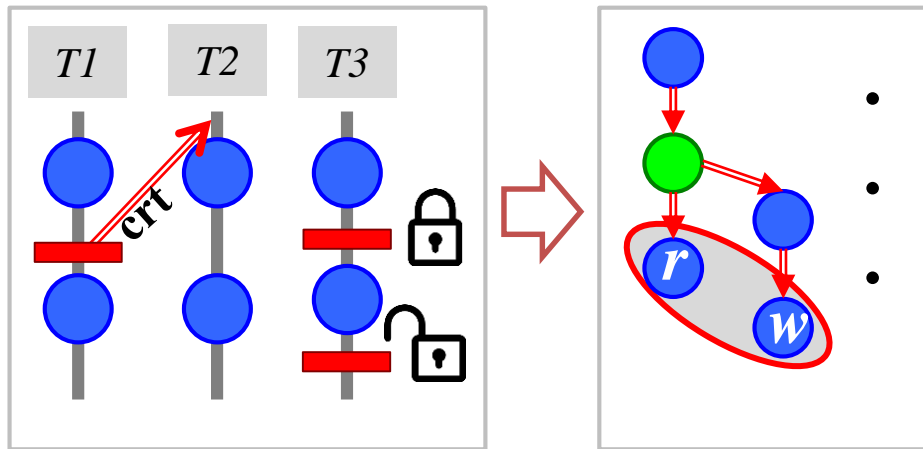
Local concurrency bug detection



C1: How to handle the huge amount of mem accesses?

Challenges

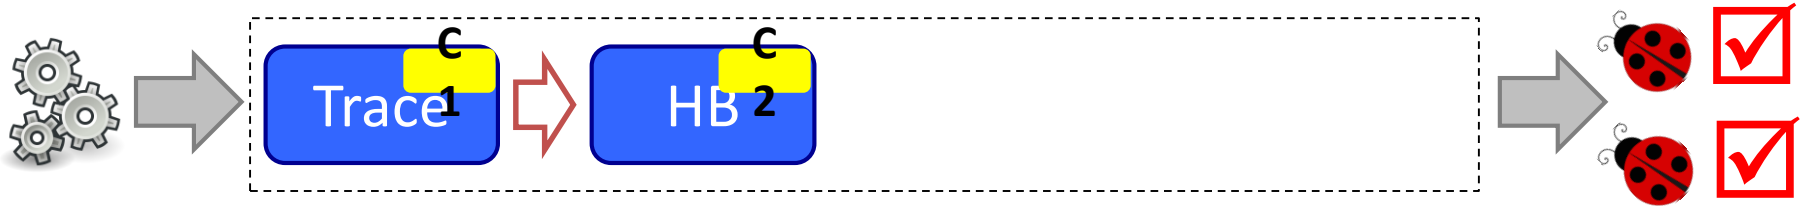
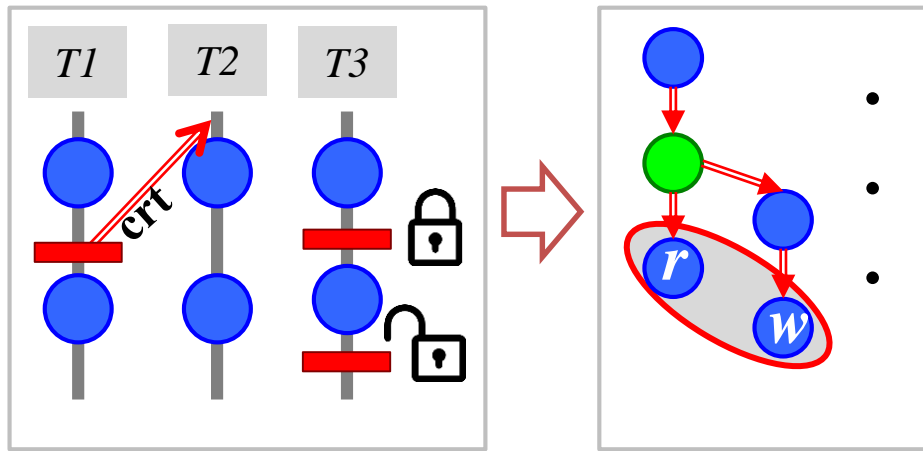
Local concurrency bug detection



C1: How to handle the huge amount of mem accesses?

Challenges

Local concurrency bug detection

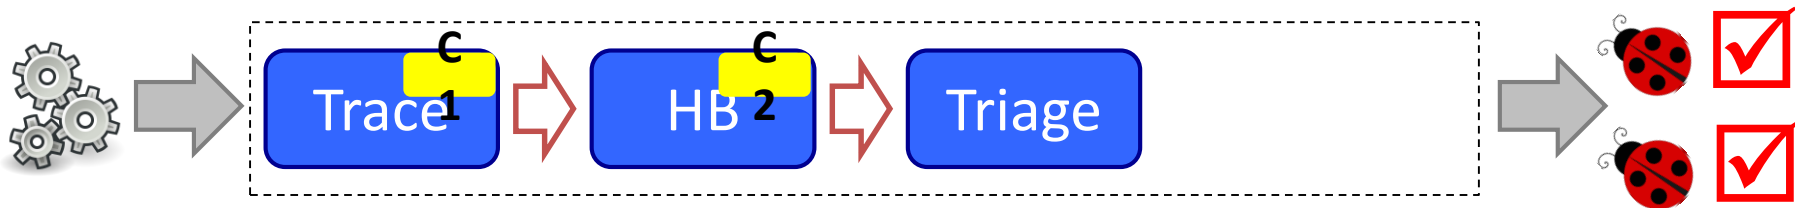
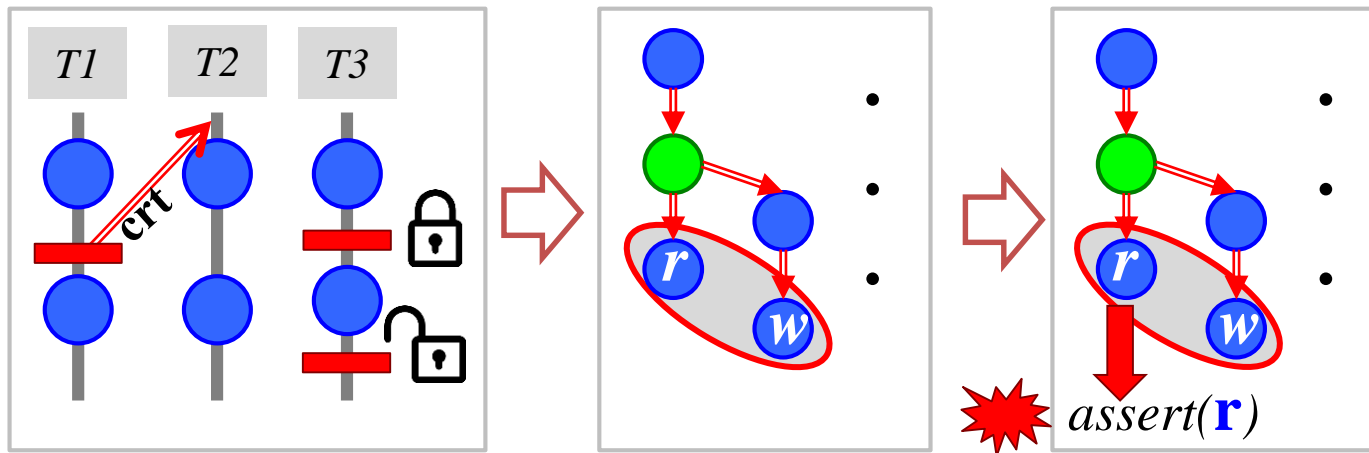


C1: How to handle the huge amount of mem accesses?

C2: What's the happens-before model?

Challenges

Local concurrency bug detection

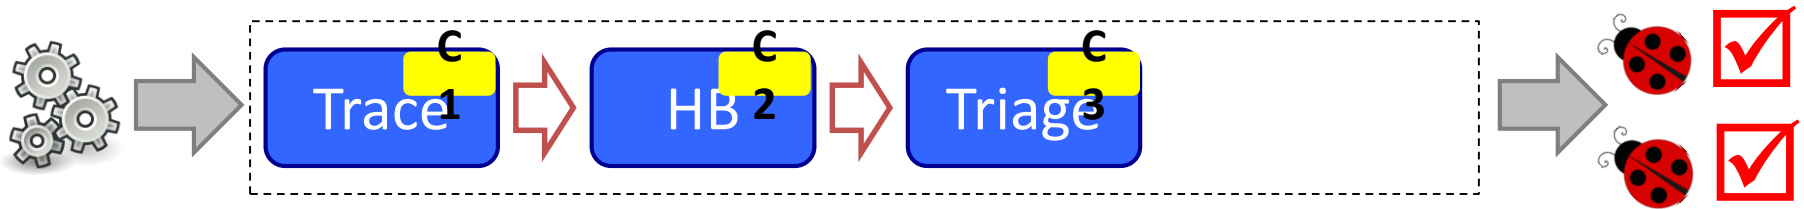
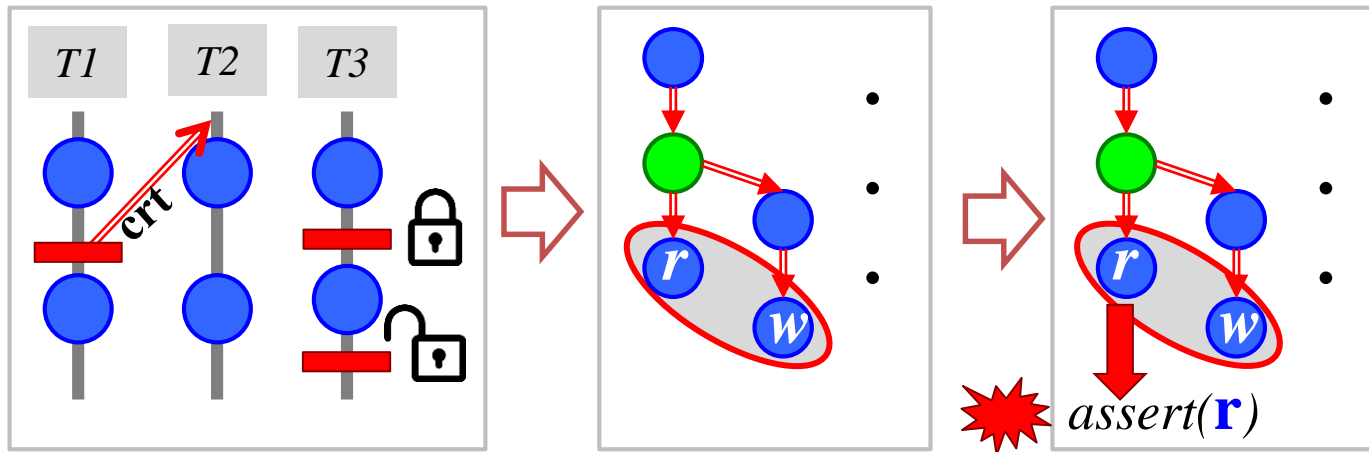


C1: How to handle the huge amount of mem accesses?

C2: What's the happens-before model?

Challenges

Local concurrency bug detection



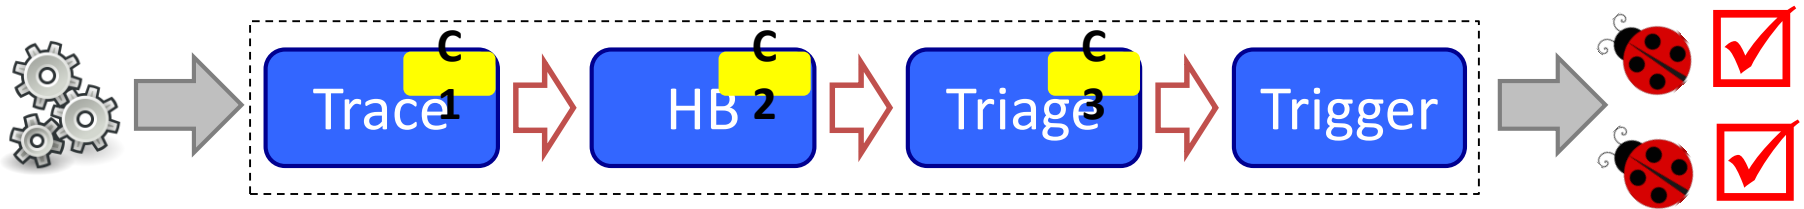
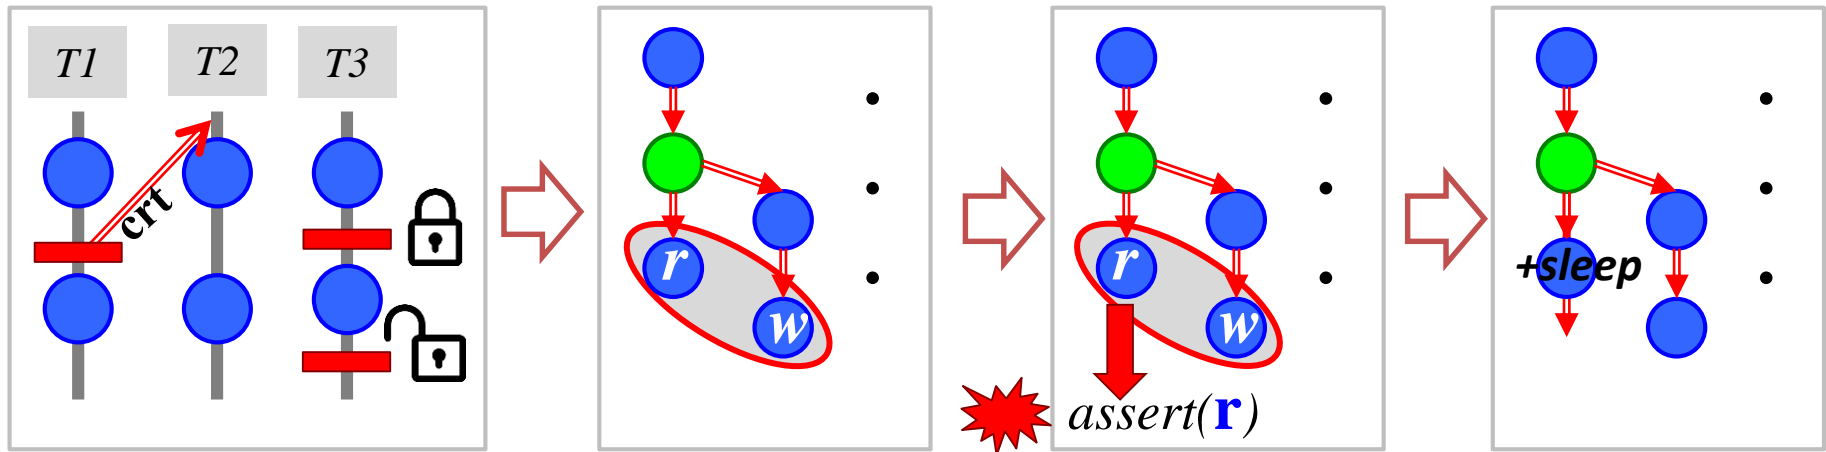
C1: How to handle the huge amount of mem accesses?

C2: What's the happens-before model?

Challenges

C3: How to estimate the distributed impact of a race?

Local concurrency bug detection



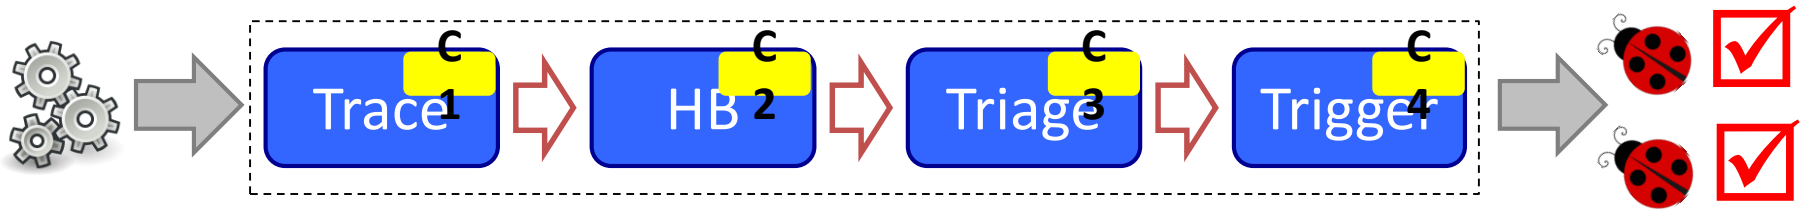
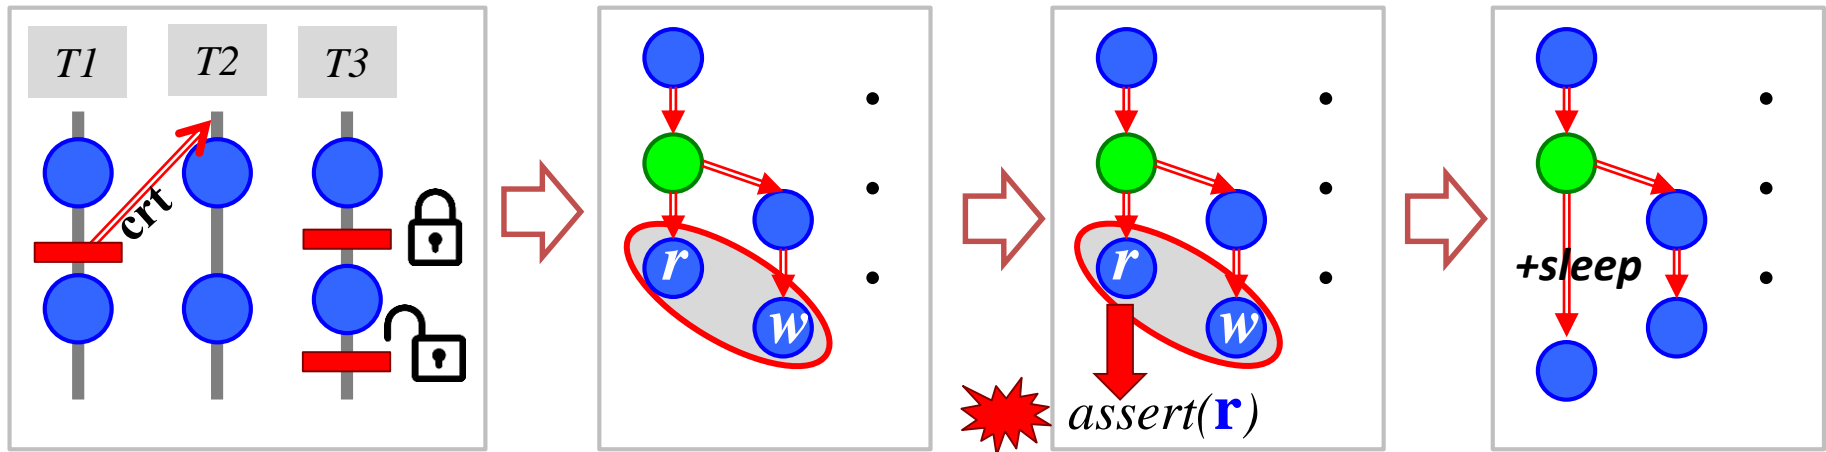
C1: How to handle the huge amount of mem accesses?

C2: What's the happens-before model?

C3: How to estimate the distributed impact of a race?

Challenges

Local concurrency bug detection



C1: How to handle the huge amount of mem accesses?

C2: What's the happens-before model?

Challenges

C3: How to estimate the distributed impact of a race?

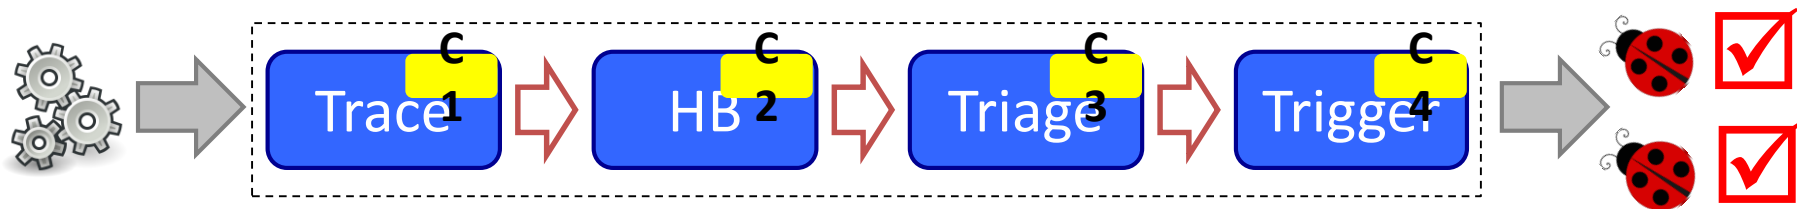
C4: How to trigger with distributed time manipulation?

Contribution

- A **comprehensive HB Model** for distributed systems

Contribution

- A **comprehensive HB Model** for distributed systems
- DCatch tool detects DCbugs from **correct runs**



C1: How to handle the huge amount of mem accesses?





C2: What's the happens-before model?

C3: How to estimate the distributed impact of a race?

C4: How to trigger with distributed time manipulation?

Challenges
Solved by
DCatch

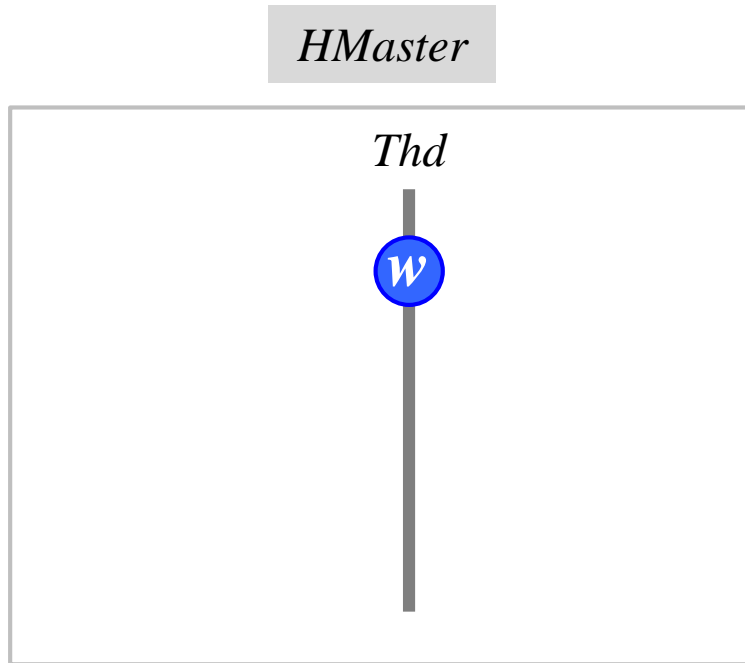
Contribution

- A **comprehensive HB Model** for distributed systems
- DCatch tool detects DCbugs from **correct runs**
- Evaluate on **4** systems    
- Report **32** DCbugs, with **20** of them being truly harmful

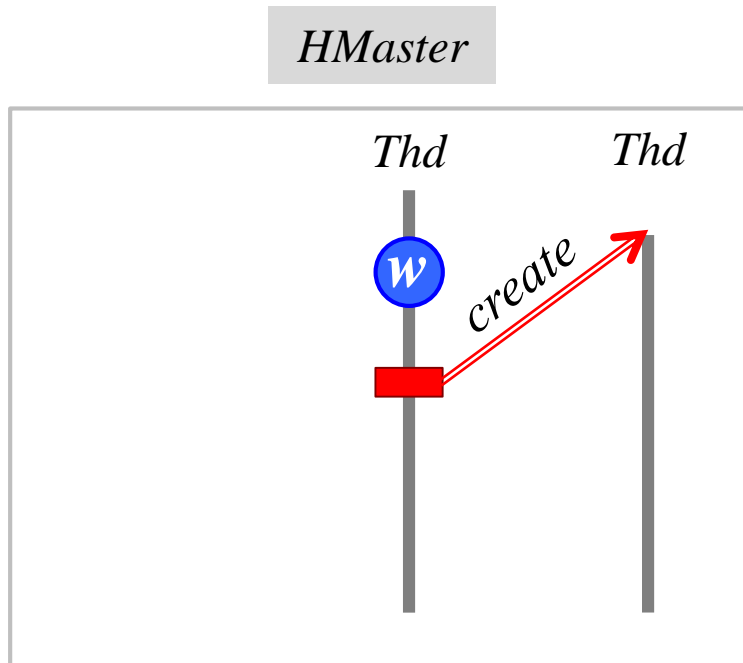
Outline

- Motivation
- **DCatch Happens-before Model**
- DCatch tool
- Evaluation
- Conclusion

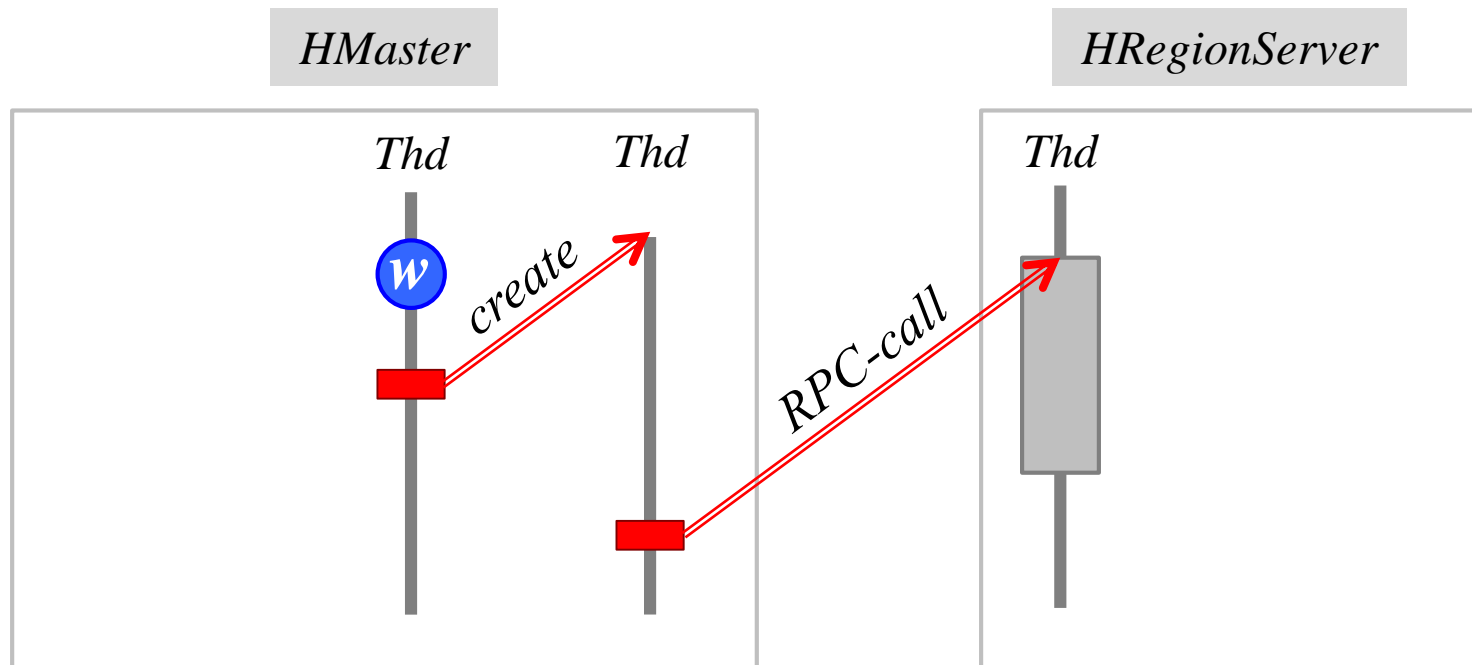
DCatch Happens-before Model



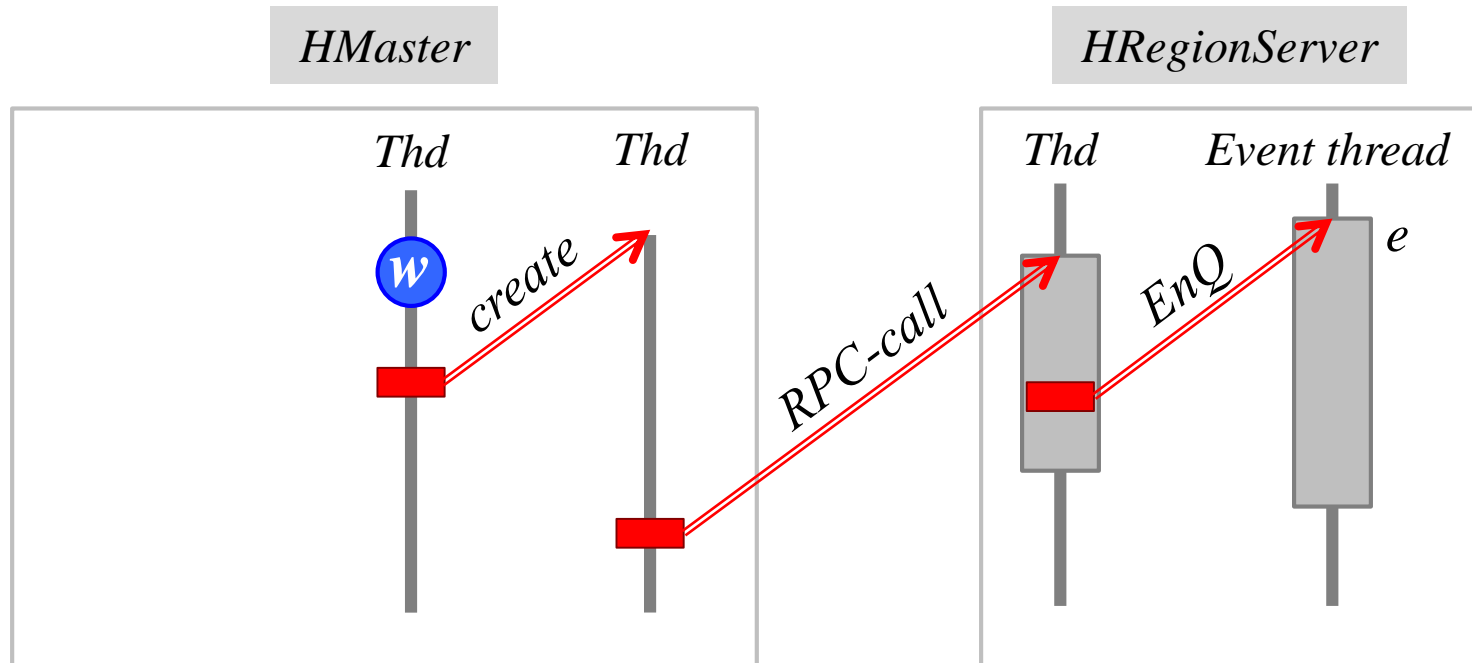
DCatch Happens-before Model



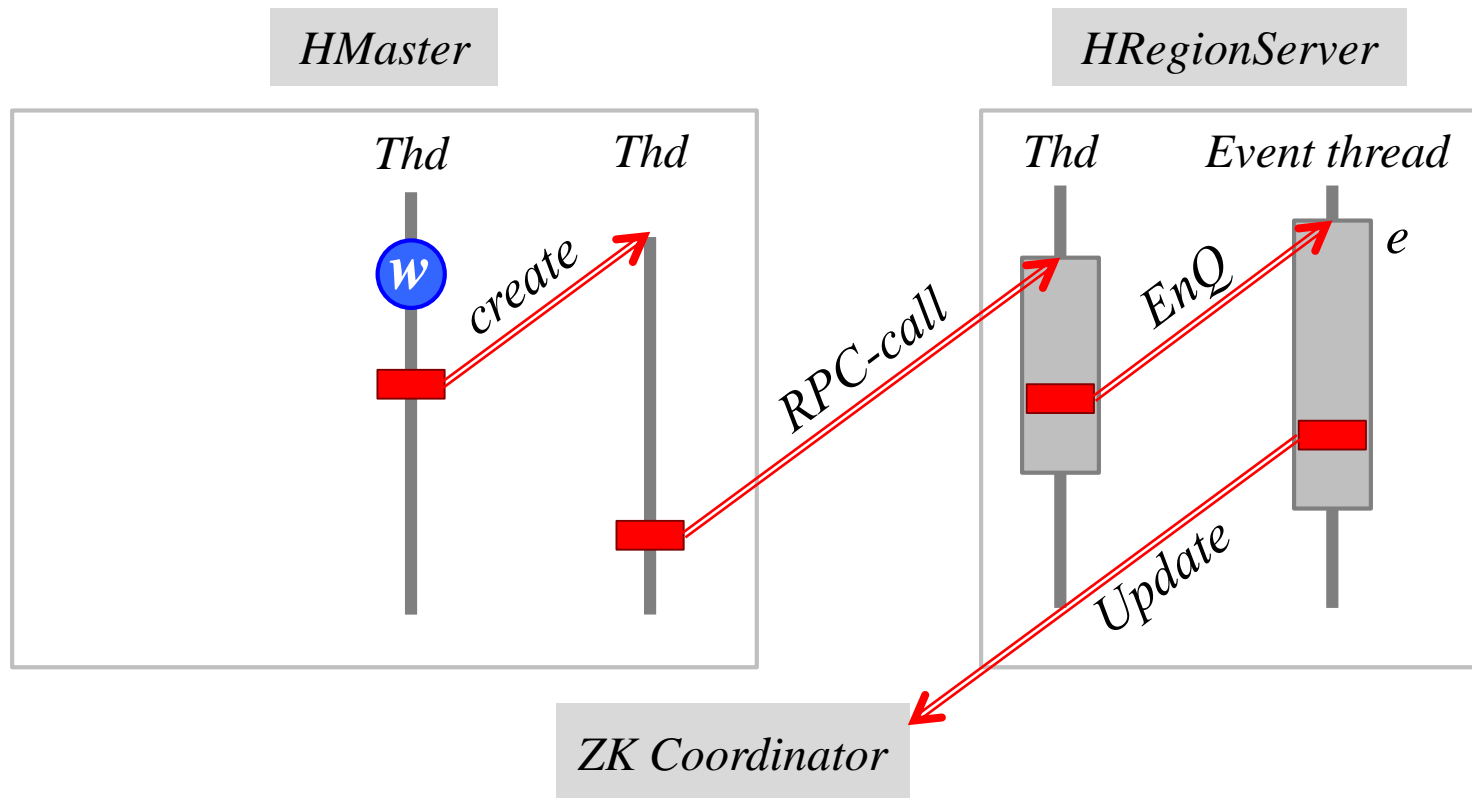
DCatch Happens-before Model



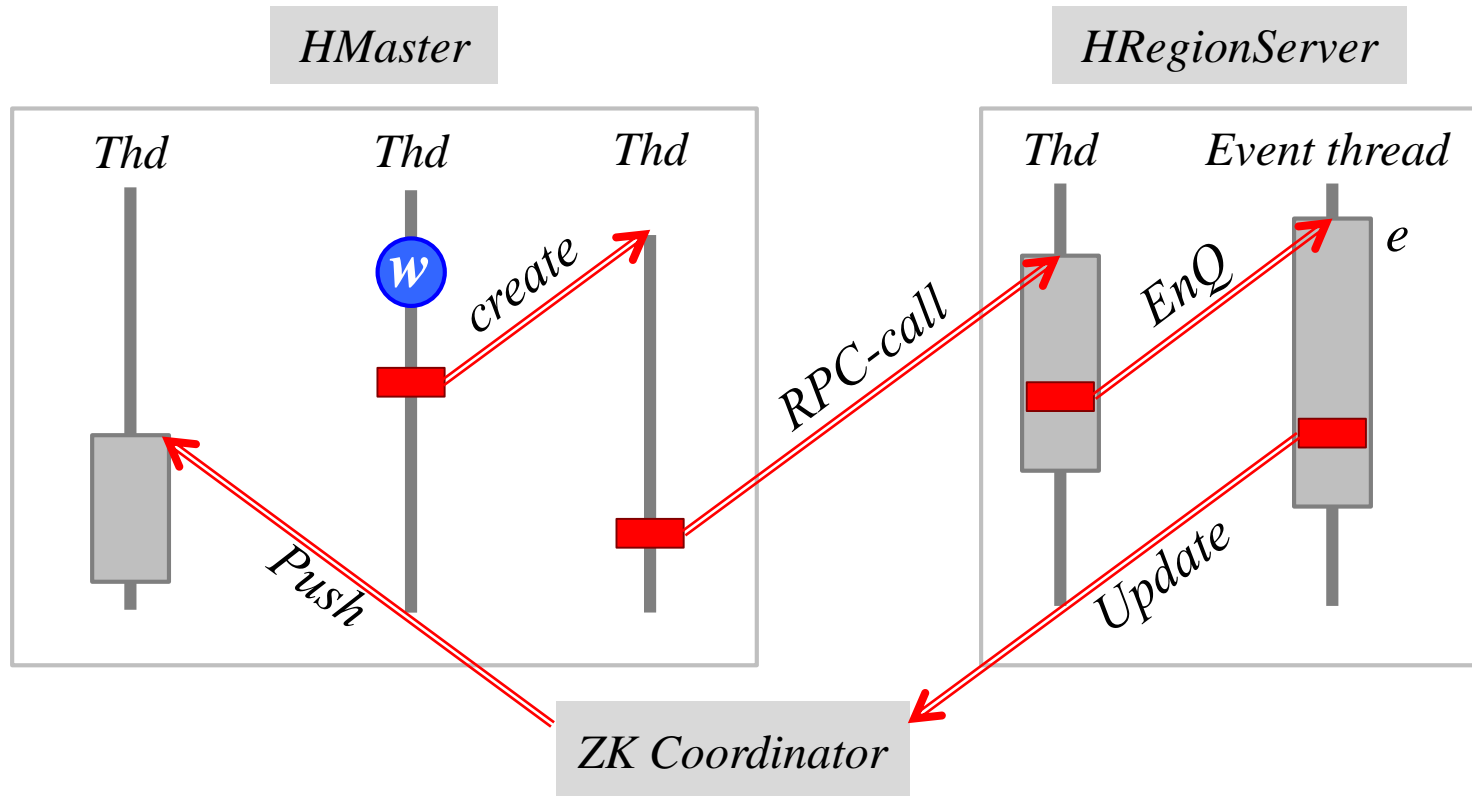
DCatch Happens-before Model



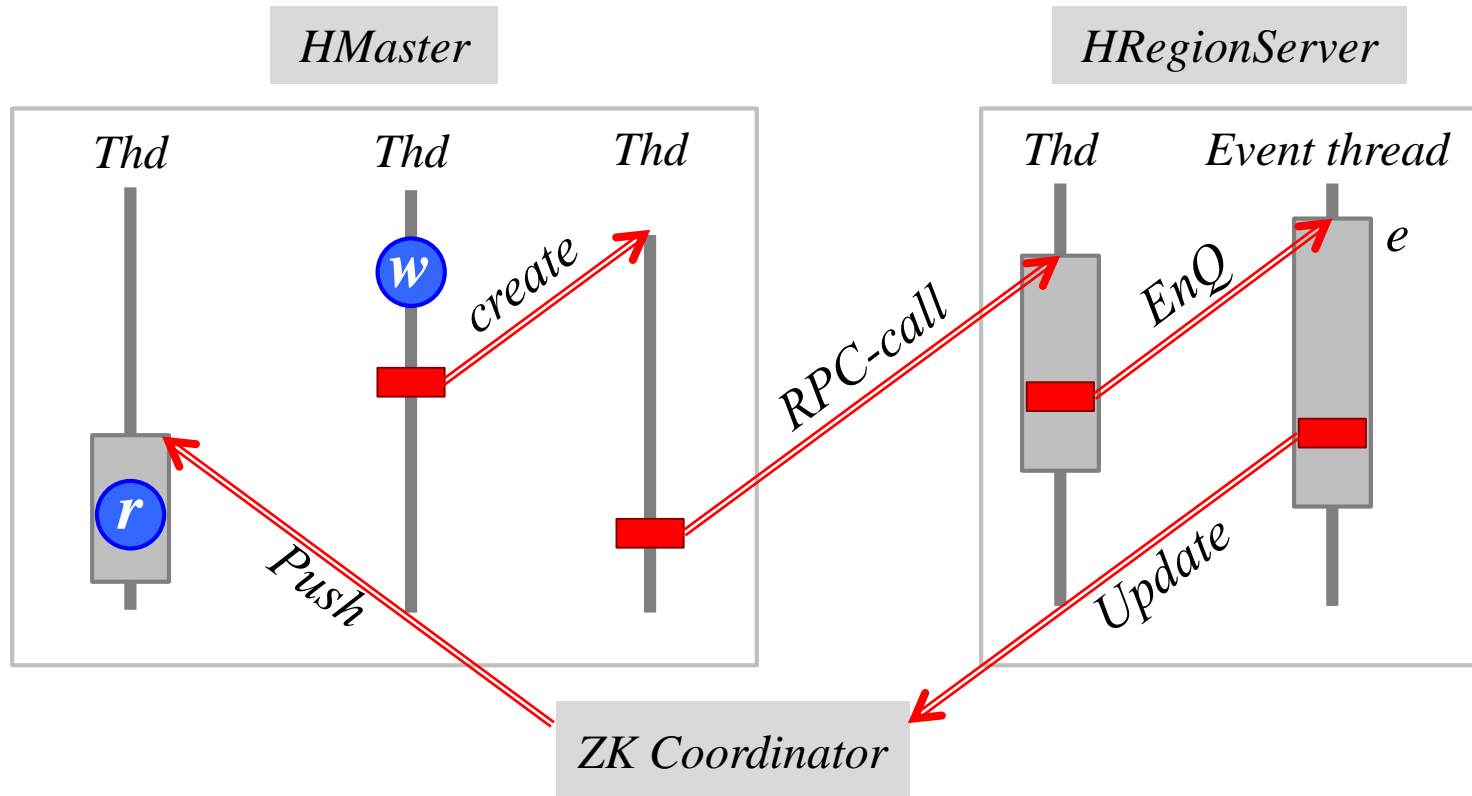
DCatch Happens-before Model



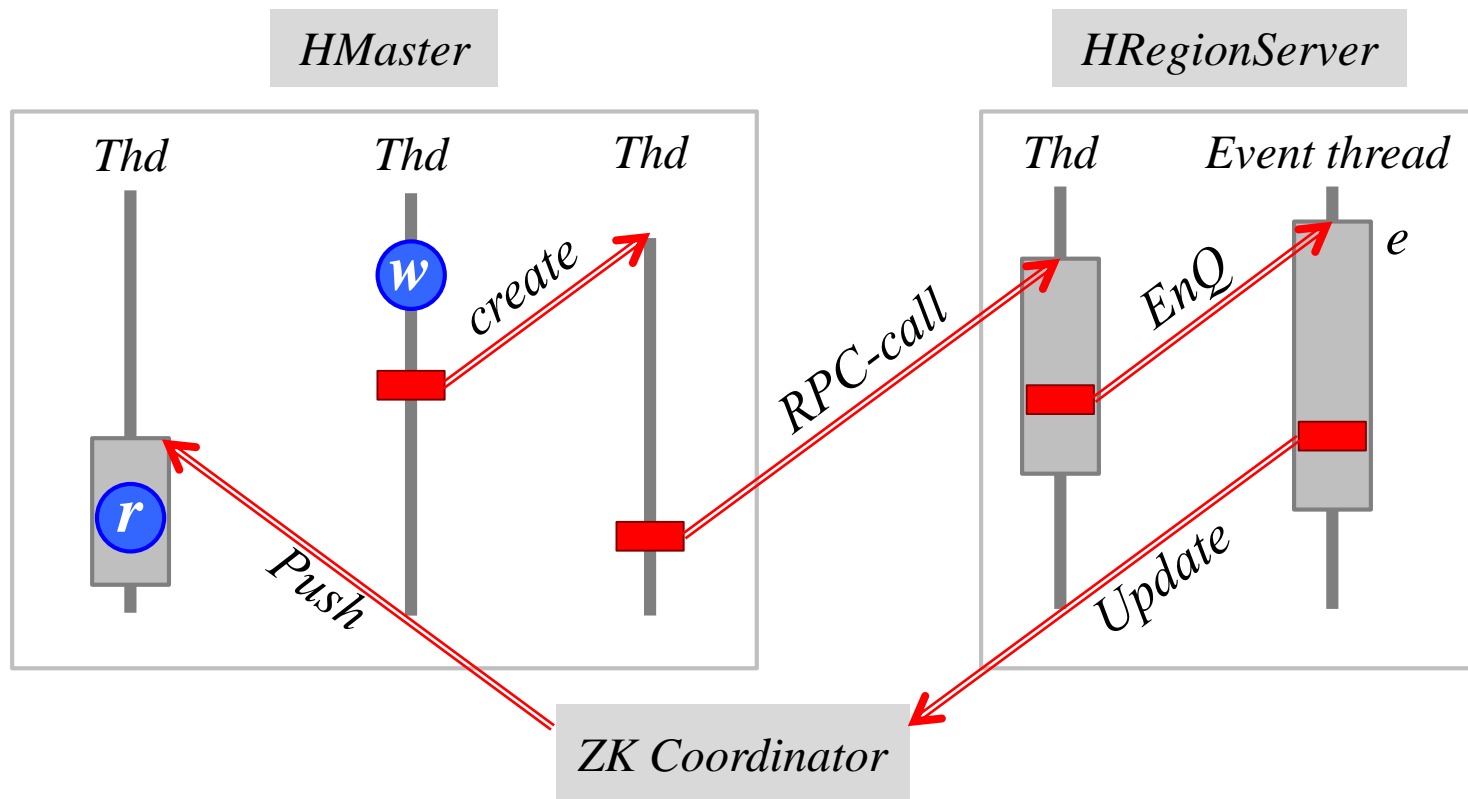
DCatch Happens-before Model



DCatch Happens-before Model

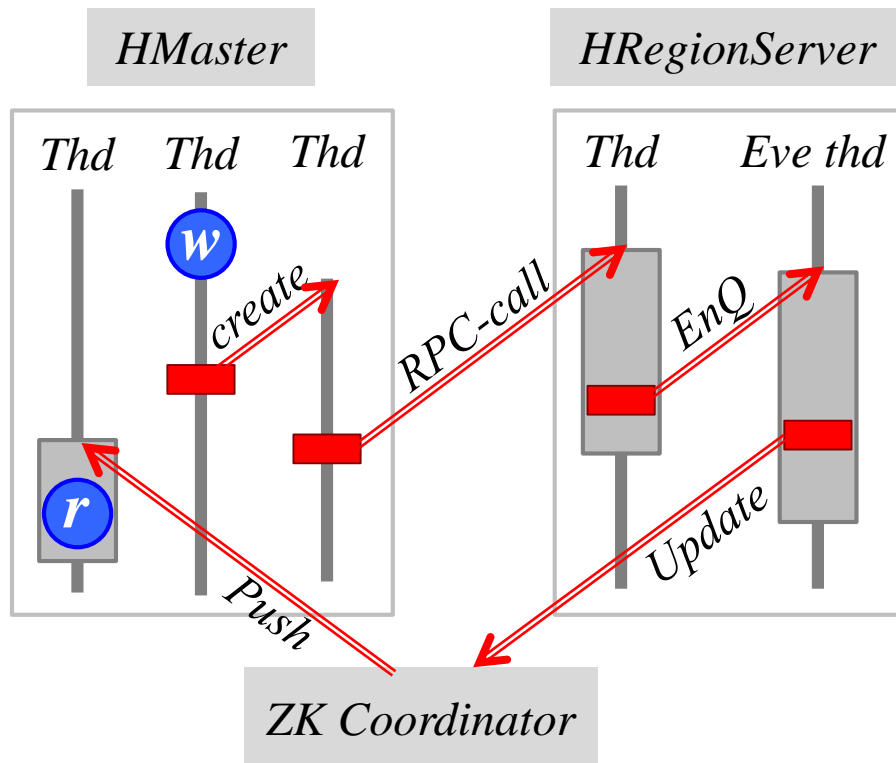


DCatch Happens-before Model

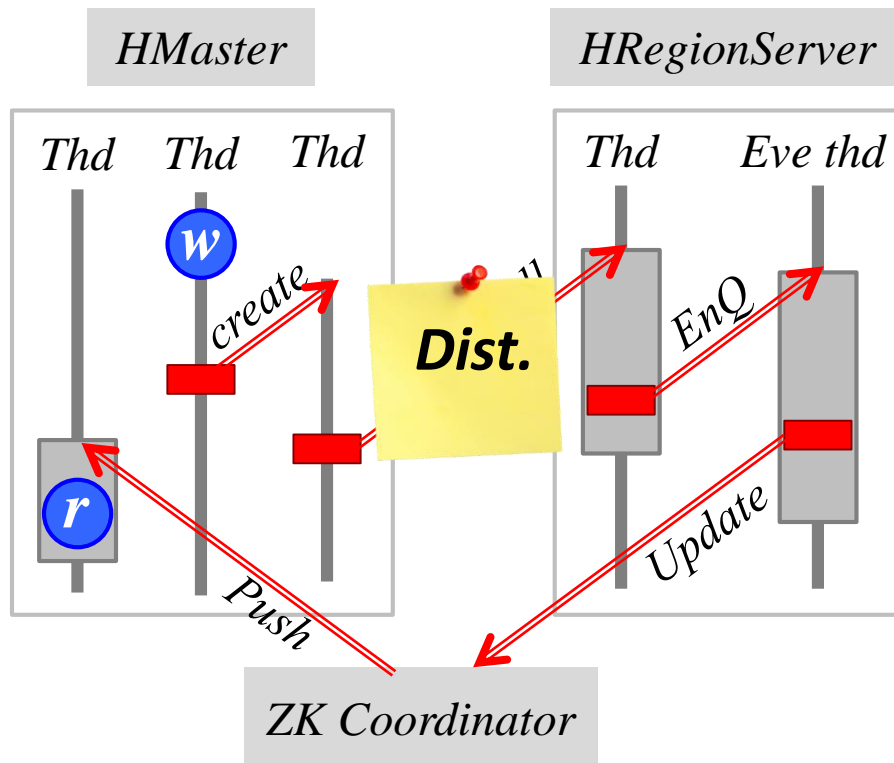


Where is HB model for distributed systems?

DCatch Happens-before Model

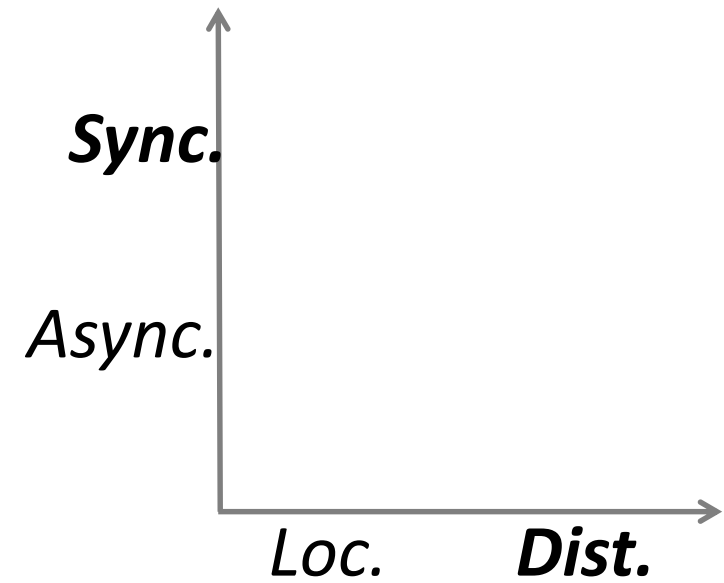
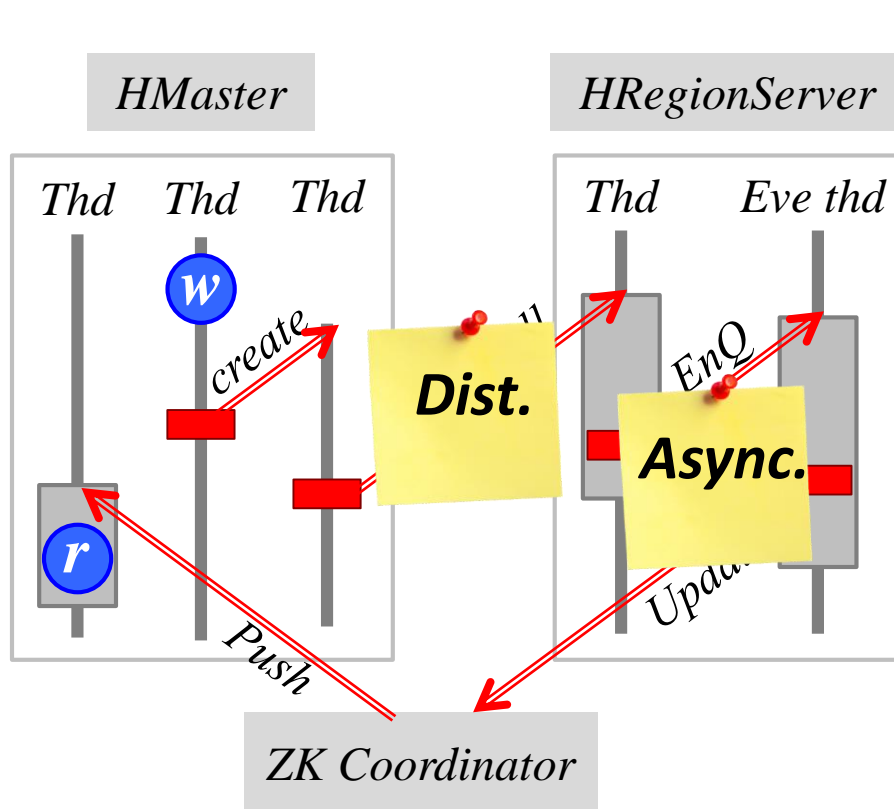


DCatch Happens-before Model

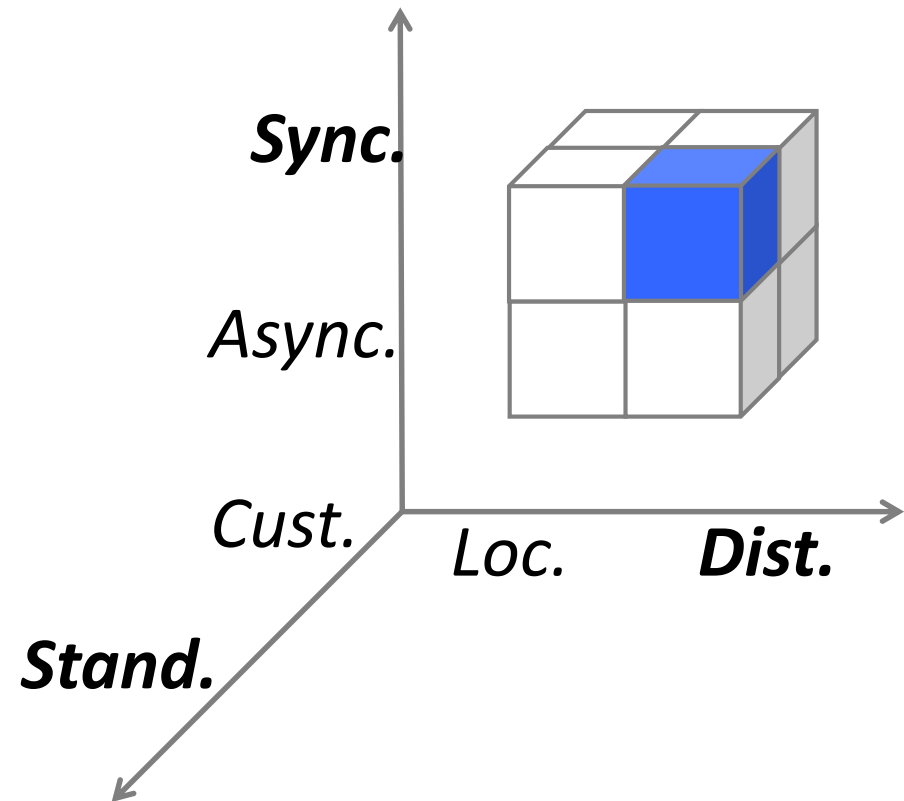
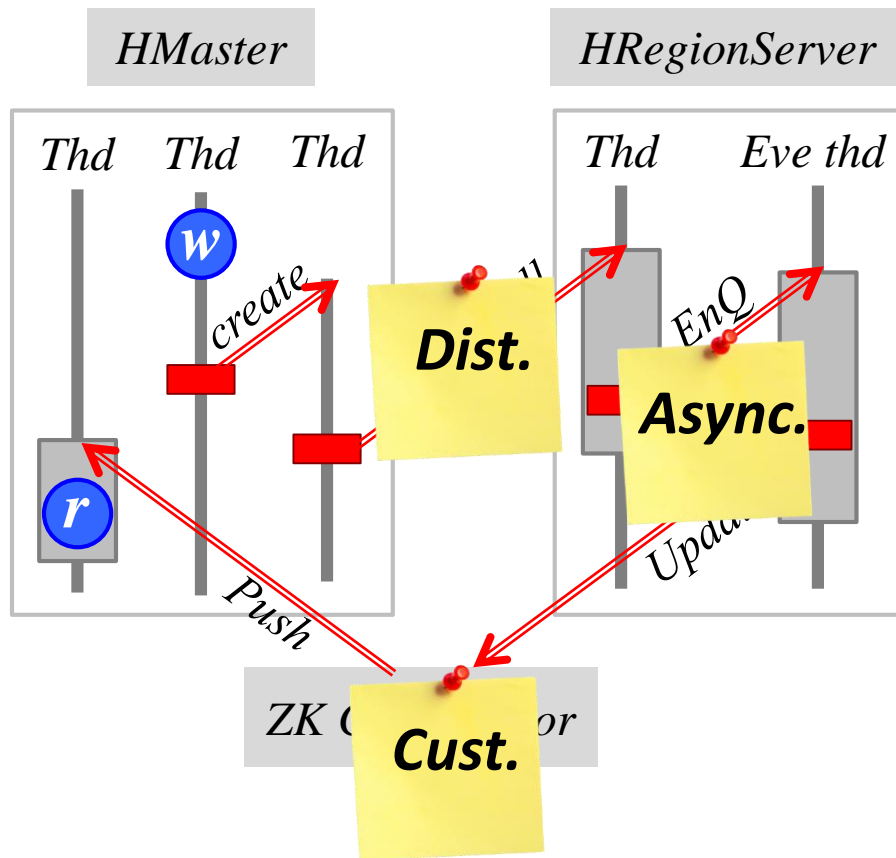


Loc. ***Dist.***

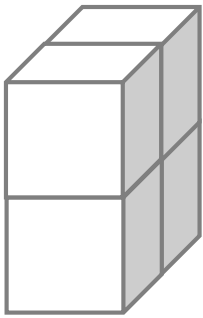
DCatch Happens-before Model



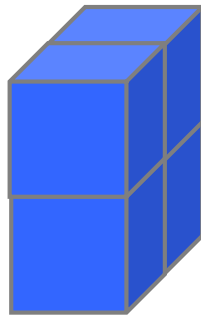
DCatch Happens-before Model



Distributed rules



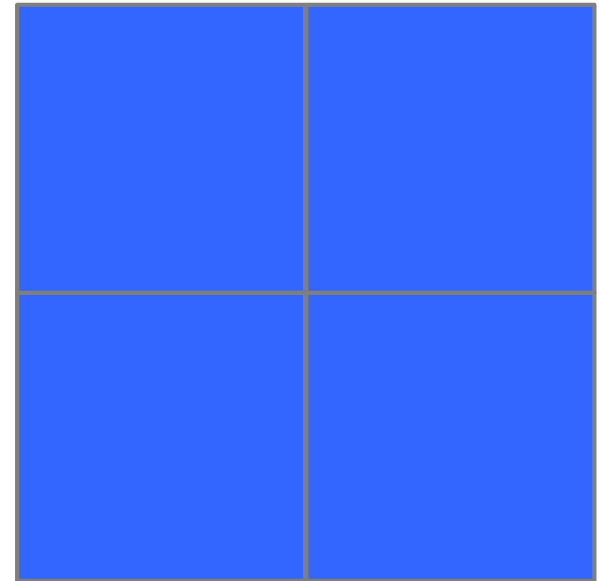
Local



Distributed

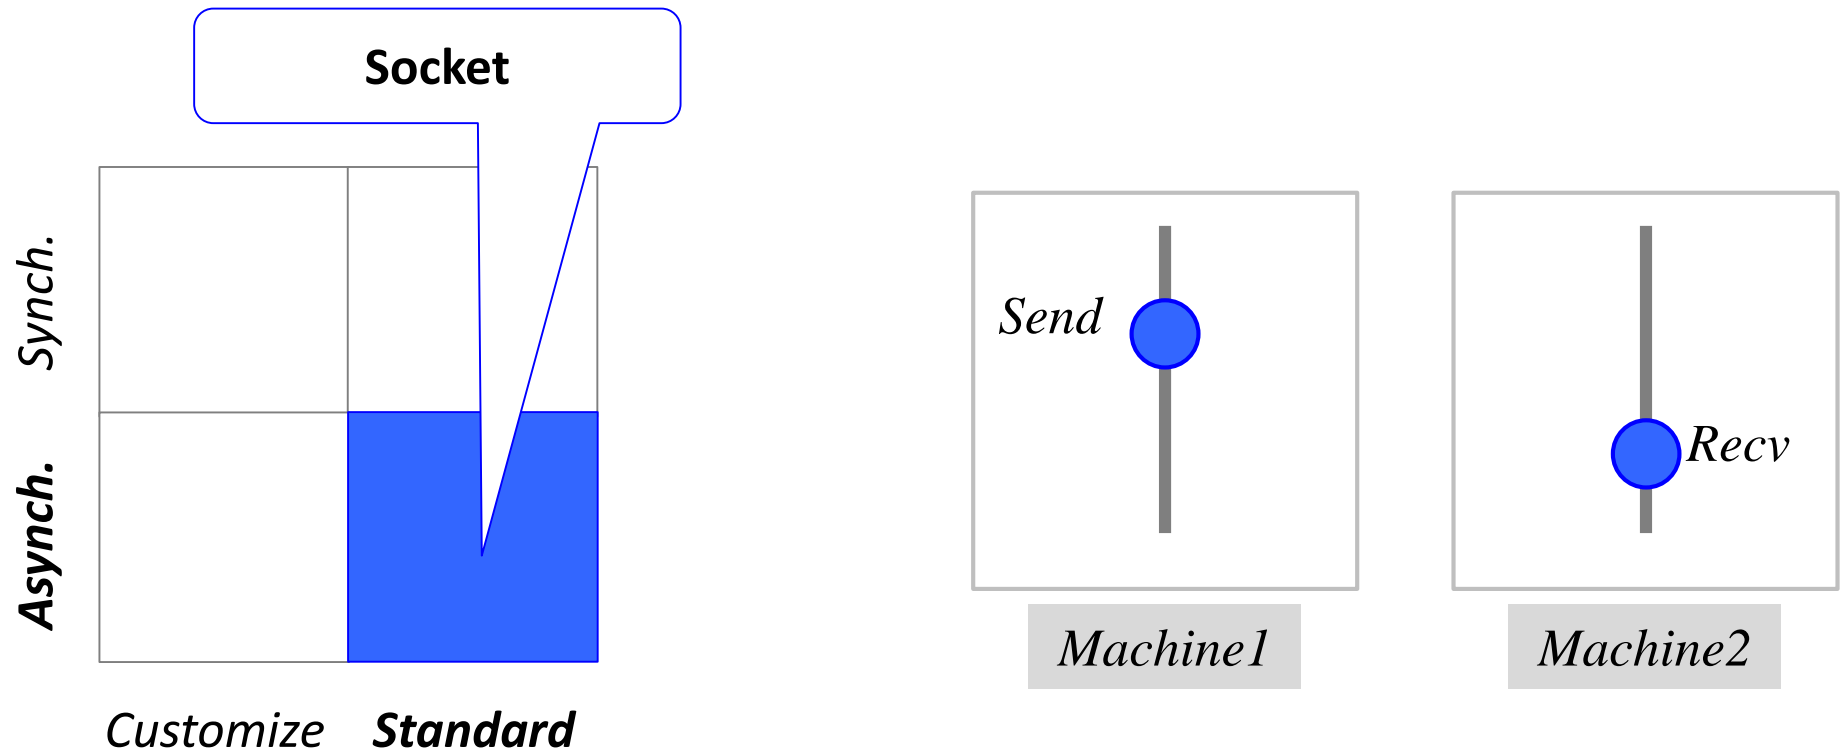
Sync.

Async.



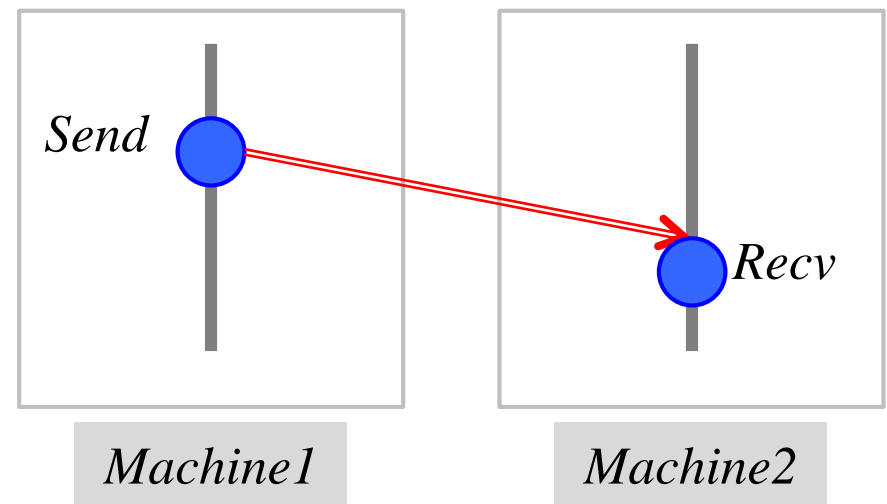
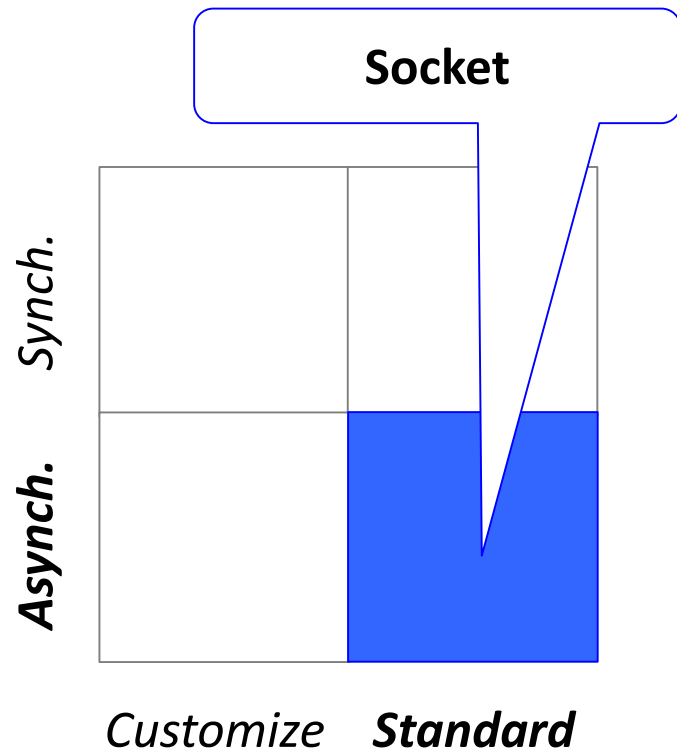
Custom Standard

Distributed rule #1



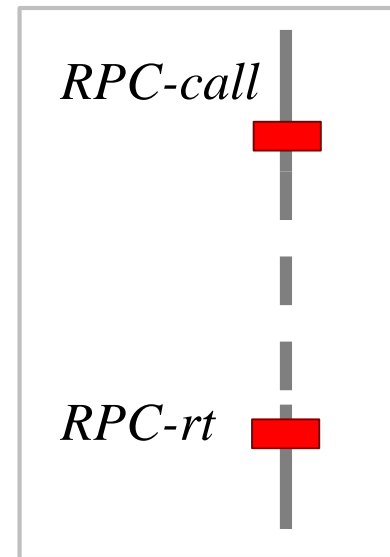
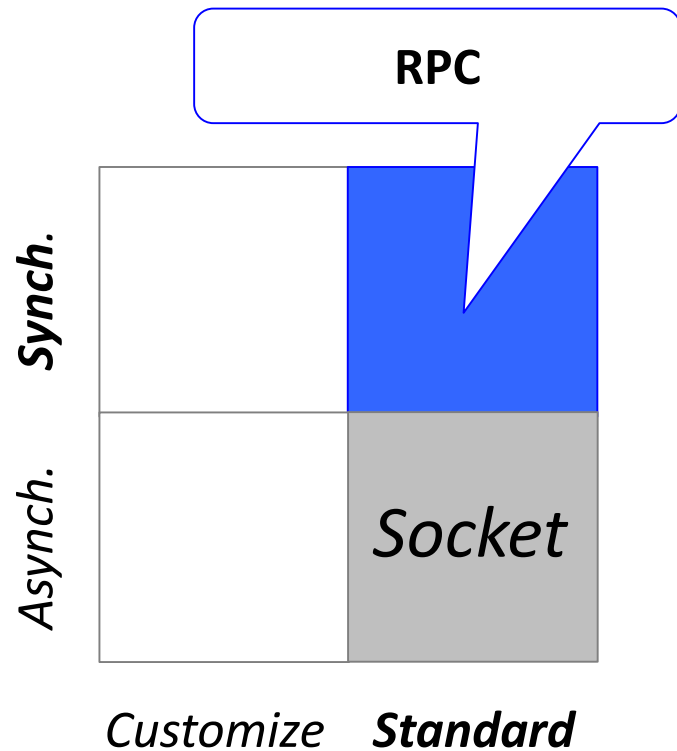
- Logical time clock (Leslie Lamport, 1978)

Distributed rule #1

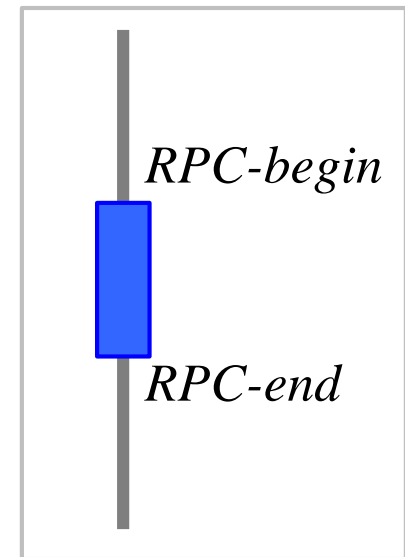


- Logical time clock (Leslie Lamport, 1978)

Distributed rule #2

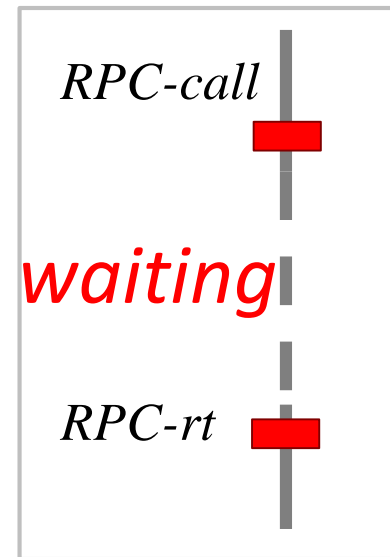
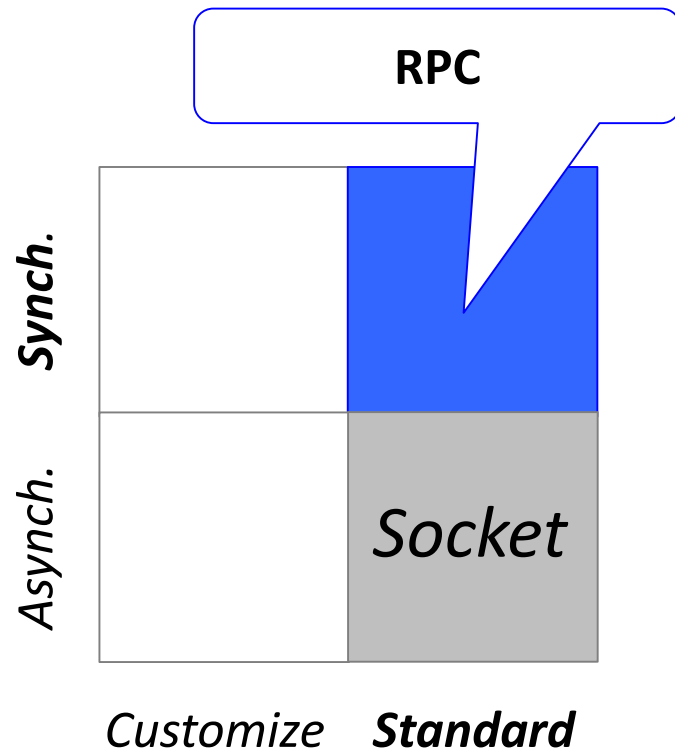


Machine1

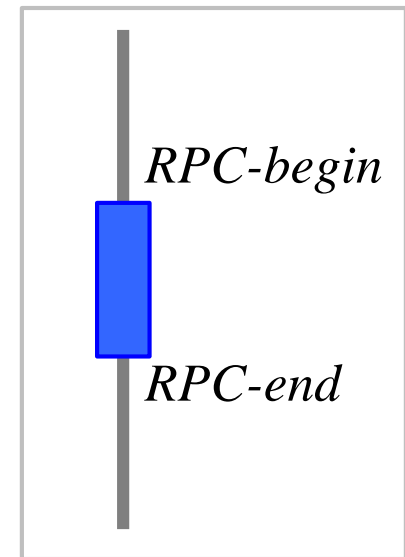


Machine2

Distributed rule #2

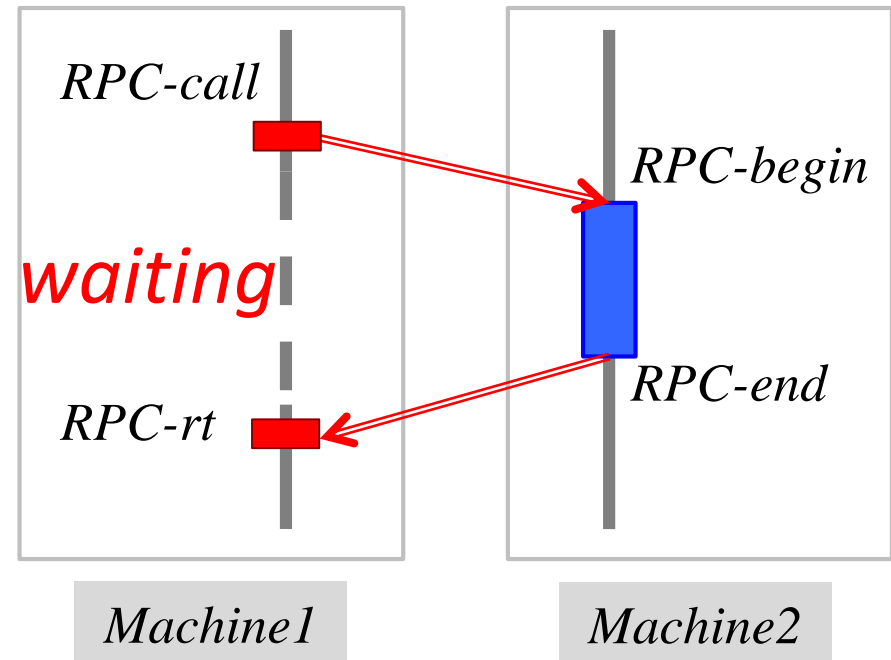
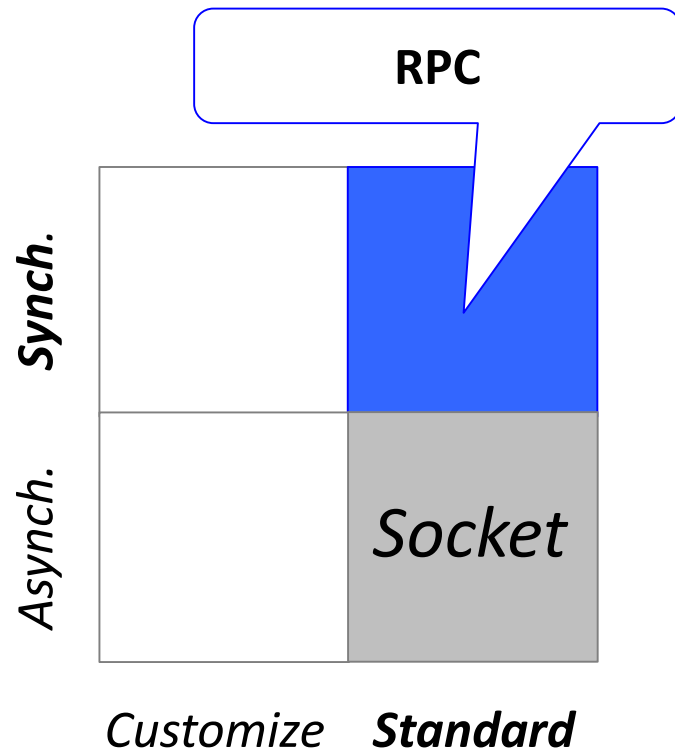


Machine 1



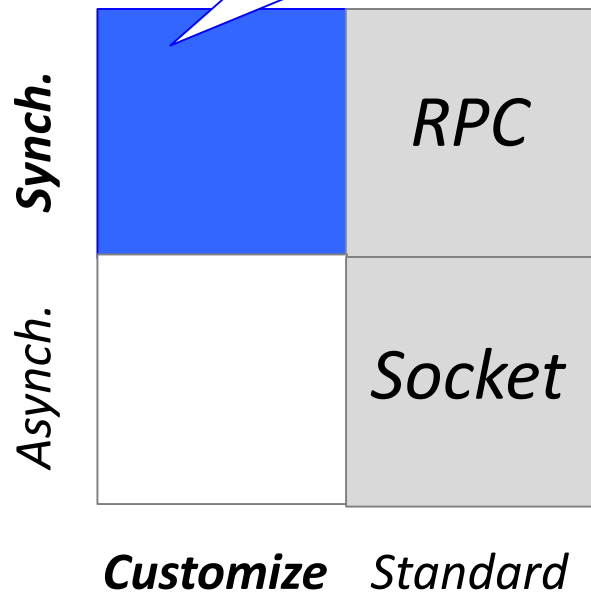
Machine 2

Distributed rule #2



Distributed rule #3

Dist. while-loop



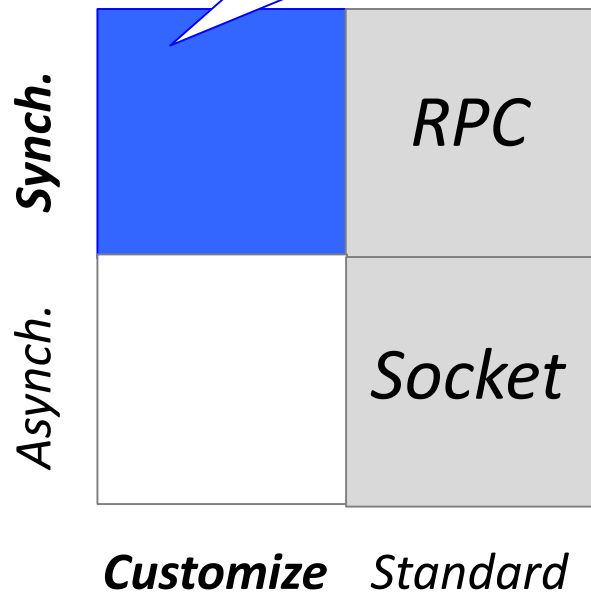
In multi-threaded systems:

```
//Thread1  
flag = True;
```

```
//Thread  
while(!flag){  
}  
...
```

Distributed rule #3

Dist. while-loop



In multi-threaded systems:

```
//Thread1  
flag = True;
```

```
//Thread  
while(!flag){  
}  
...
```

Distributed rule #3

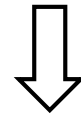
Dist. while-loop

Synch.		RPC
		Socket
Asynch.		
	Customize	Standard

In multi-threaded systems:

```
//Thread1
flag = True;

//Thread
while(!flag){
}
...
```



In distributed systems:

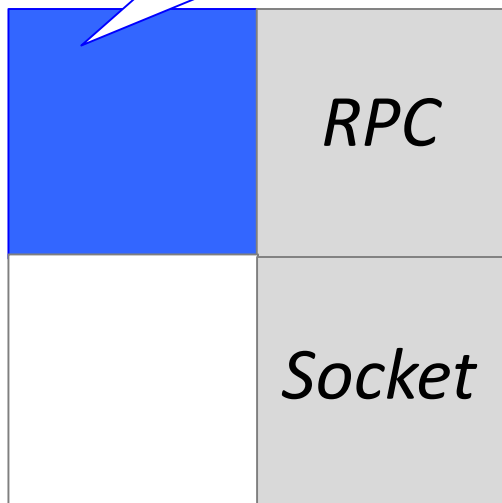


Distributed rule #3

Dist. while-loop

Synch.

Asynch.



Customize Standard

In distributed systems:

Machine A

```
//Thread1  
flag = True;
```

```
//Thread2  
bool getFlag() {  
    return flag;  
}
```

Machine B

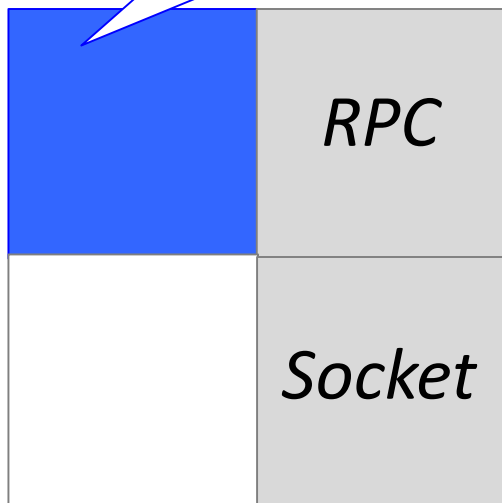
```
//Thread  
while(!getFlag()) {  
}  
...
```

Distributed rule #3

Dist. while-loop

Synch.

Asynch.



Customize

Standard

In distributed systems:

Machine A

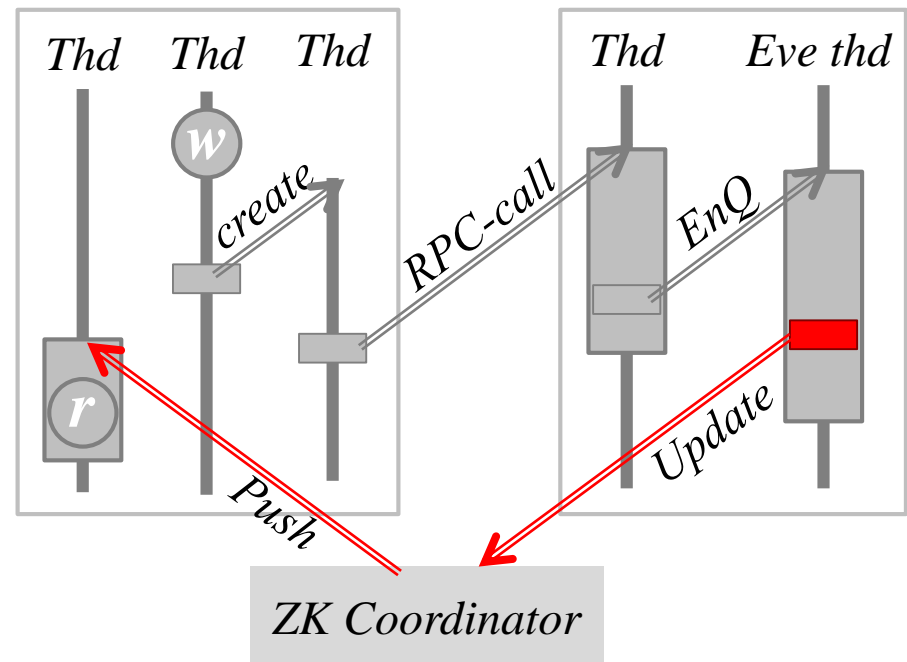
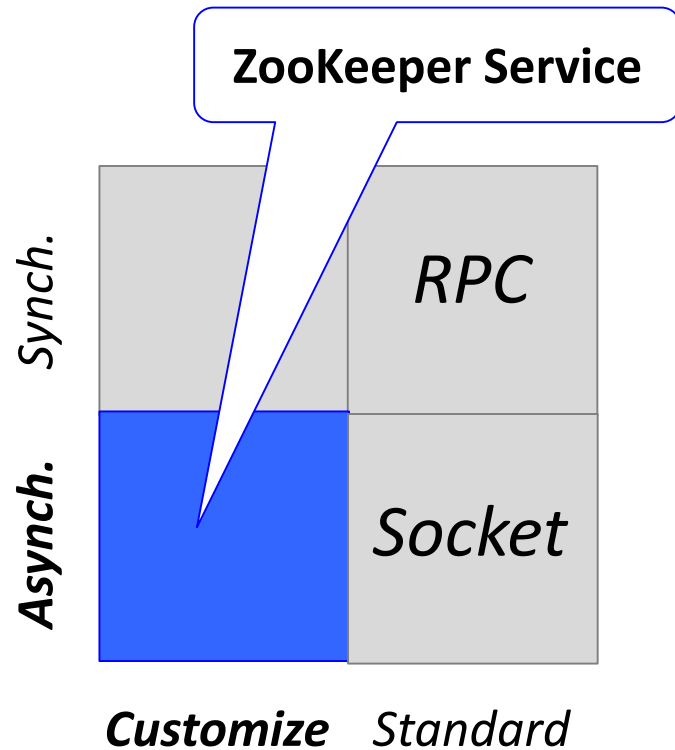
```
//Thread1
flag = True;
```

```
//Thread2
bool getFlag() {
    return flag;
}
```

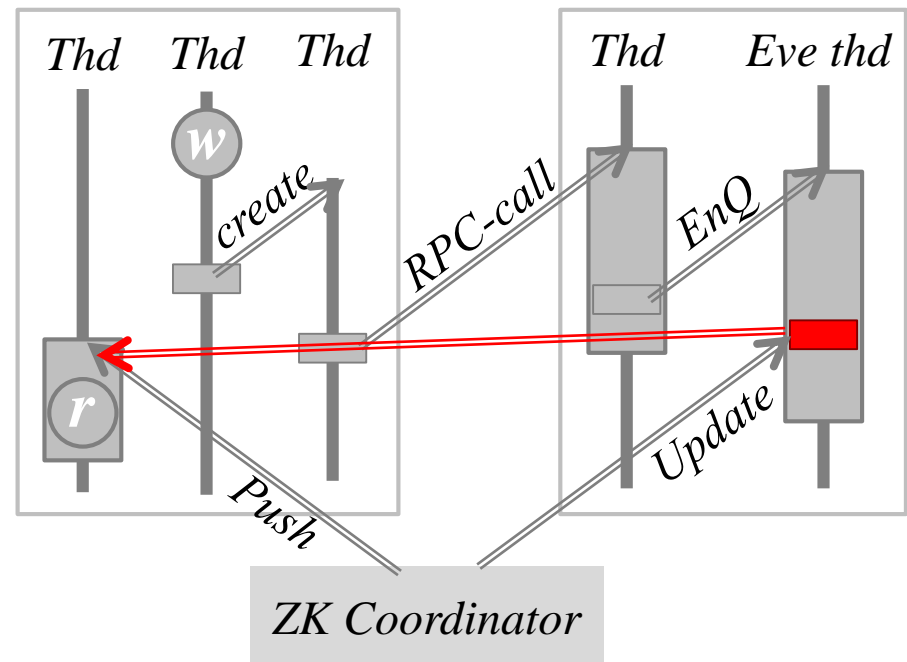
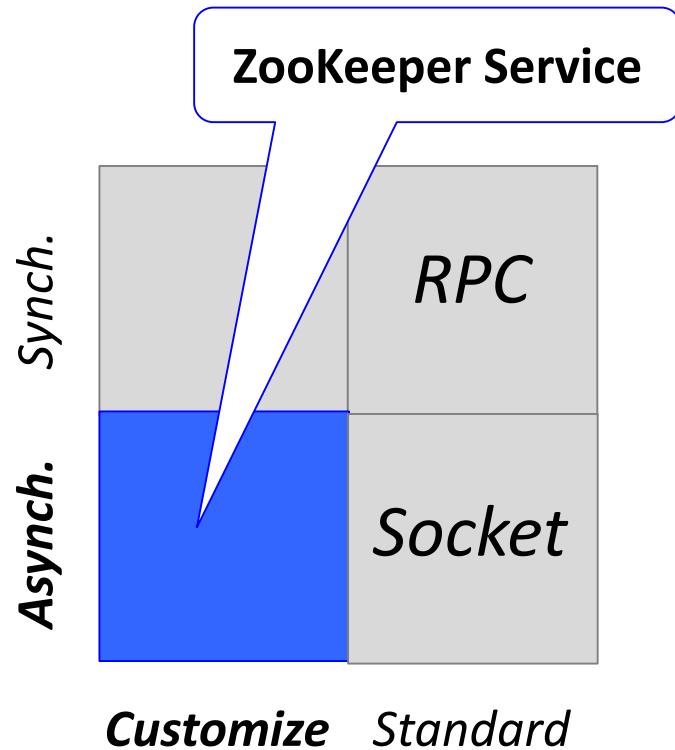
Machine B

```
//Thread
while(!getFlag()) {
}
...
```

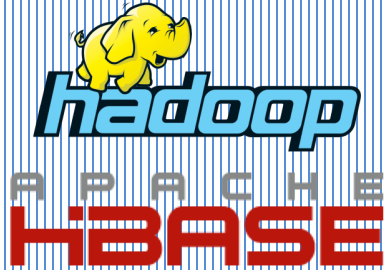
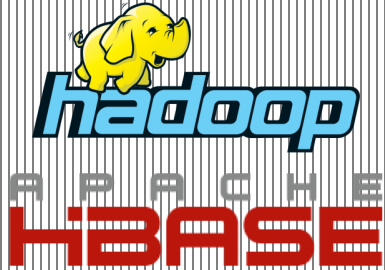


Distributed rule #4



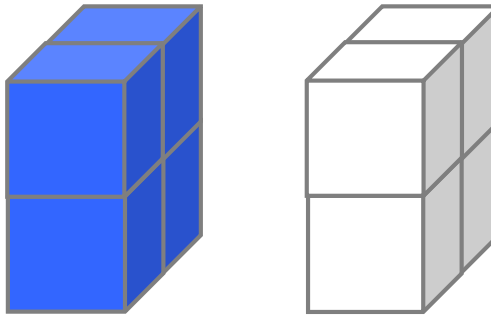
Distributed rule #4



Distributed rules

Synchronous	<p><i>Dist. While-loop</i></p> 	<p><i>RPC</i></p> 
	<p><i>Zookeeper Service</i></p> 	<p><i>Socket</i></p> 
Asynchronous	<i>Customized</i>	<i>Standard</i>

Local rules



Local Distributed




Local rules

Synchronous		
Asynchronous	<i>n/a</i>	<i>Event-related</i>
	<i>Customized</i>	<i>Standard</i>

Local rules

Synchronous	<i>While-loop</i>	<i>Thread fork/join</i>
Asynchronous	<i>n/a</i>	<i>Event-related</i>
	<i>Customized</i>	<i>Standard</i>

Local rules

Synchronous	<p><i>While-loop</i></p> 	<p><i>Thread fork/join</i></p> 
	<p><i>n/a</i></p>	<p><i>Event-related</i></p> 
Asynchronous	<p><i>Customized</i></p>	<p><i>Standard</i></p>

Outline

- Motivation
- DCatch Happens-before Model
- **DCatch tool**
- Evaluation
- Conclusion

Trace

HB

Triage

Trigger

Trace

HB

Triage

Trigger

C1: How to handle the huge amount of mem accesses?

C2: What's the happens-before model?

C3: How to estimate the distributed impact of a race?

C4: How to trigger with distributed time manipulation?

Challenges

Selective tracing: only mem accesses in Event/message handlers and their callee.

Trace

HB

Triage

Trigger

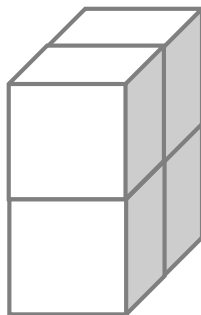
C1: How to handle the huge amount of mem accesses?

C2: What's the happens-before model?

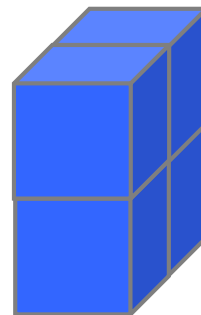
C3: How to estimate the distributed impact of a race?

C4: How to trigger with distributed time manipulation?

Challenges



Local



Distributed

Trace

HB

Triage

Trigger

C1: How to handle the huge amount of mem accesses?

C2: What's the happens-before model?

C3: How to estimate the distributed impact of a race?

C4: How to trigger with distributed time manipulation?

Challenges

Machine B

```
//RPC thread
Task getTask(jID) {
    ...
    return jMap.get(jID);
}
```

```
//UnReg thread
void unReg(jID) {
    jMap.remove(jID);
    ....
}
```

Trace

HB

Triage

Trigger

C1: How to handle the huge amount of mem accesses?

C2: What's the happens-before model?

C3: How to estimate the distributed impact of a race?

C4: How to trigger with distributed time manipulation?

Challenges

Machine A

```
while(!getTask(jID)) {  
}
```

Machine B

```
//RPC thread  
Task getTask(jID) {  
  ...  
  return jMap.get(jID);  
}
```

```
//UnReg thread  
void unReg(jID) {  
  jMap.remove(jID);  
  ....  
}
```

Trace

HB

Triage

Trigger

C1: How to handle the huge amount of mem accesses?

C2: What's the happens-before model?

C3: How to estimate the distributed impact of a race?

C4: How to trigger with distributed time manipulation?

Challenges

Machine A

```
while(!getTask(jID)) {  
}
```

Machine B

```
//RPC thread  
Task getTask(jID) {  
    ...  
    return jMap.get(jID);  
}
```

```
//UnReg thread  
void unReg(jID) {  
    jMap.remove(jID);  
    ....  
}
```

Trace

HB

Triage

Trigger

C1: How to handle the huge amount of mem accesses?

C2: What's the happens-before model?

C3: How to estimate the distributed impact of a race?

C4: How to trigger with distributed time manipulation?

Challenges

Machine A

Machine B

```
while (!getTask(jID)) {
}
```


hang

```
//RPC thread
Task getTask(jID) {
    ...
    return jMap.get(jID);
}
```

```
//UnReg thread
void unReg(jID) {
    jMap.remove(jID);
    ....
}
```

Trace

HB

Triage

Trigger

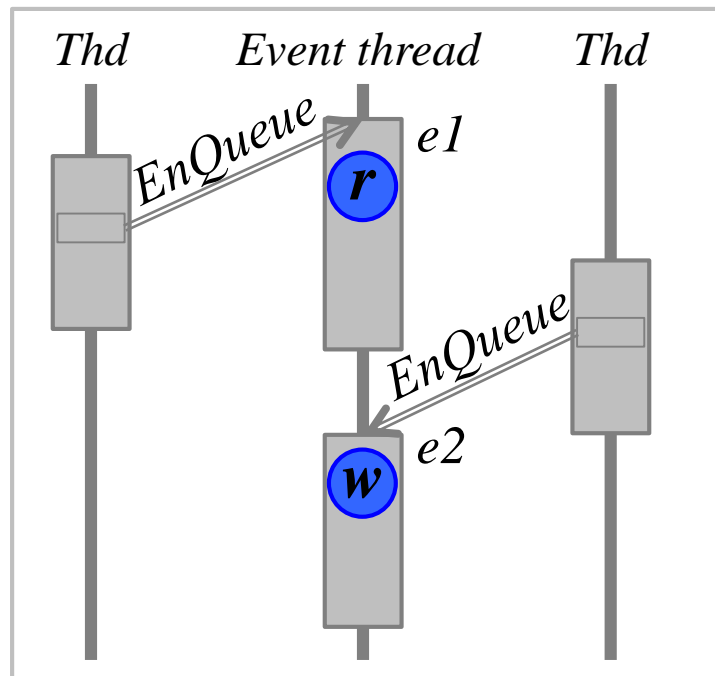
C1: How to handle the huge amount of mem accesses?

C2: What's the happens-before model?

C3: How to estimate the distributed impact of a race?

C4: How to trigger with distributed time manipulation?

Challenges



Trace

HB

Triage

Trigger

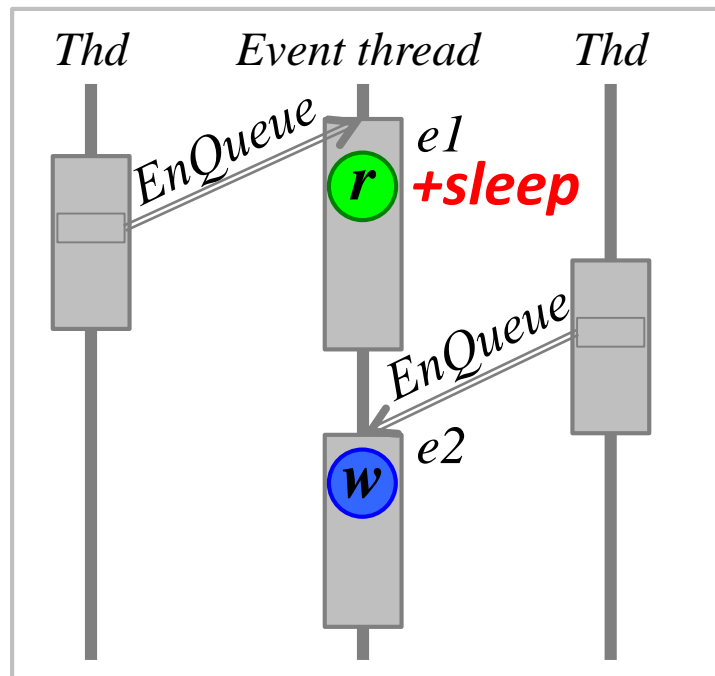
C1: How to handle the huge amount of mem accesses?

C2: What's the happens-before model?

C3: How to estimate the distributed impact of a race?

C4: How to trigger with distributed time manipulation?

Challenges



Trace

HB

Triage

Trigger

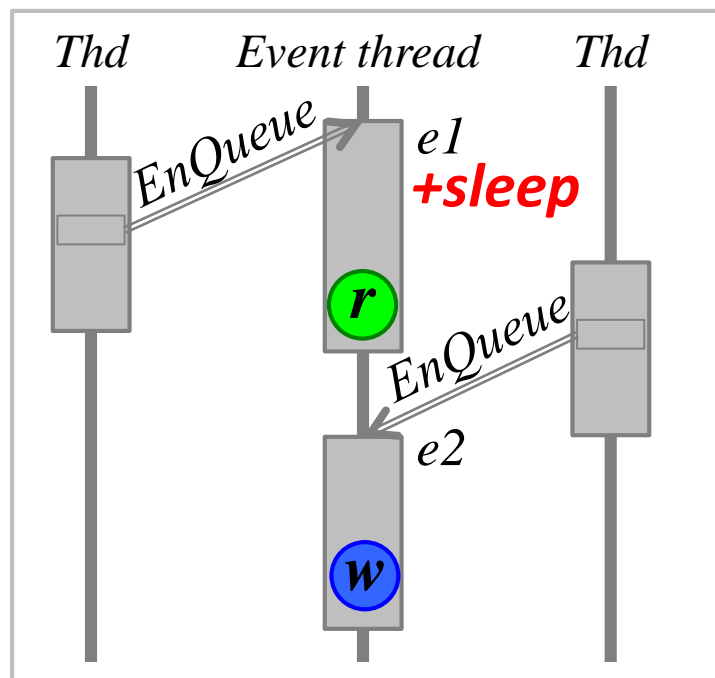
C1: How to handle the huge amount of mem accesses?

C2: What's the happens-before model?

C3: How to estimate the distributed impact of a race?

C4: How to trigger with distributed time manipulation?

Challenges



Trace

HB

Triage

Trigger

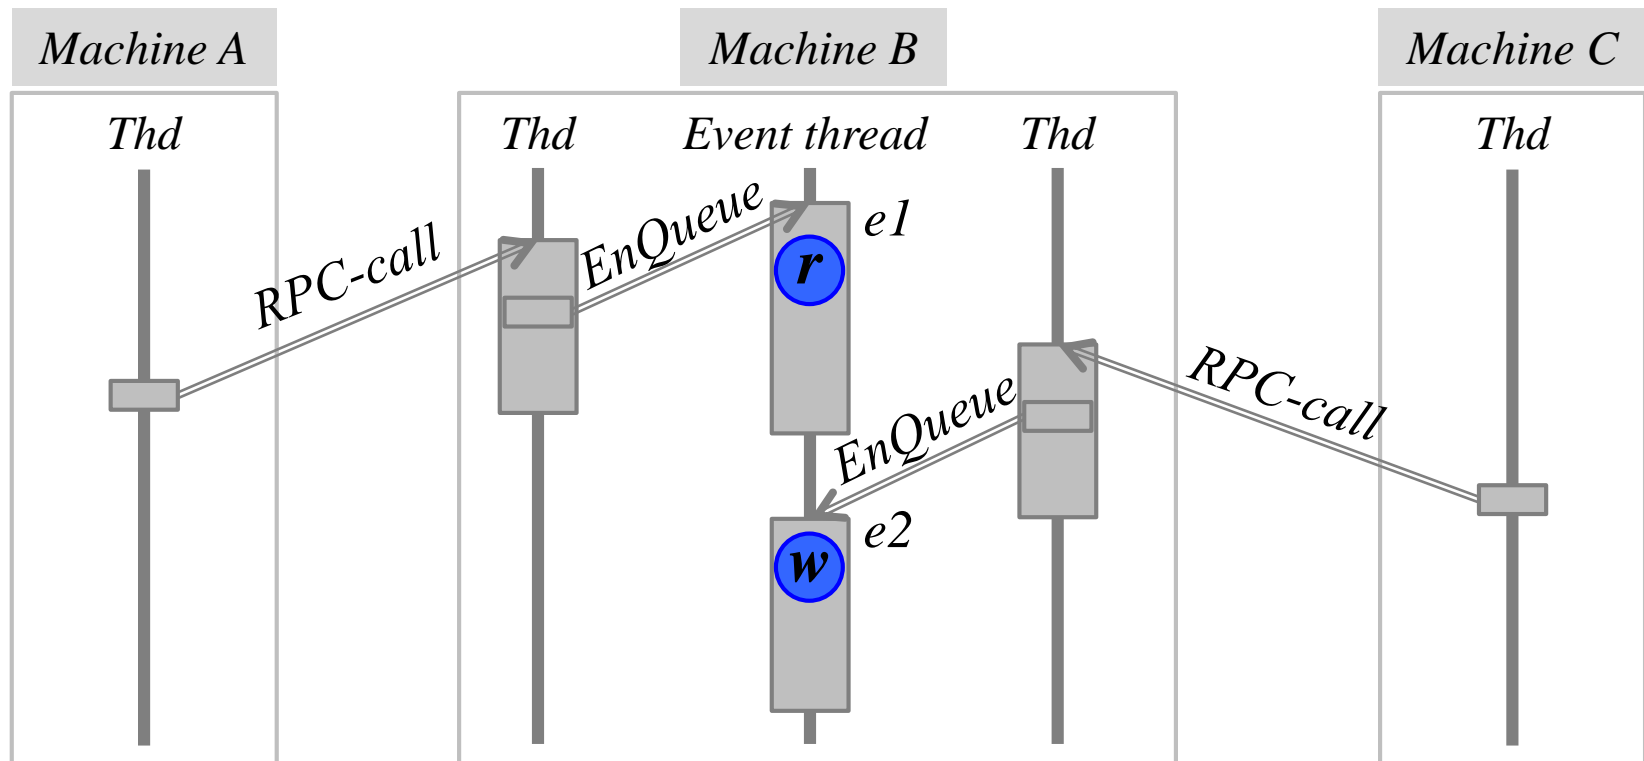
C1: How to handle the huge amount of mem accesses?

C2: What's the happens-before model?

C3: How to estimate the distributed impact of a race?

C4: How to trigger with distributed time manipulation?

Challenges



Trace

HB

Triage

Trigger

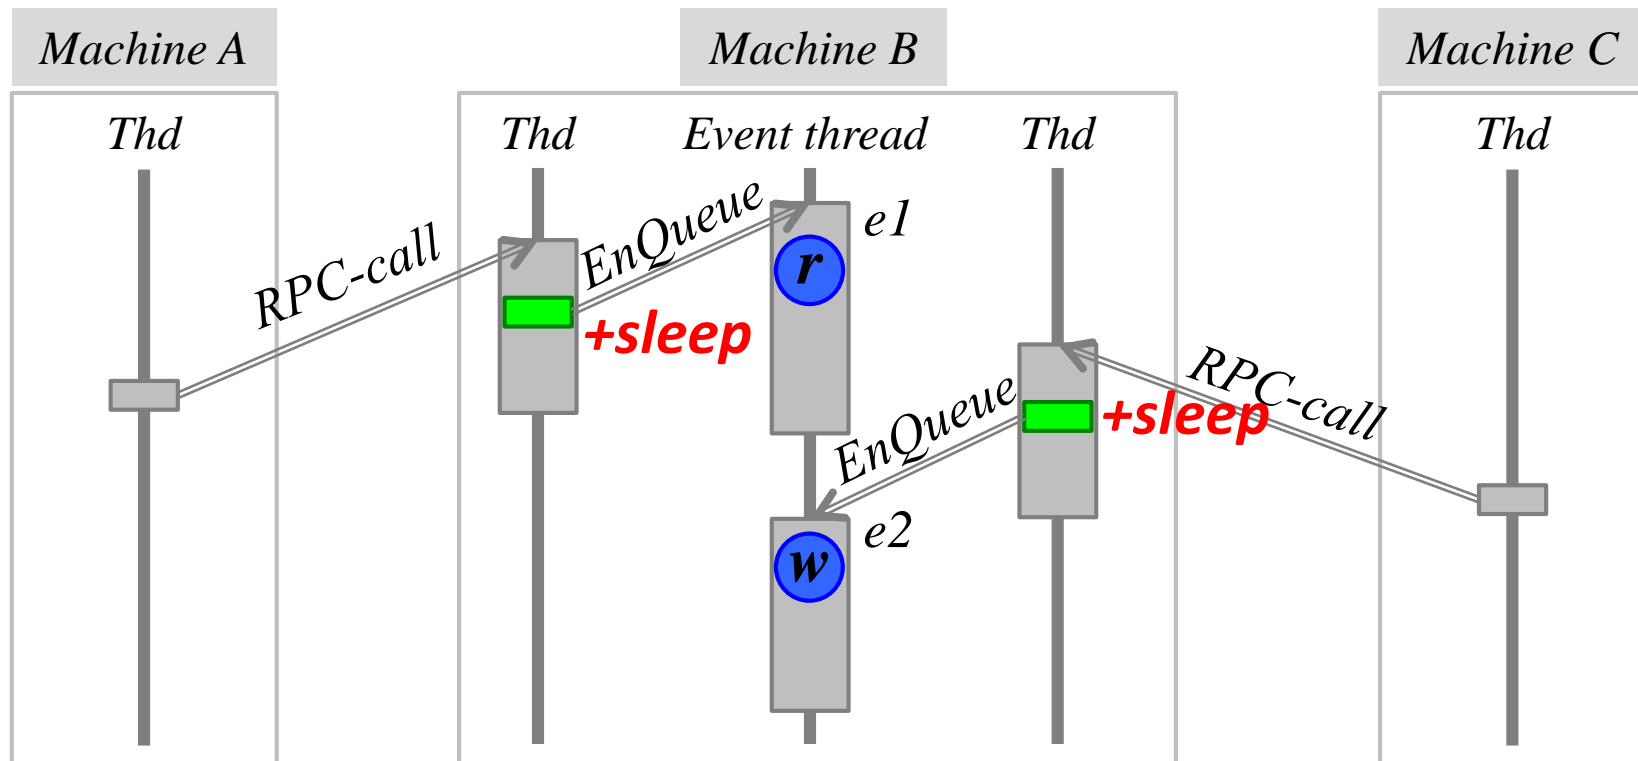
C1: How to handle the huge amount of mem accesses?

C2: What's the happens-before model?

C3: How to estimate the distributed impact of a race?

C4: How to trigger with distributed time manipulation?

Challenges



Trace

HB

Triage

Trigger

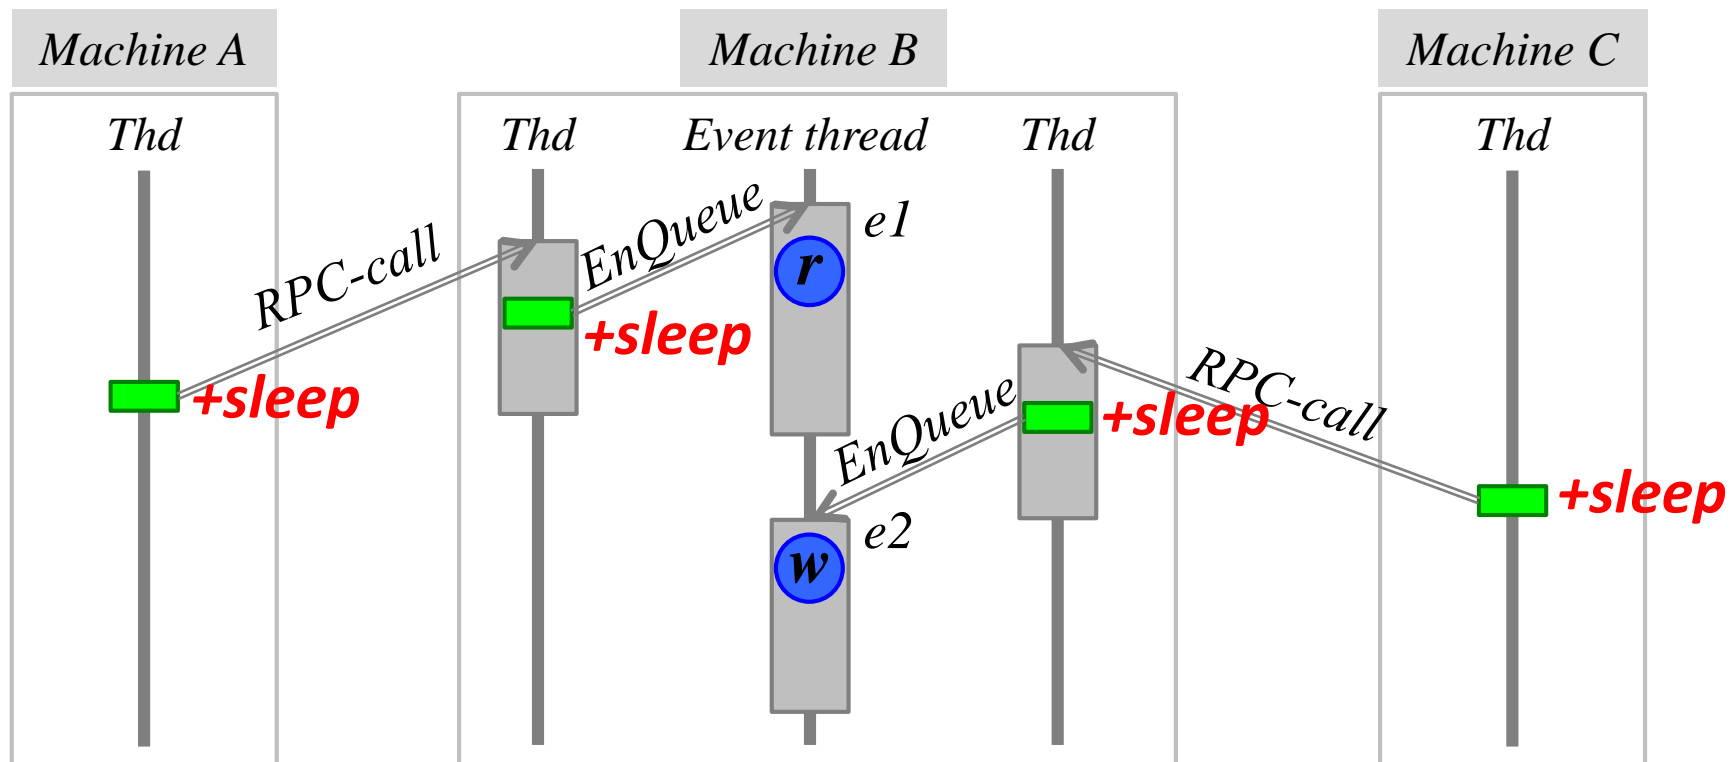
C1: How to handle the huge amount of mem accesses?

C2: What's the happens-before model?

C3: How to estimate the distributed impact of a race?

C4: How to trigger with distributed time manipulation?

Challenges



Outline

- Motivation
- DCatch Happens-before Model
- DCatch tool
- **Evaluation**
- Conclusion

Methodology

- Benchmarks:
 - 7 real-world DCbugs from TaxDC[1]
 - 4 distributed systems



[1] Leesatapornwongsa. TaxDC. In ASPLOS'16

Overall results

BugID	Detected?	#. Bugs	#. Benign	#. false-pos
<i>CA-1011</i>	✓	3	0	0
<i>HB-4539</i>	✓	3	0	1
<i>HB-4729</i>	✓	4	1	0
<i>MR-3274</i>	✓	2	0	4
<i>MR-4637</i>	✓	1	2	4
<i>ZK-1144</i>	✓	5	1	1
<i>ZK-1270</i>	✓	6	2	0

Overall results

BugID	Detected?	#. Bugs	#. Benign	#. false-pos
CA-1011	✓	3	0	0
HB-4539	✓	3	0	1
HB-4729	✓	4	1	0
MR-3274	✓	2	0	4
MR-4637	✓	1	2	4
ZK-1144	✓	5	1	1
ZK-1270	✓	6	2	0

Overall results

BugID	Detected?	#. Bugs	#. Benign	#. false-pos
CA-1011	✓	3	0	0
HB-4539	✓	3	0	1
HB-4729	✓	4	1	0
MR- 3274	✓	2	0	4
MR- 4637	✓	1	2	4
ZK-1144	✓	5	1	1
ZK-1270	✓	6 = 12 + 8	2	0

Overall results

BugID	Detected?	#. Bugs	#. Benign	#. false-pos
<i>CA-1011</i>	✓	3	0	0
<i>HB-4539</i>	✓	3	0	1
<i>HB-4729</i>	✓	4	1	0
<i>MR-3274</i>	✓	2	0	4
<i>MR-4637</i>	✓	1	2	4
<i>ZK-1144</i>	✓	5	1	1
<i>ZK-1270</i>	✓	6	2	0

Other results in our paper

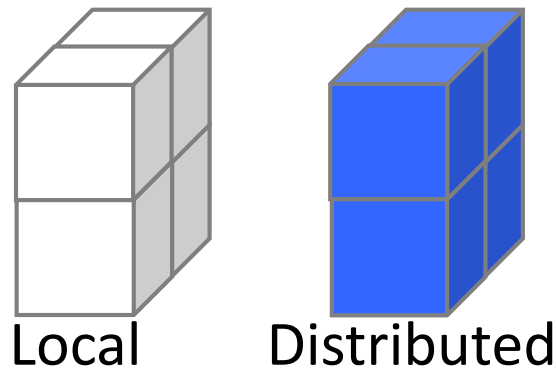
- Performance overhead
- Trace compositions
- HB model impact
 - False-positive
 - False-negatives
- ...

Outline

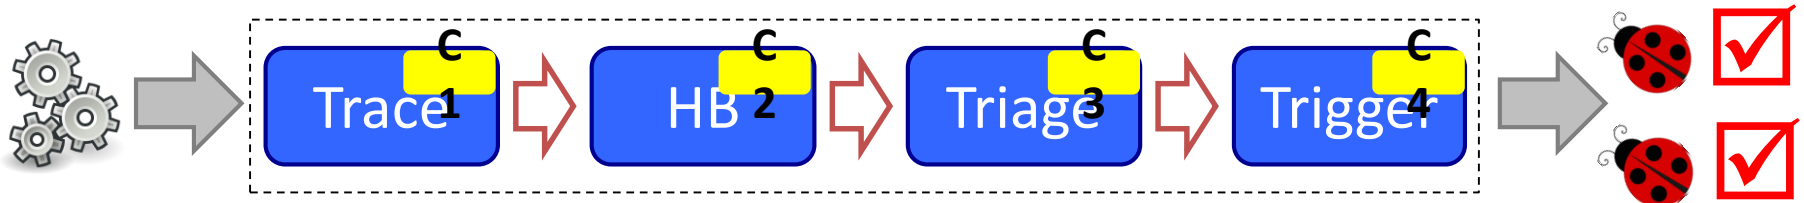
- Motivation
- DCatch Happens-before Model
- DCatch tool
- Evaluation
- Conclusion

Conclusion

- A **HB Model** for distributed systems



- DCatch** detects DCbugs from **correct runs** with **low false positive** rates.



Thank you!
Q&A