


Finding and Evaluating the Performance Impact of Redundant Data Access for Applications that are Developed Using Object-Relational Mapping Frameworks



Tse-Hsun(Peter) Chen



Weiyl Shang



Zhen Ming Jiang



Ahmed E. Hassan



Mohamed Nasser, Parminder Flora



\$1.6 billion revenue loss for a one-second slowdown

A close-up photograph of a brown cardboard box. The Amazon.com logo is printed on the box in a dark blue or black ink. The logo consists of the text "amazon.com" in a sans-serif font, with a curved arrow underneath the "a" that points towards the "n". Below the main text, the phrase "and you're done." is printed in a smaller, lighter font.

amazon.com
and you're done.

Slow database access is often the performance bottleneck



Object-Relational Mapping abstracts relational database as objects



Benefits of using ORM

- Less **boilerplate code**
- Automatically **manage object-DB translations**

Using ORM results in much less code and shorter development time.

Over 67% of Java developers use Object-Relational Mapping (Hibernate) to access databases

Two most popular Java database access frameworks

2001~Now



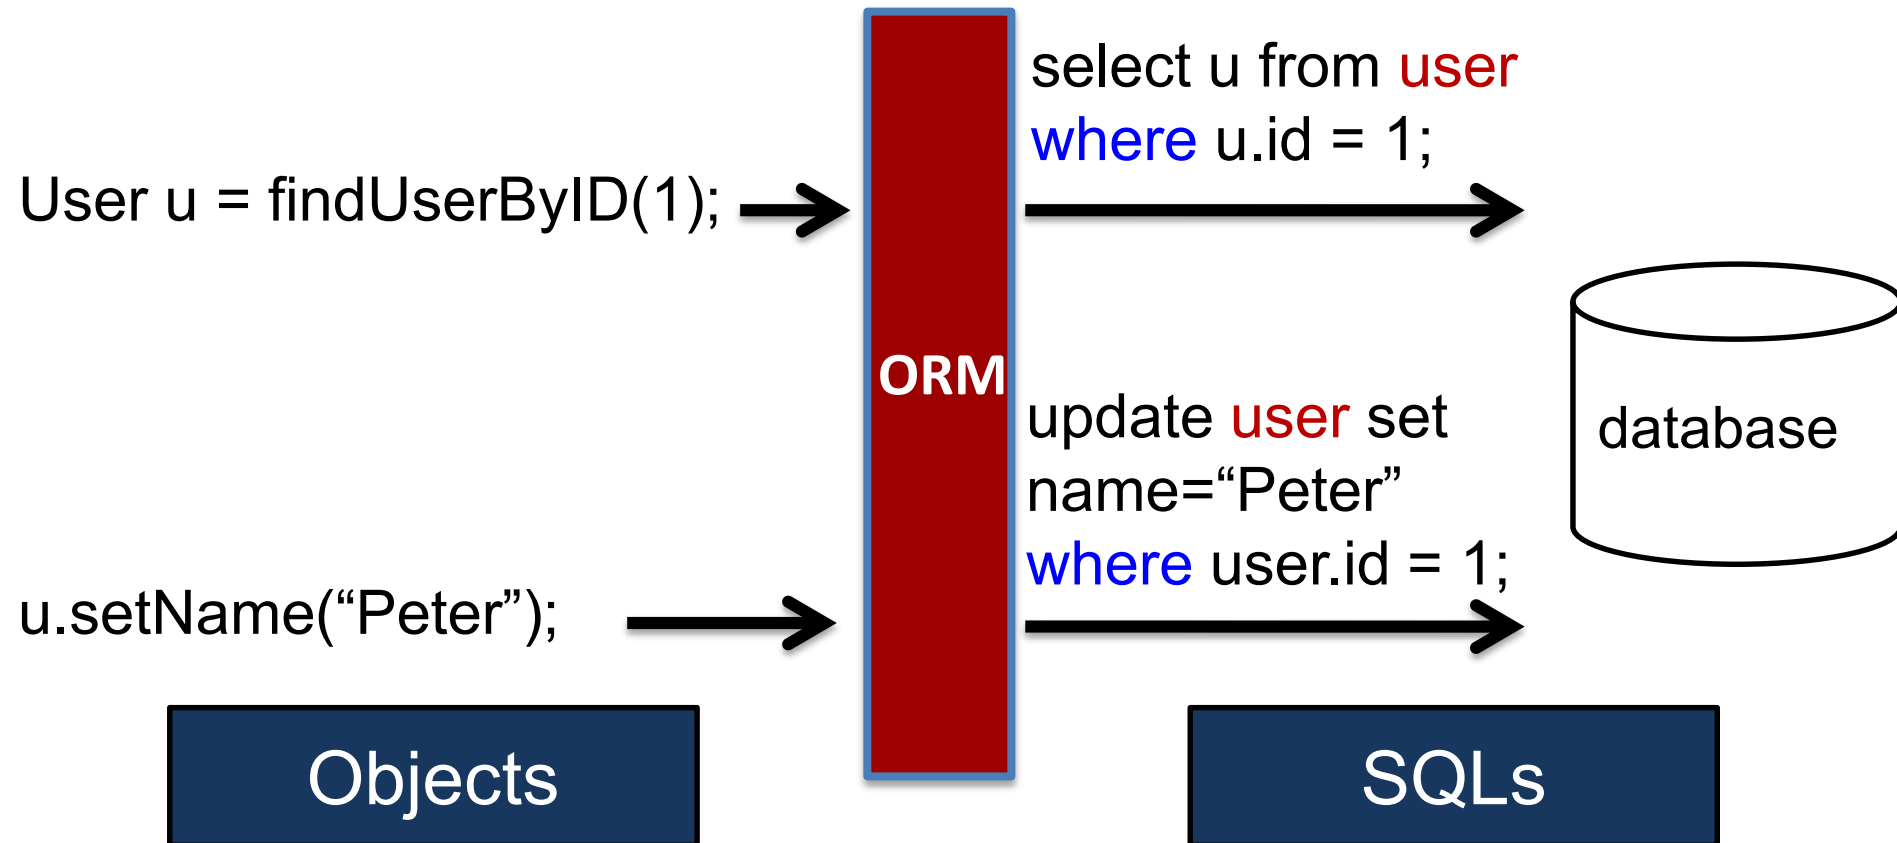
67%

1997~Now




22%

Accessing the database using ORM



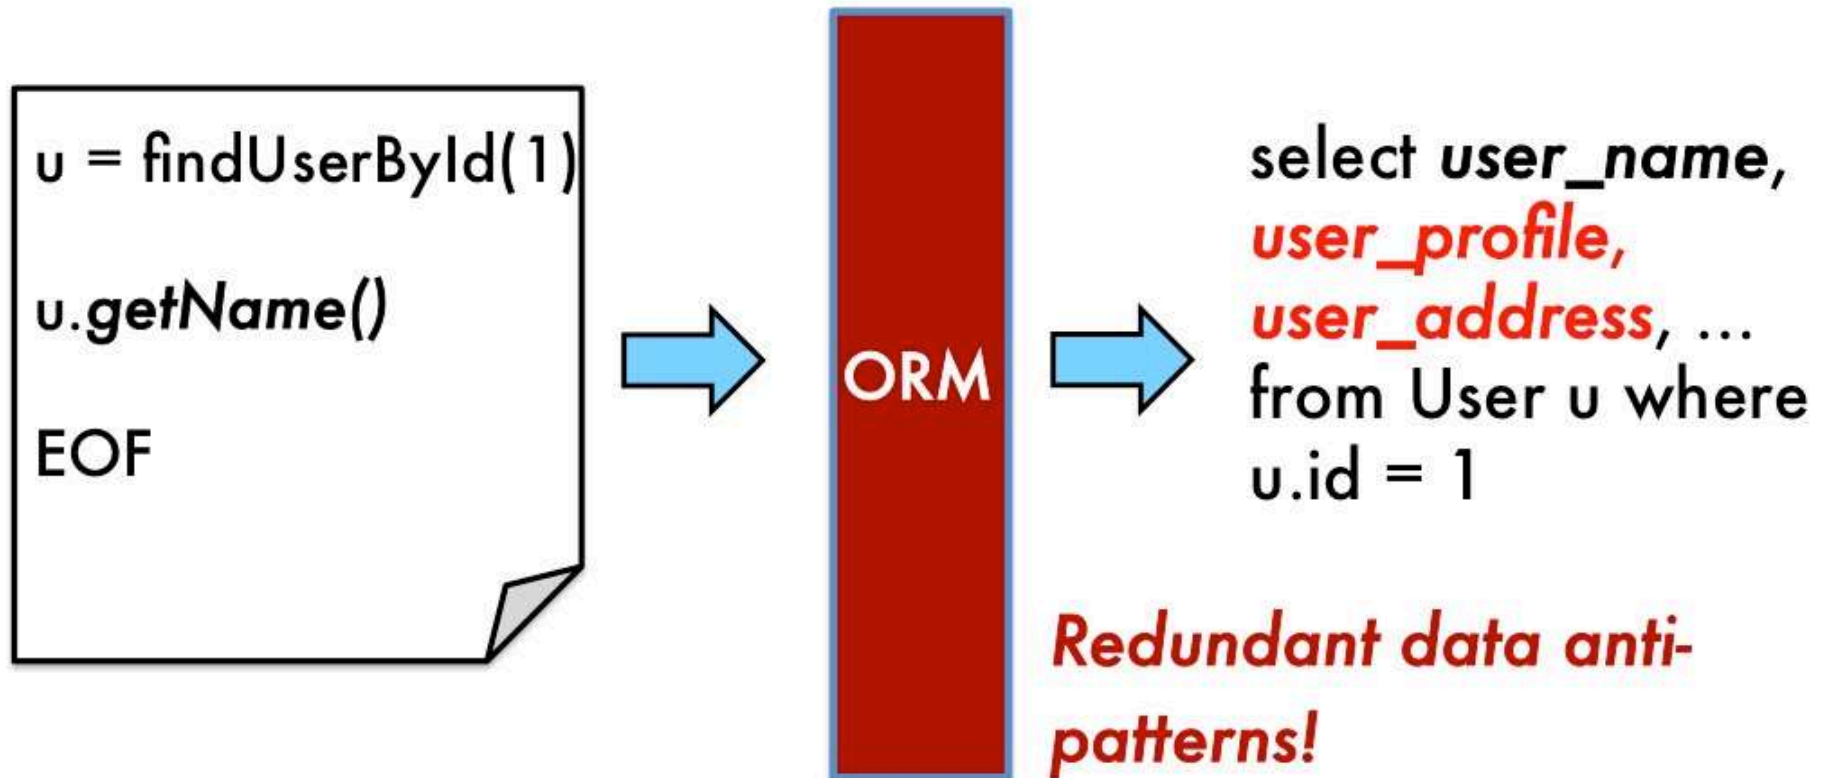
Redundant data is caused by ORM's lack of knowledge on the business logic

Redundant data is common across ORMs!




We need to understand how ORM APIs are used in the application in order to detect redundant data anti-patterns.

Redundant data is caused by ORM's lack of knowledge on the business logic

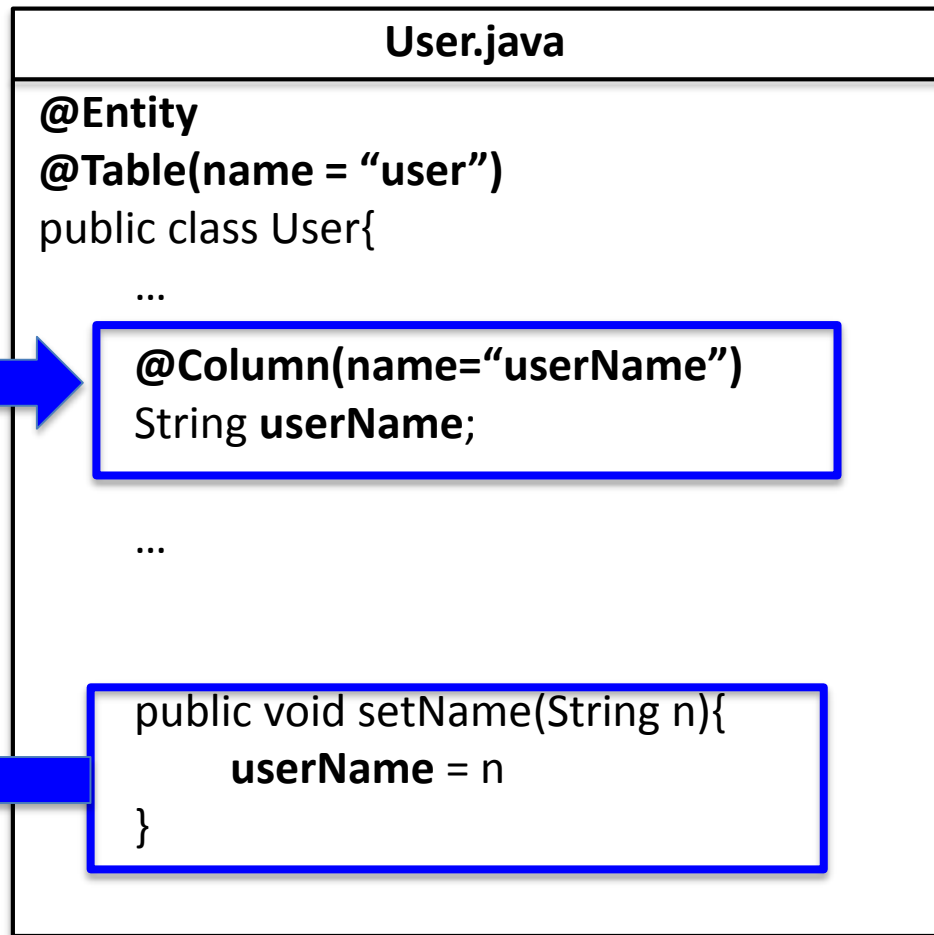


We need to understand how ORM APIs are used in the application in order to detect redundant data anti-patterns.

Our approach for finding redundant data anti-patterns

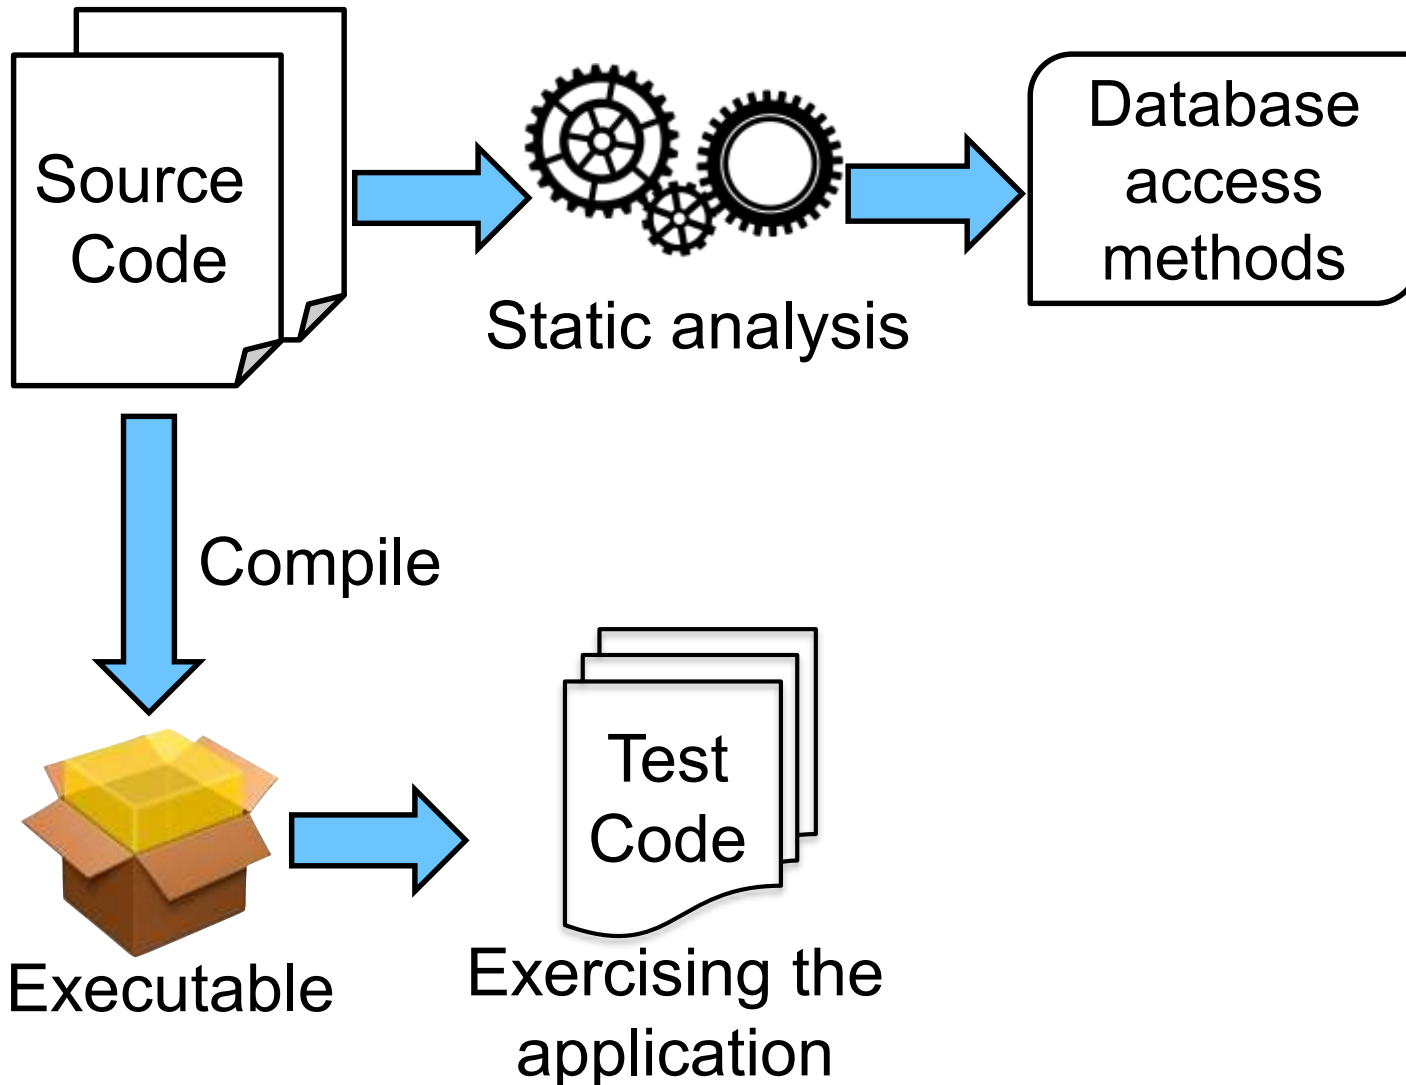


Finding database access methods

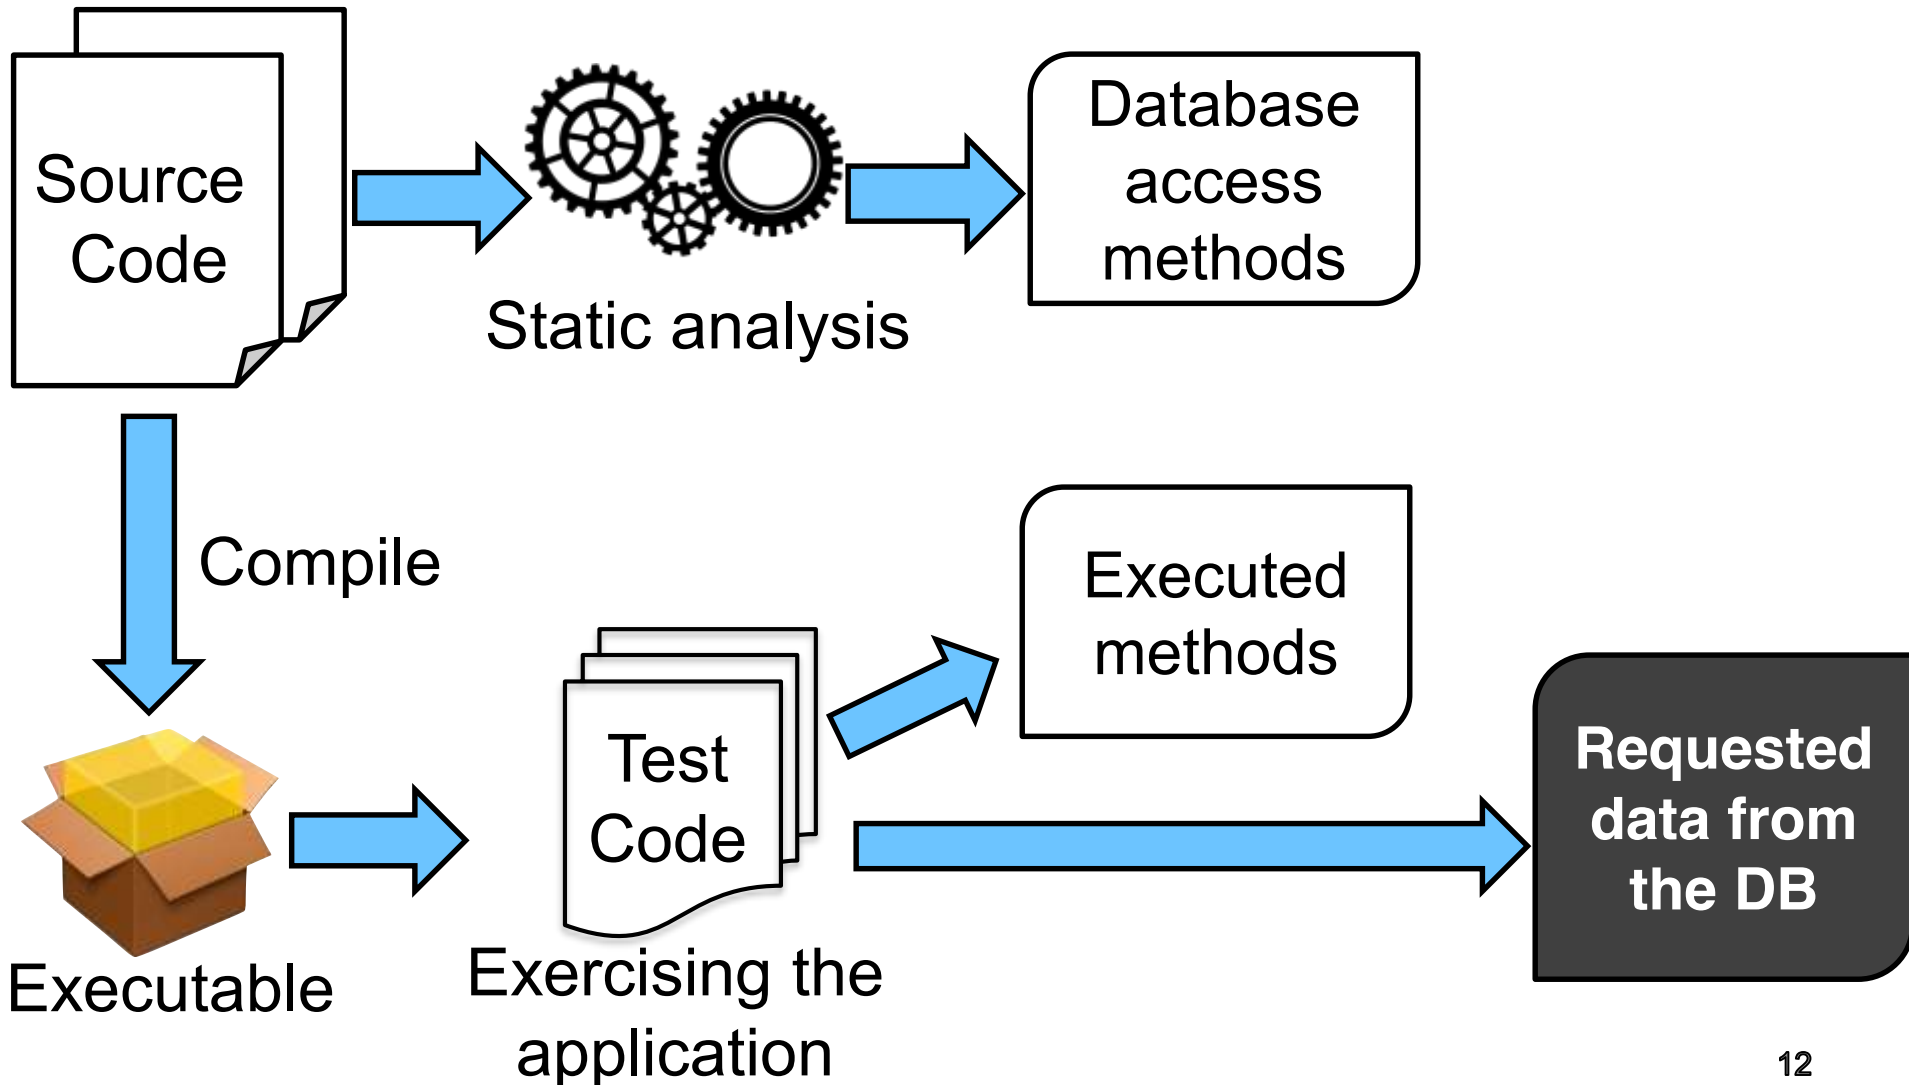


Using static analysis to find which database column a method is reading or modifying.


Our approach for finding redundant data anti-patterns



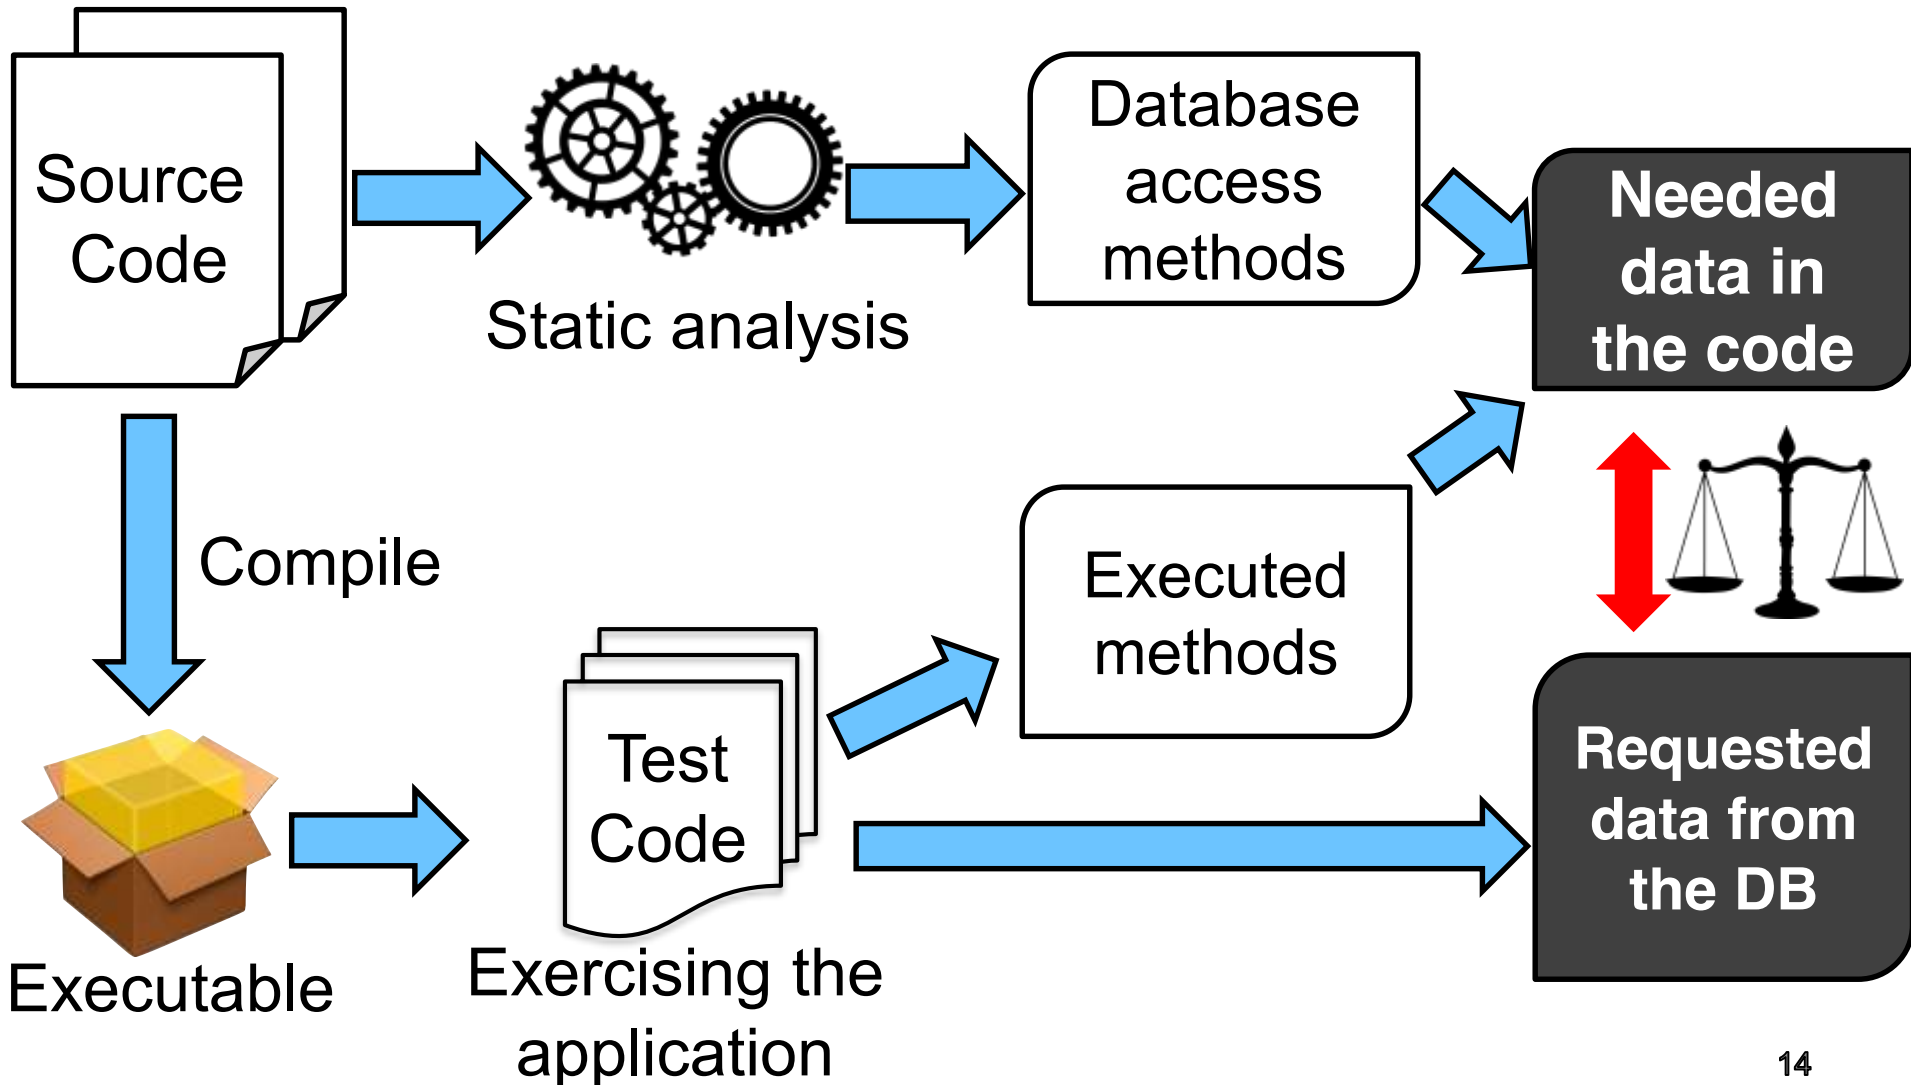
Our approach for finding redundant data anti-patterns



Understanding application execution using bytecode instrumentation



Our approach for finding redundant data anti-patterns



Finding redundant data anti-patterns by comparing needed and requested data

```
<transaction>  
  <methodCall>  
    user.getUserName()  
  </methodCall>
```

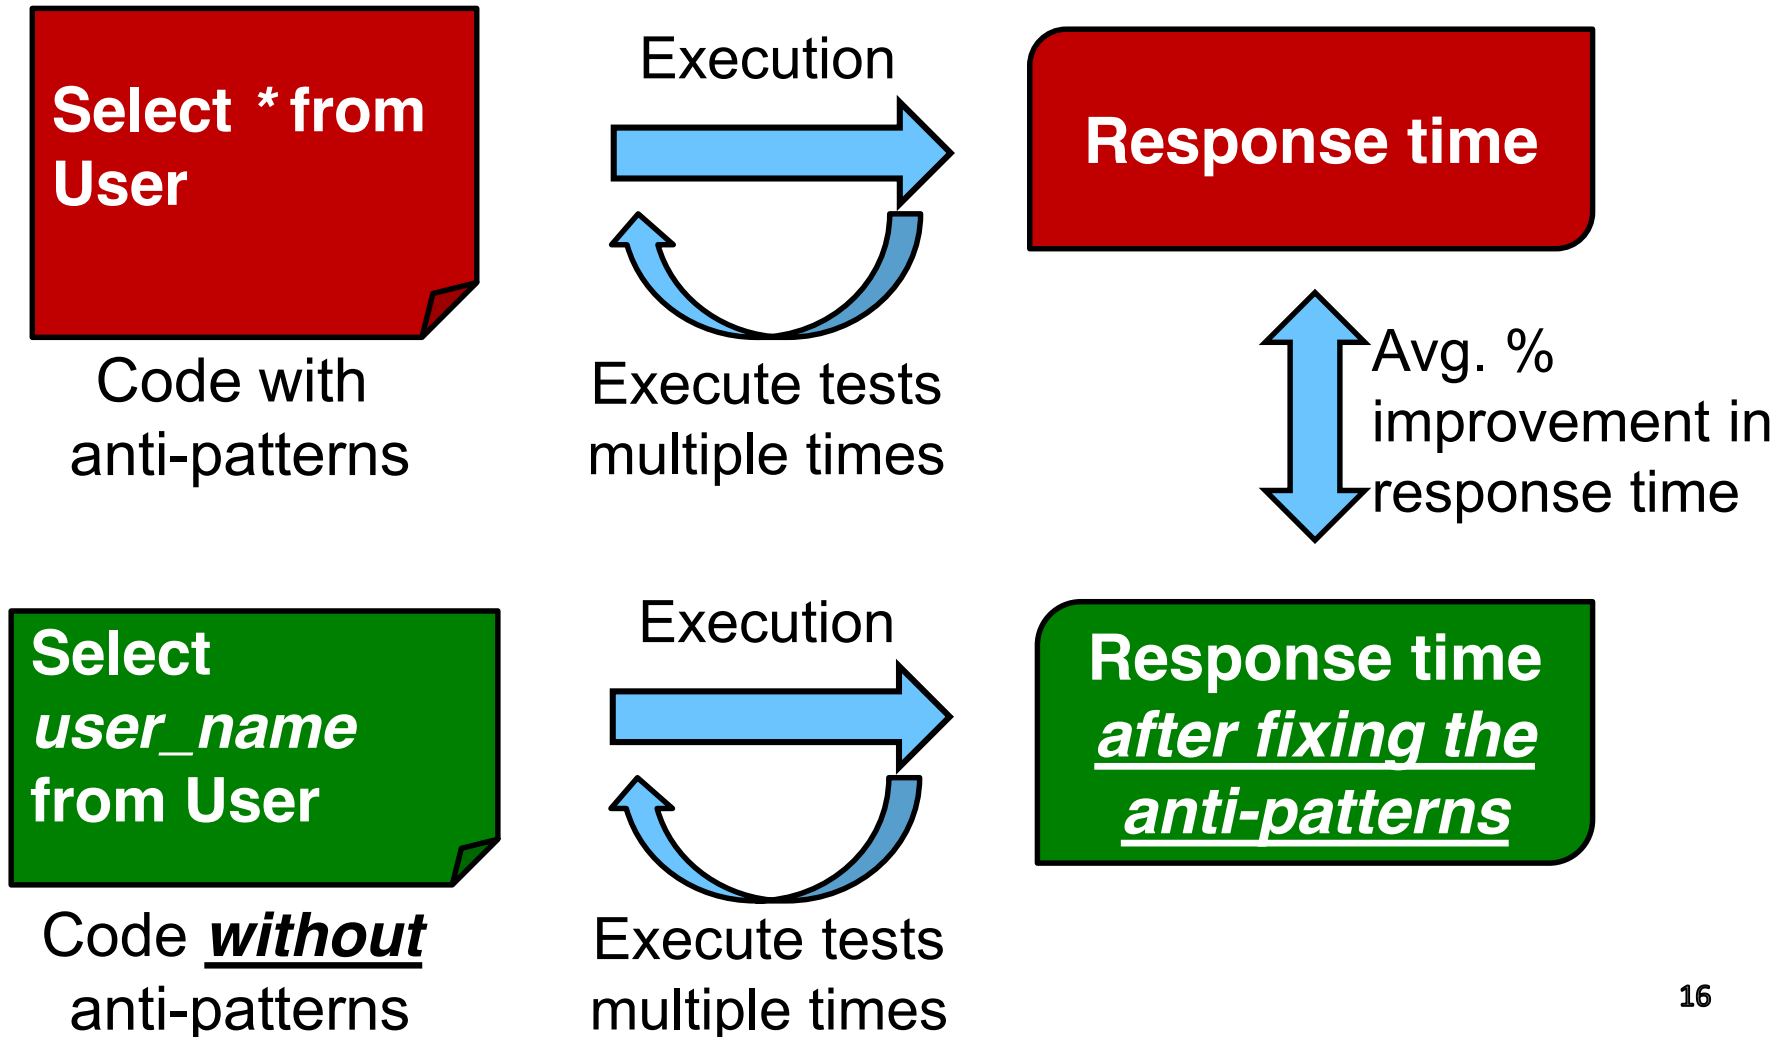
Only needs *name* in the application logic.

```
<ORM-generated SQL>  
  select u.id, u.name, u.address, u.phone_number from  
  User u where u.id = 1  
</ORM-generated SQL>  
</transaction>
```

The requested **id**, **address**, **phone number** are not needed.

We detect redundant data anti-patterns for all types of database access (e.g., select, update, join).

Assessing redundant data anti-pattern impact by fixing the anti-patterns



Studied applications



Large open-source
e-commerce system
> 1,700 files
> 206K LOC




Enterprise application
> 5,000 files
> 10M LOC



Open-source pet clinic
application
51 files
3.3K LOC

Removing redundant data anti-patterns improves response time




Slow database access is often the performance bottleneck



3

Redundant data is caused by ORM's lack of knowledge on the business logic



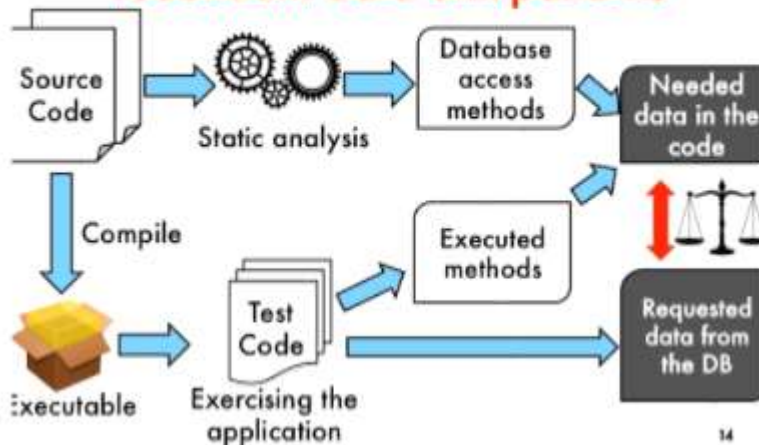
We need to understand how ORM APIs are used in the application in order to detect redundant data anti-patterns.

7

<http://petertsehsun.github.io/>

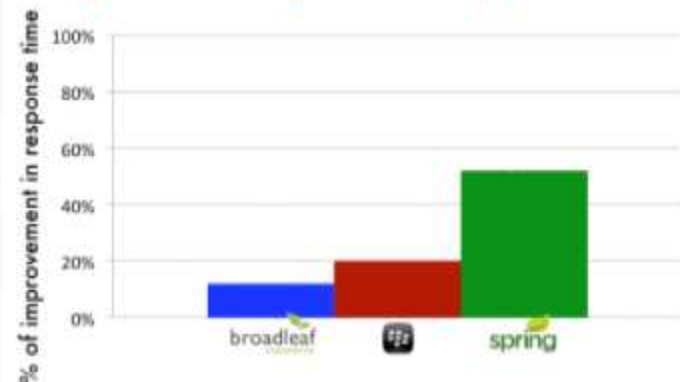
 @petertsehsun

Our approach for finding redundant data anti-patterns



14

Removing redundant data anti-patterns improves response time



18